ENTITY LINKING WITH A KG

Craig Knoblock

DSCI 558/CSCI 663

What is entity linking?

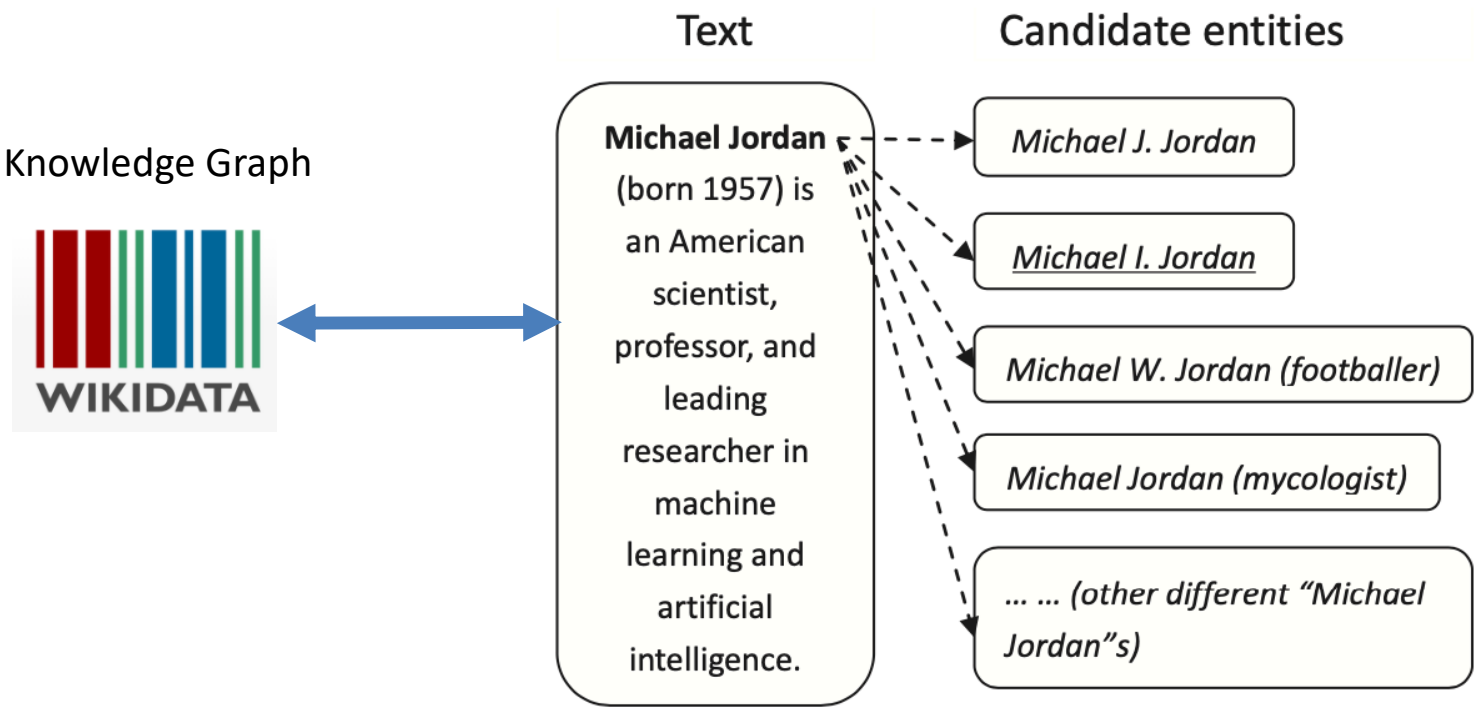

Sometimes called Named Entity Disambituation (NED)

From Shen et al., Entity Linking with a Knowledge Base

Task Description

- Given a KB with a set of entities E and a text collection with a set of entity mentioned M:
  goal is to map each m ϵ M to the corresponding e ϵ E

- Some mentions may be unlinkable because the corresponding e is not in the KB
  - In that case m should be labeled as NIL

- This task is typically preceded by Named Entity Recognition to find the entity mentions
  - Commonly implemented using a conditional random field (CRF) using in information extraction
  - Publicly available tools for NER: Stanford NER, OpenNLP, and LingPipe

Relationship of Entity Linking to Other Methods

- Entity coreference resolution
  - Determining the entity mentions within a document and across documents
  - Reduces to this case if there is no knowledge base
  - Entity linking is able to leverage the other information in the KB
- Word sense disambiguation
  - Determine the sense of a word from a catalog of senses e.g., Wordnet
  - Assumes the catalog of senses is complete
    - Not the case in entity linking
- Record linkage
  - Compare different records using a set of attributes using similarity scores
  - Entity linking there may be no other attributes or attributes missing and/or difficult to extract

General Approach to Entity Linking

Applications

- Information extraction
- Information retrieval
- Content analysis
- Question answering
- Knowledge base population

Candidate Entity Generation

- Name dictionary built from Wikipedia:
  - Entity pages – one for each entity
  - Redirect pages – aliases for an entity
  - Disambigutation pages – lists all of the entities with the same name
  - Bold phrases -- in the first paragraph are often aliases
  - Hyperlinks – use the anchor text to linked articles
- Exact matching and partial matching against any of these names
  - Some methods are do some form of spell checking
- Surface form expansion
  - E.g., DOD and Department of Defense both appear in the same document, DOD would get expanded
- Google search – select top Wikipedia pages

Candidate Entity Ranking

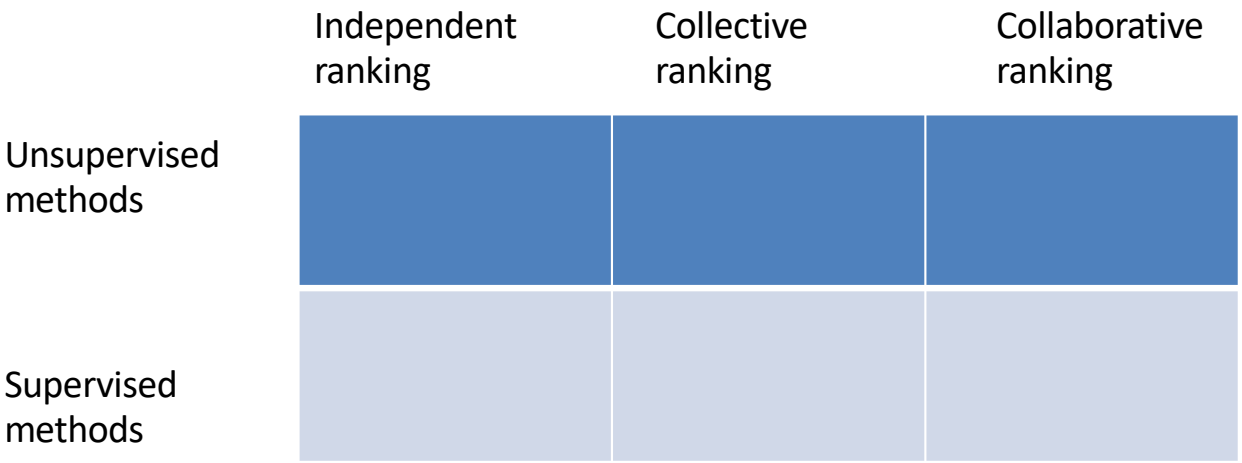

Entity ranking features (context independent)

- Name string comparison
  - Whether the entity mention exactly matches the candidate entity name.
  - Whether the candidate entity name starts with or ends with the entity mention.
  - Whether the candidate entity name is the prefix of or postfix of the entity mention.
  - Whether the entity mention is wholly contained in the candidate entity name, or vice-versa.
- Entity popularity
  - Count of the entity references
  - In-degree, out-degree, and length of page
- Entity type
  - Extracted from the NER

Entity ranking features (context dependent)

- Bag of words – uses TF-IDF
- Concept vector – using cosine similarity or embeddings
- Coherence between mapping entities
  - E.g., the link structure of Wikipedia

Supervised entity ranking

- Binary classification
  - uses ML method such as SVM
  - Multiple entities could be positive
- Learning to rank
  - learns a total order on all pairwise comparisons
- Probabilistic methods
  - Using graphical models
- Graph-based approaches
  - Graph-based collective entity linking
- Ensemble methods
  - Any combination of the above

Unsupervised entity ranking

- Vector-space-model (VSM) based methods
  - Compares the vector model of the entity article with the candidate entity
- IR-based methods
  - Similar to what search engines do
- Graph embeddings
  - Exploits the collective structure

Unlinkable Mention Prediction

- Approaches
  - Ignore the problem -- assume all entities are linkable
  - Only return NIL if no candidates are found
  - Assign scores to a match and require the match to be above a threshold
  - Learn a binary classifier to decide if an entity link is correct
    - Typically employ an SVM across a variety of features
  - Incorporate into the entity ranking process where NIL is a choice
    - If NIL is the top-ranked choice, the entity is unlinkable

Evaluation

$$precision = \frac{|\{correctly\ linked\ entity\ mentions\}|}{|\{linked\ mentions\ generated\ by\ system\}|}$$

$$recall = \frac{|\{correctly\ linked\ entity\ mentions\}|}{|\{entity\ mentions\ that\ should\ be\ linked\}|}$$

$$F_1 = \frac{2 \cdot precision \cdot recall}{precision + recall}$$

Summary

- Entity linking is the problem of identifying the corresponding entity in a KG
- Researchers have linked to Wikipedia, DBPedia, Wikidata, Yago, etc
- Typically, three steps
  - Candidate entity generation
    - Performed with a dictionary or google search
  - Candidate entity ranking
    - Range of unsupervised and supervised methods for performing these task
  - Unlinkable mention prediction
    - Typically done with a threshold or learned classifier
- Useful in a variety of applications
  - From content analysis to KG population