



Fast Entity Linking via Graph Embeddings

Alberto Parravicini

Rhicheek Patra

Davide Bartolini

Marco Santambrogio

2019-05-{17-31}, NGCX



POLITECNICO
MILANO 1863



POLITECNICO MILANO 1863

ORACLE®

Entity Linking

Entity Linking (EL): connecting words of interest to unique identities (e.g. Wikipedia Page)



Use Cases

Component of applications that require high-level representations of text:

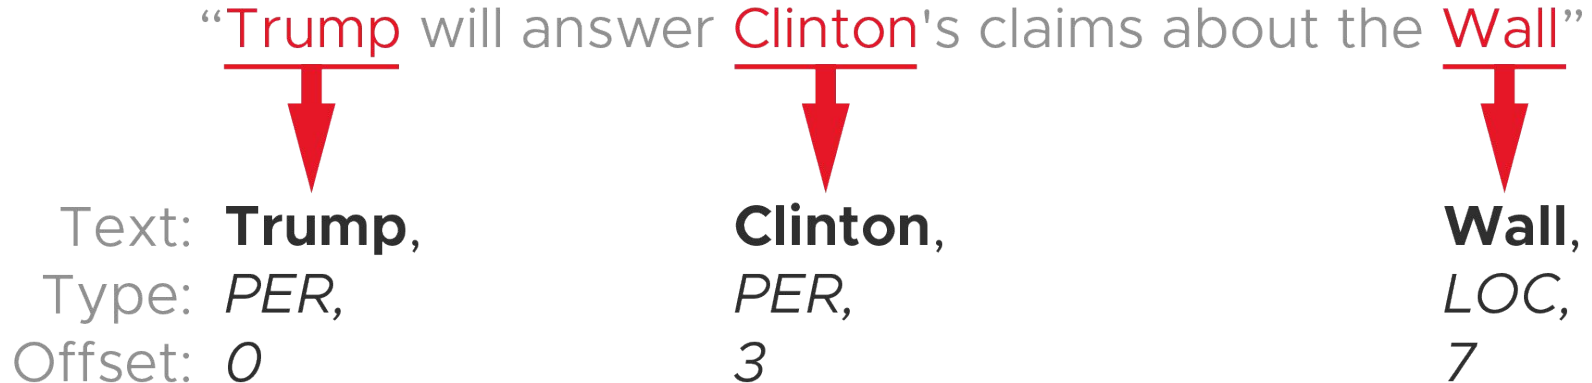
1. **Search Engines**, for semantic search
2. **Recommender Systems**, to retrieve documents similar to each other
3. **Chat bots**, to understand intents and entities

The EL Pipeline (1/2)

An EL system requires 2 steps:

1. Named Entity Recognition (NER): spot **mentions** (a.k.a. Named Entities)

- High-accuracy in the state-of-the-art^[1]

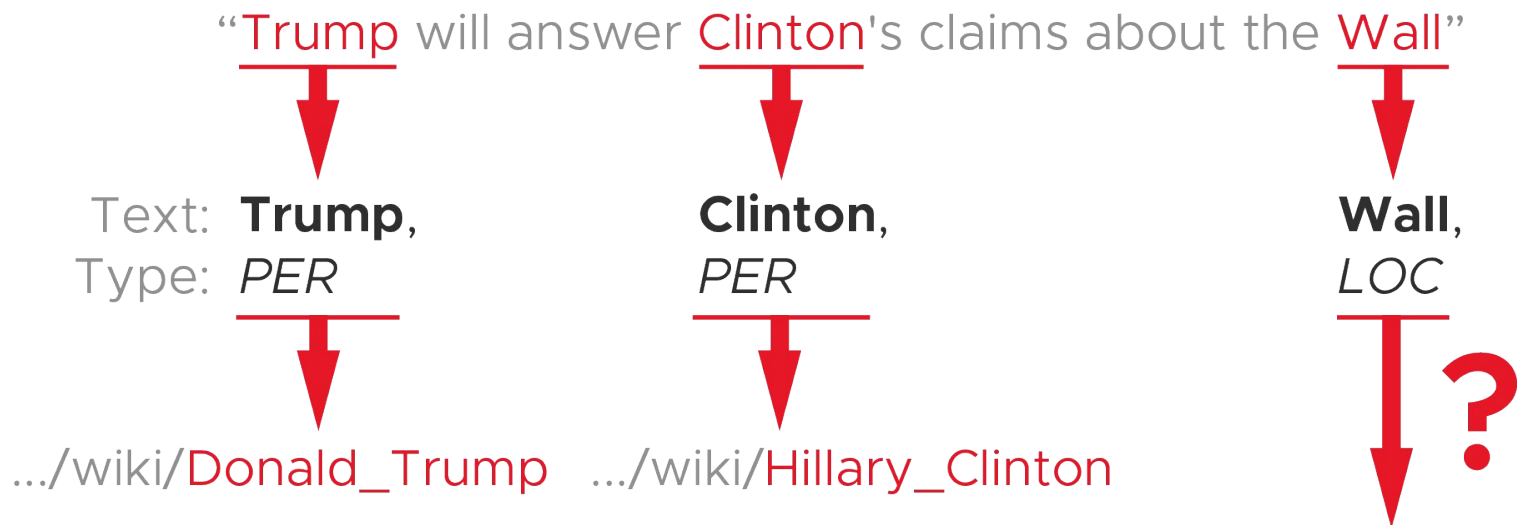


[1] Huang, Zhiheng, Wei Xu, and Kai Yu. "Bidirectional LSTM-CRF models for sequence tagging."

The EL Pipeline (2/2)

An EL system requires 2 steps:

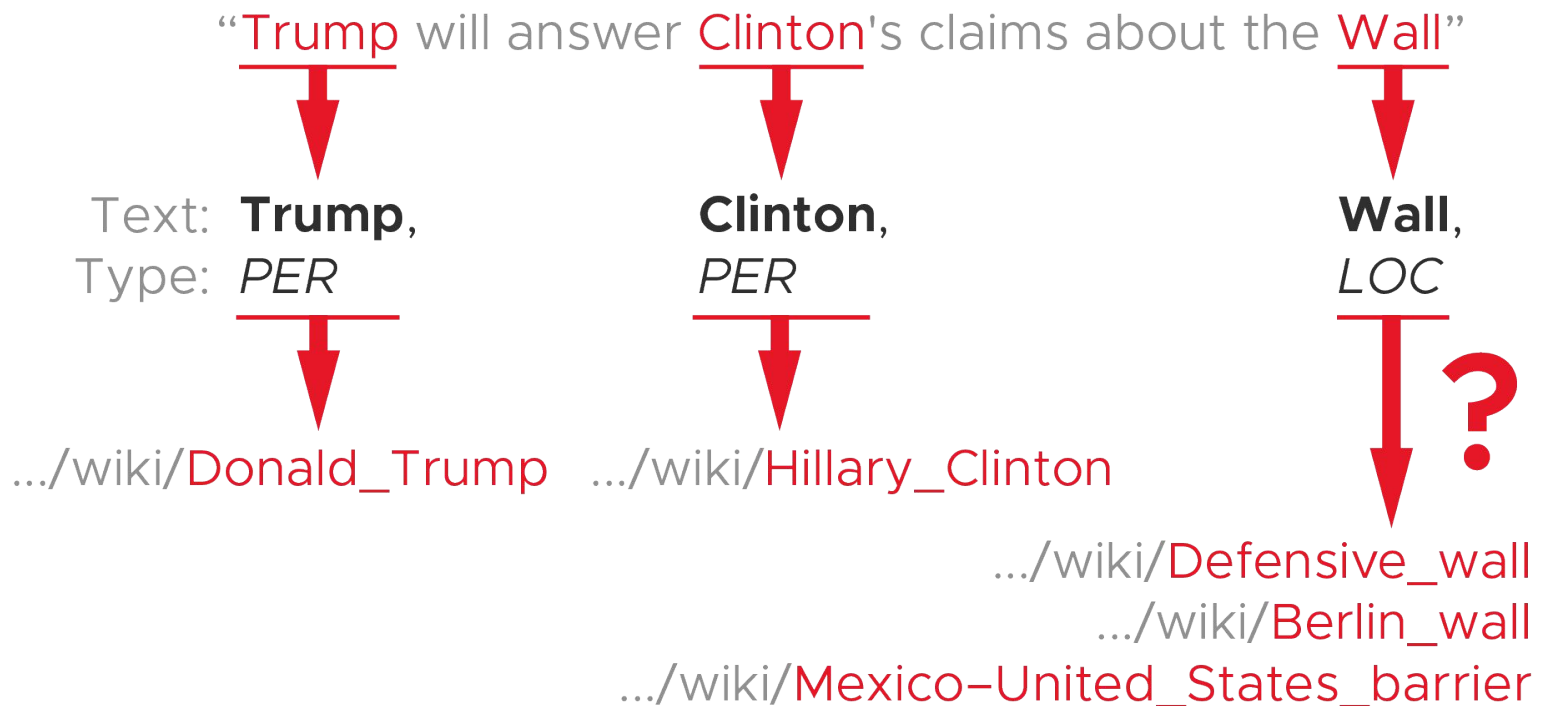
2. Entity Linking: connect mentions to entities



The EL Pipeline (2/2)

An EL system requires 2 steps:

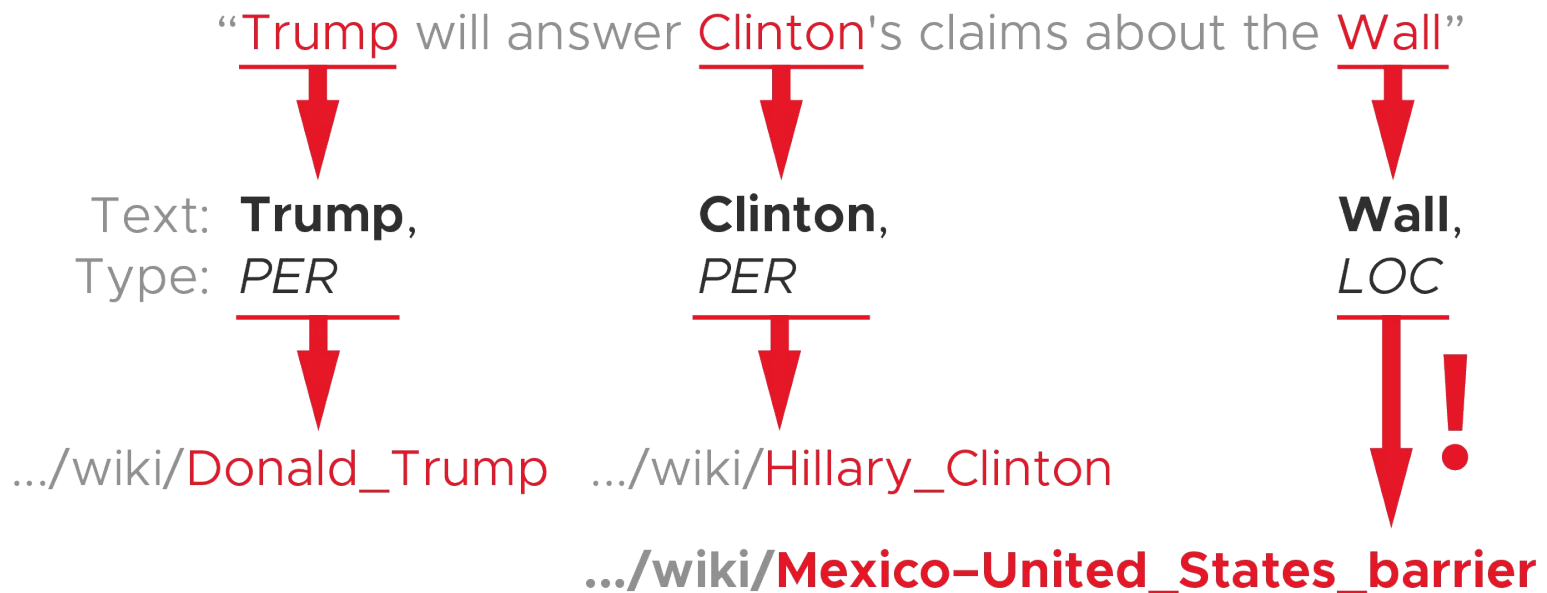
2. Entity Linking: connect mentions to entities



The EL Pipeline (2/2)

An EL system requires 2 steps:

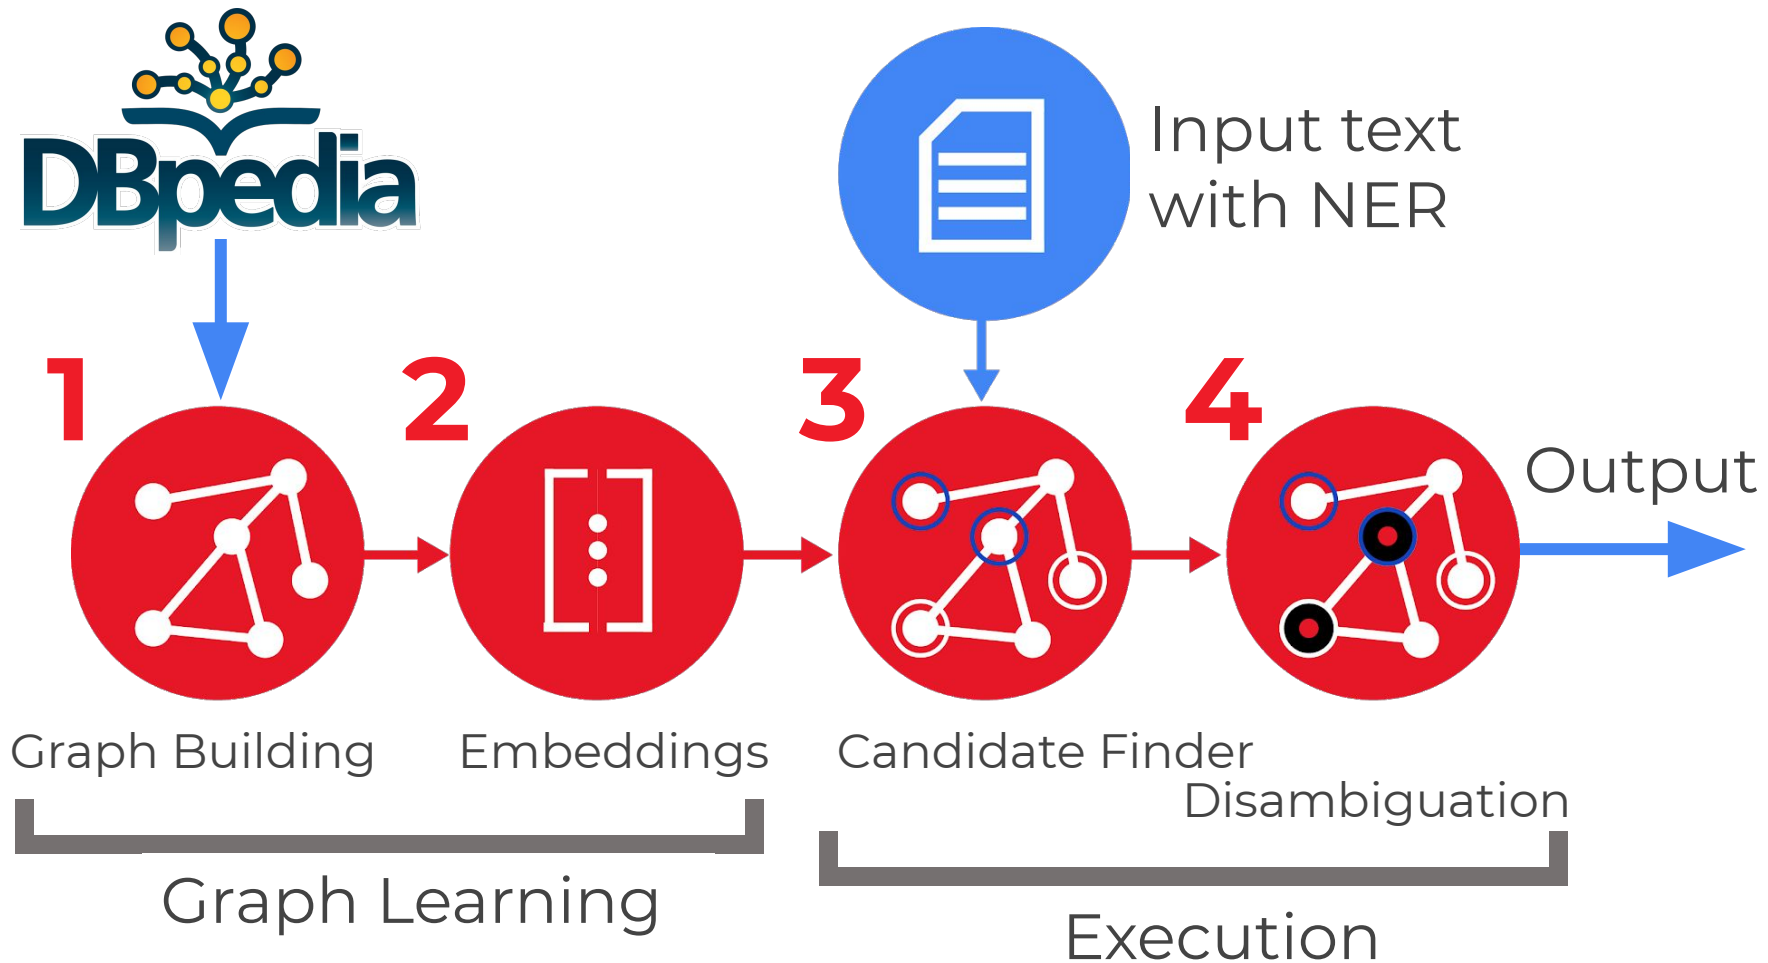
2. Entity Linking: connect mentions to entities



Our contributions

- Novel **unsupervised** framework for EL
 - No dependency on NLP
- First EL algorithm to use **graph embeddings**
 - Accuracy similar to supervised SoA techniques
- Highly scalable and **real-time** execution time
 - < 1 sec to process text with 30+ mentions

Our EL Pipeline



Graph Creation

We obtain a large graph from DBpedia

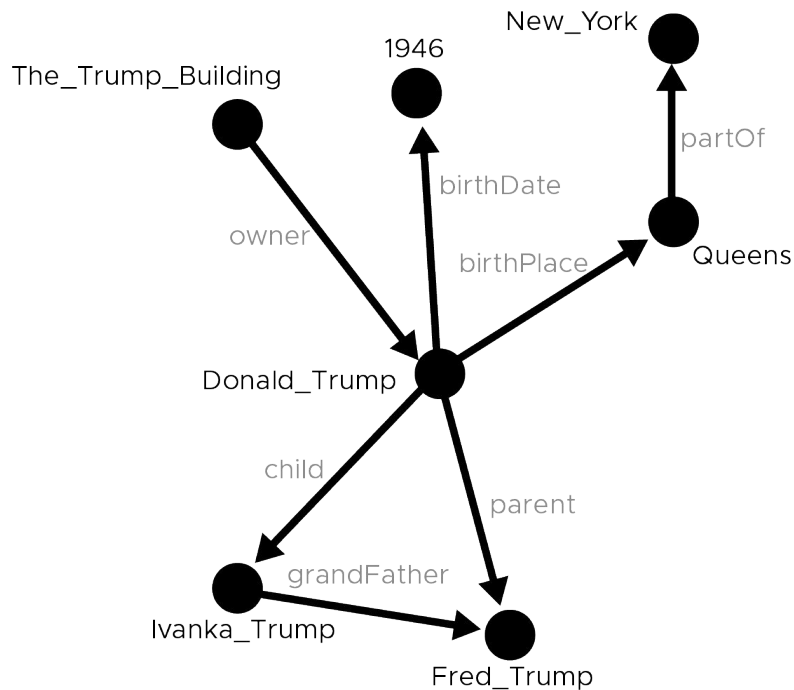
- All the information of Wikipedia, stored as triples
- 12M entities, 170M links




Graph Creation

From DBPedia, we build two graphs

Property Graph




Redirects Graph



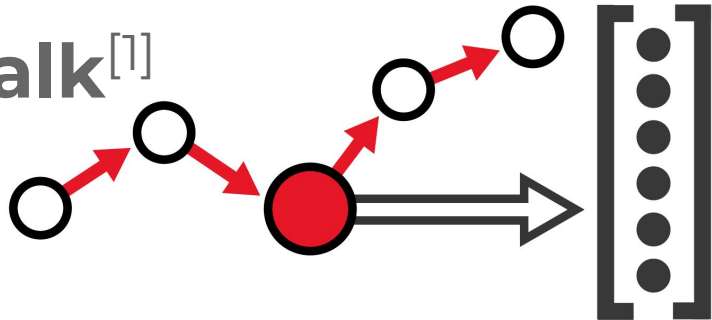
Embeddings Creation (1/2)

- Graph embeddings encode **vertices as vectors**
 - “Similar” vertices have “similar” embeddings
- **Idea:** entities with the same context should have low embedding distance



12

Embeddings Creation (2/2)

- In our work, we use **DeepWalk**^[1]
- Like **word2vec**^[2], it leverages **random walks** (i.e. vertex sequences) to create embeddings
 - Embedding size 170, walk length 8
- DeepWalk uses only the graph **topology**
 - Simple baseline, we can use better algorithms and leverage graph features

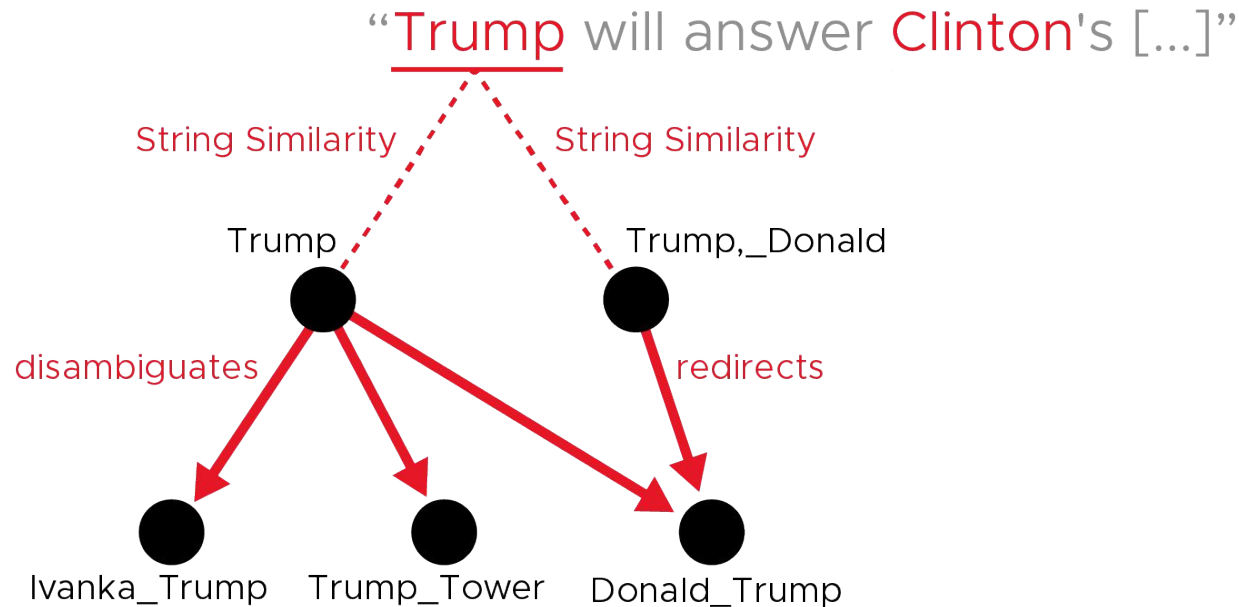
[1] Bryan Perozzi et al. 2014. Deepwalk

[2] Tomas Mikolov et al. 2013. Distributed representations of words and phrases and their compositionality

Candidates Finder

Idea: for each mention, select a few candidate vertices with index- based **string similarity**


- Solve ambiguity following **redirect** and **disambiguation** edges



Candidates Finder


Idea: for each mention, select a few candidate vertices with index- based **string similarity**

- Solve ambiguity following **redirect** and **disambiguation** edges



Disambiguation (1/3)

- We want to pick the “best” candidate for each mention
 - In a “good” solution, candidates are related to each other (e.g. *Donald Trump*, *Hillary Clinton*)
- **Observation**: a good tuple of candidates has embeddings close to each other
- Evaluating all combinations is infeasible
 - 10 mentions with 100 candidates $\rightarrow 100^{10}$



Disambiguation (2/3)

- We use an heuristic state-space search algorithm to maximize:


$$Best\ Tuple = \operatorname{argmax}_T \left(\sum_{t_i \in T} \boxed{Local(t_i)} + \boxed{Global(T)} \right)$$

Sum of string similarities

Sum of embedding cosine similarities w.r.t. tuple mean

$$\bar{\mathbf{e}}(T) = \frac{1}{|T|} \sum_{t_i \in T} \mathbf{e}(t_i)$$

$$Global(T) = \sum_{t_i \in T} \frac{\langle \mathbf{e}(t_i), \bar{\mathbf{e}}(T) \rangle}{\|\mathbf{e}(t_i)\|_2 \|\bar{\mathbf{e}}(T)\|_2}$$



Disambiguation (3/3)

- Iterative state-space heuristic

Function `optimizer(candidates):`

`T = find_initial_state(candidates)`

while *stop condition not met* **do**

*// Create num_children new tuples by modifying
random elements of T.*

`T = new_tuples(T, num_children)`

`T = optimize_tuple(T)`

`s = compute_score(T)`

if *curr_score* \geq *best_score* **then**

`best_tuple = T`

`best_score = s`

end

end

return *best_tuple*, *best_score*

Greedy
iterative
procedure

$$\operatorname{argmax}_T \left(\sum_{t_i \in T} \operatorname{Local}(t_i) + \operatorname{Global}(T) \right)$$

Results: accuracy



- We compared against 6 SoA EL algorithms, on 5 datasets
- Our Micro-averaged **F1** score is comparable with SoA supervised algorithms

Data Set	Ours	DoSeR	WK	AIDA	WAT	BB	SL
ACE2004	0.84	0.90	0.83	0.81	0.80	0.56	0.71
AQUAINT	0.86	0.84	0.86	0.53	0.77	0.65	0.71
MSNBC	0.92	0.91	0.85	0.78	0.78	0.60	0.51
N3-Reuters	0.82	0.85	0.70	0.60	0.64	0.53	0.58
N3-RSS-500	0.72	0.75	0.73	0.71	0.68	0.63	0.62

WK is Wikifier, BB is Babelify, SL is Spotlight

Results: exec. time

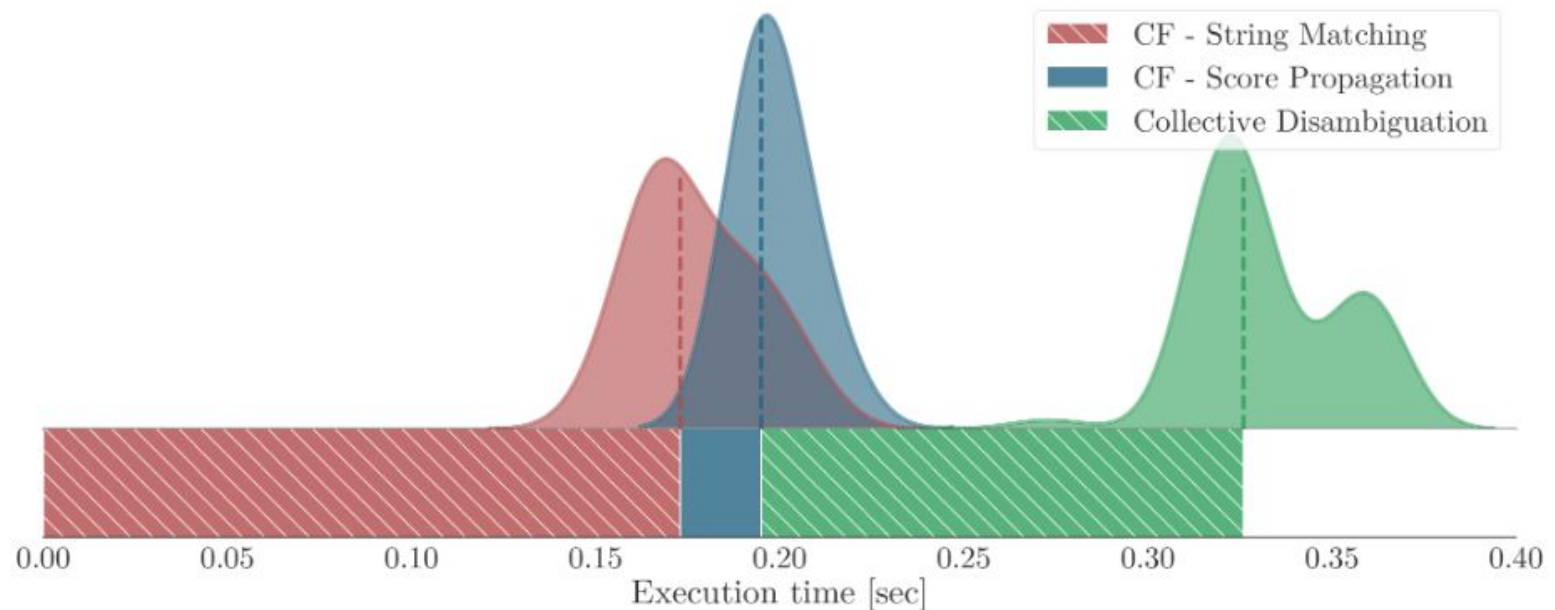
- Different settings enable **real-time EL**, with minimal loss in accuracy
 - E.g. number of iterations, early stop



Results: exec. time of single steps

- Execution time is well divided between **Candidate Finder** and **Disambiguation**

Time distribution - Fast Configuration



Thank you!

Fast Entity Linking via Graph Embeddings

- Novel unsupervised framework for EL
- First EL algorithm to use graph embeddings
 - Accuracy similar to supervised SoA techniques
- Real-time execution time

Alberto Parravicini, alberto.parravicini@polimi.it

Rhicheek Patra

Davide Bartolini

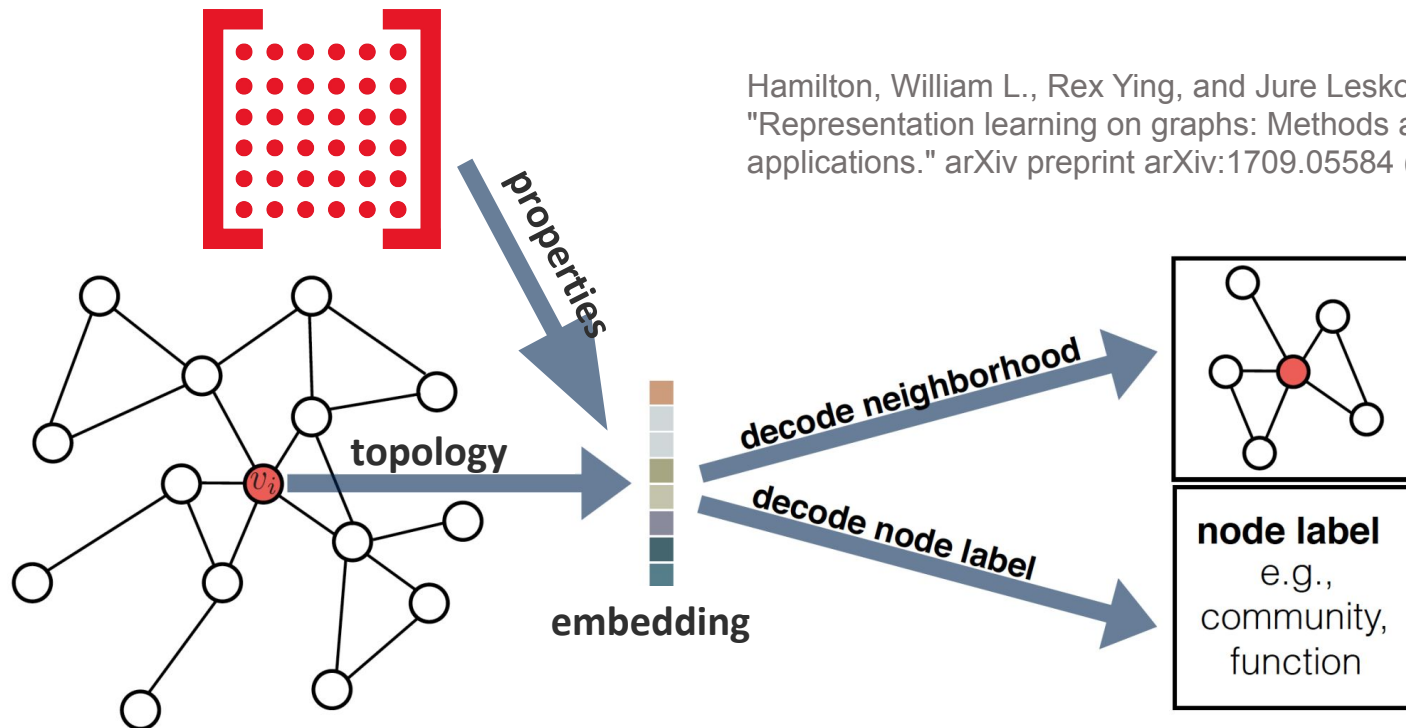
Marco Santambrogio

2019-05-{17-31}, NGCX



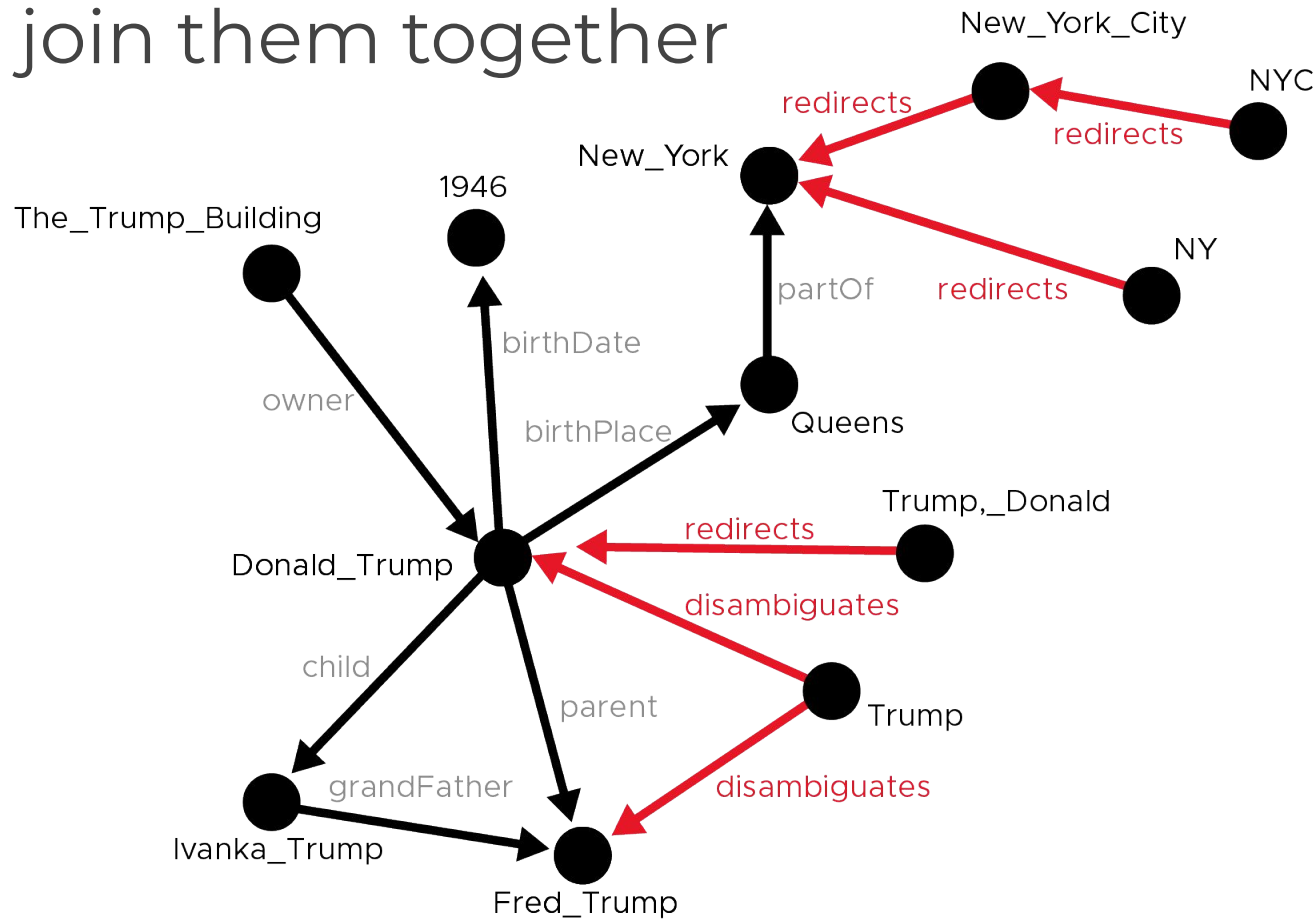
Embeddings

- Turn topology and properties of each vertex into a vector
- “Similar” vertices have “similar” embeddings



Graph Creation

... and join them together



Candidates Finder

Idea: for each mention, select a small number of candidate vertices with **string similarity**

- We use a simple **index-based** string search
- Fuzzy matching with **2-grams and 3-grams**
- This provide a simple baseline (60-70% accuracy)