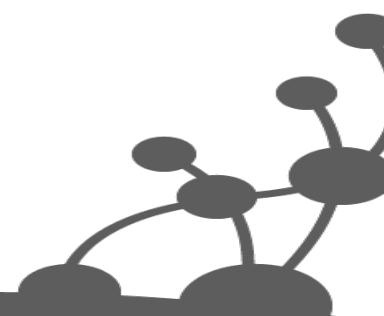
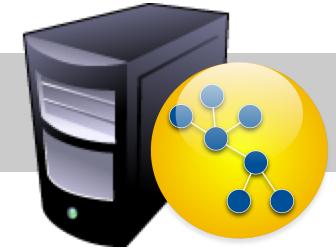




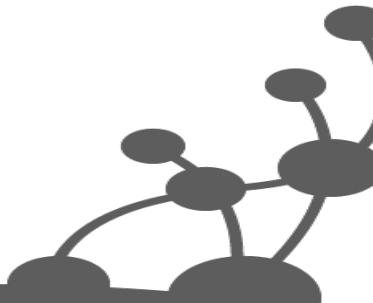
Data Modeling With Neo4j



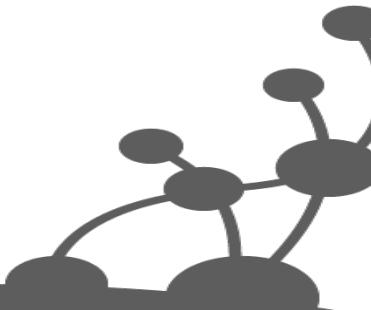
Topics



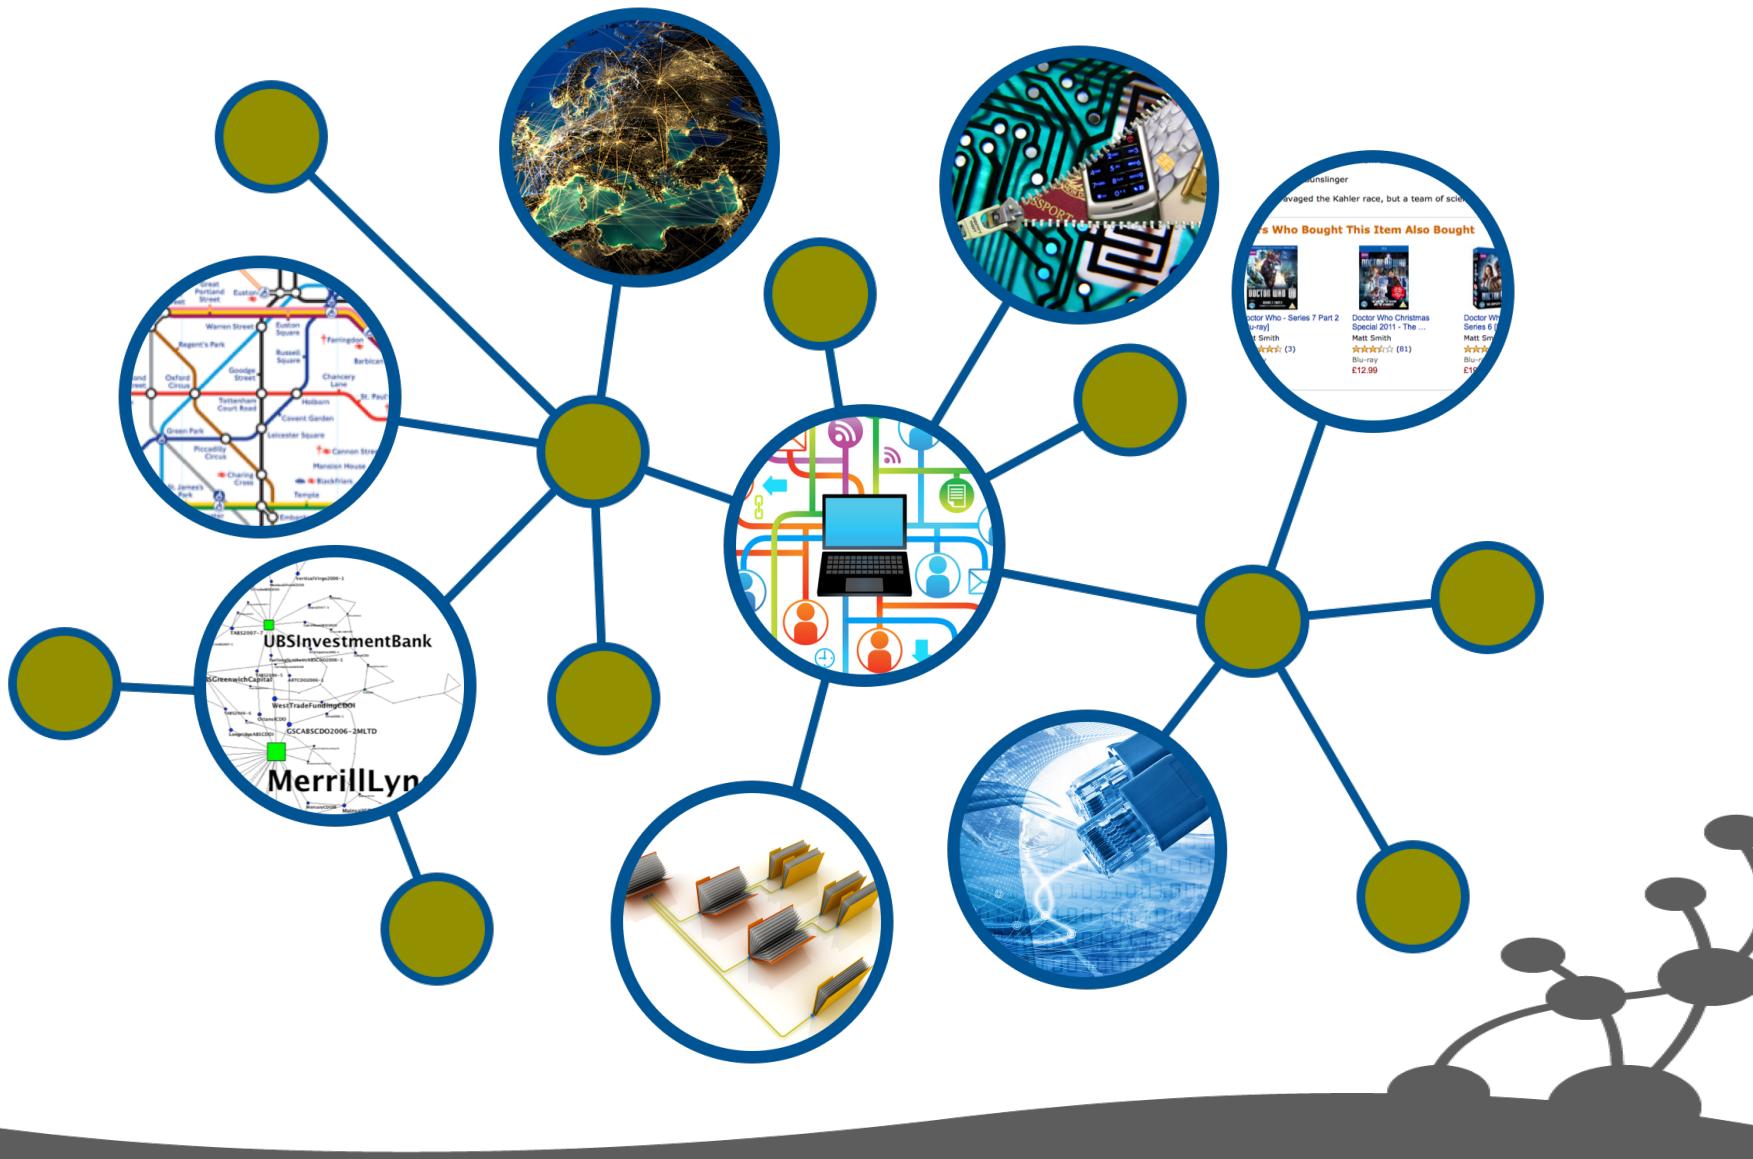
- Data complexity
- Graph model building blocks
- Quick intro to Cypher
- Modeling guidelines
- Common graph structures
- Evolving a graph model



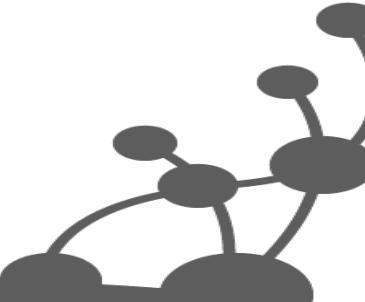
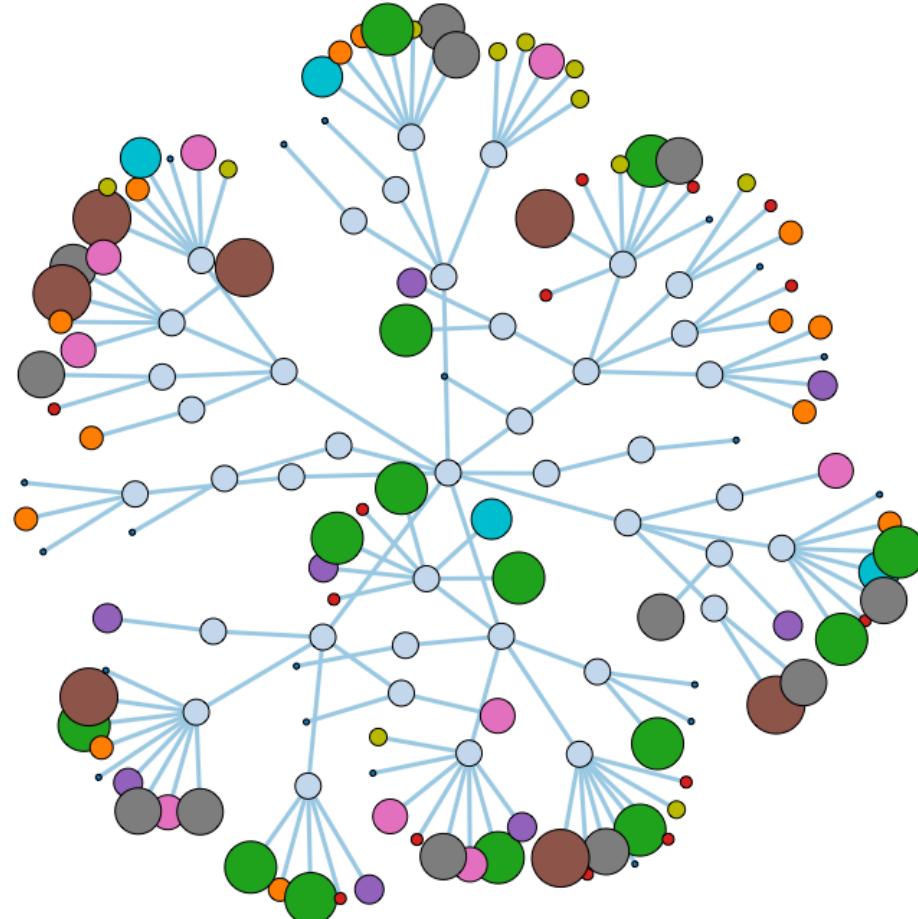
Addressing Data Complexity With Graphs



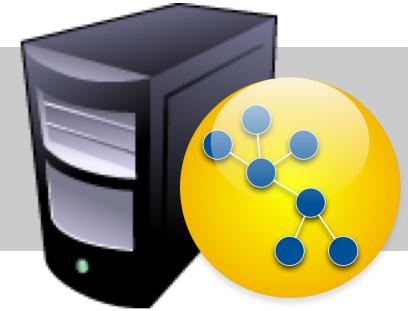
complexity = f(size, variable structure, connectedness)



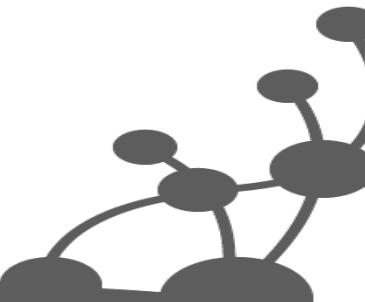
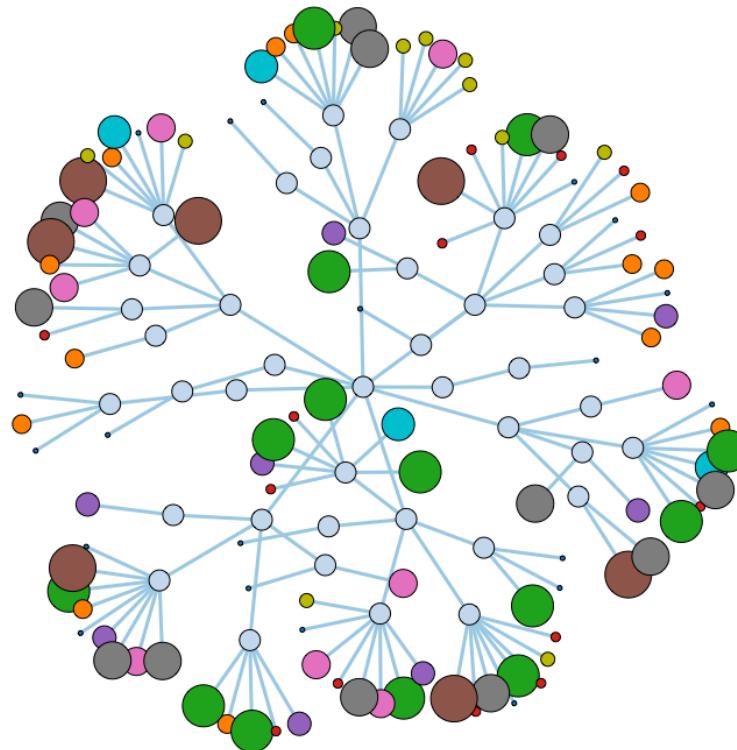
Graphs Are Everywhere



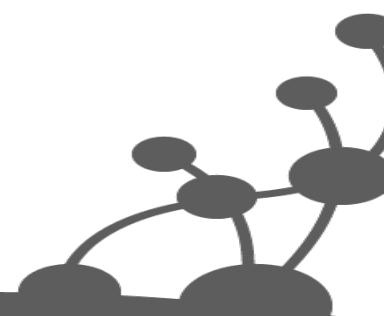
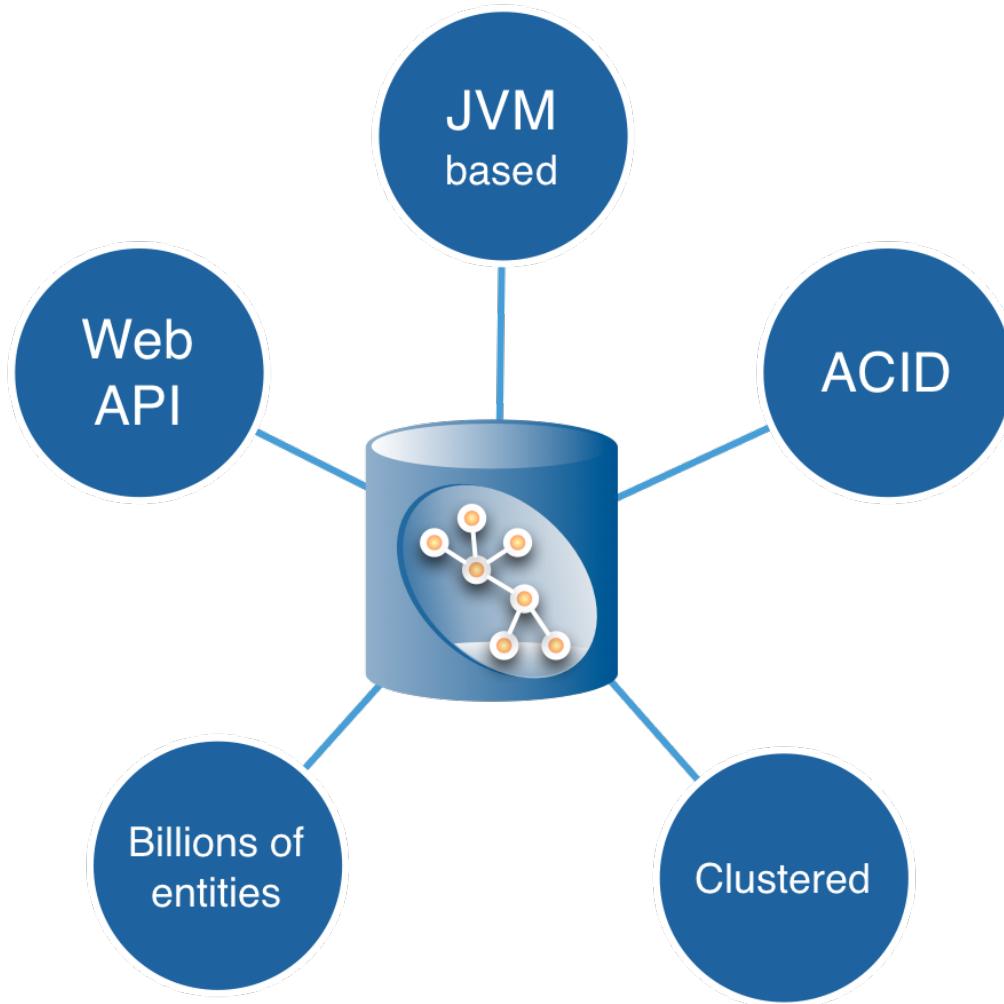
Graph Databases



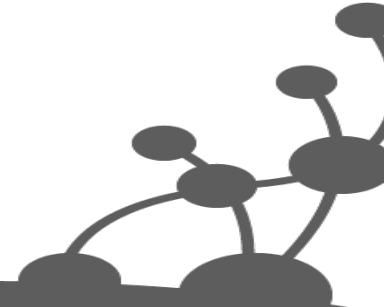
- Store
- Manage
- Query



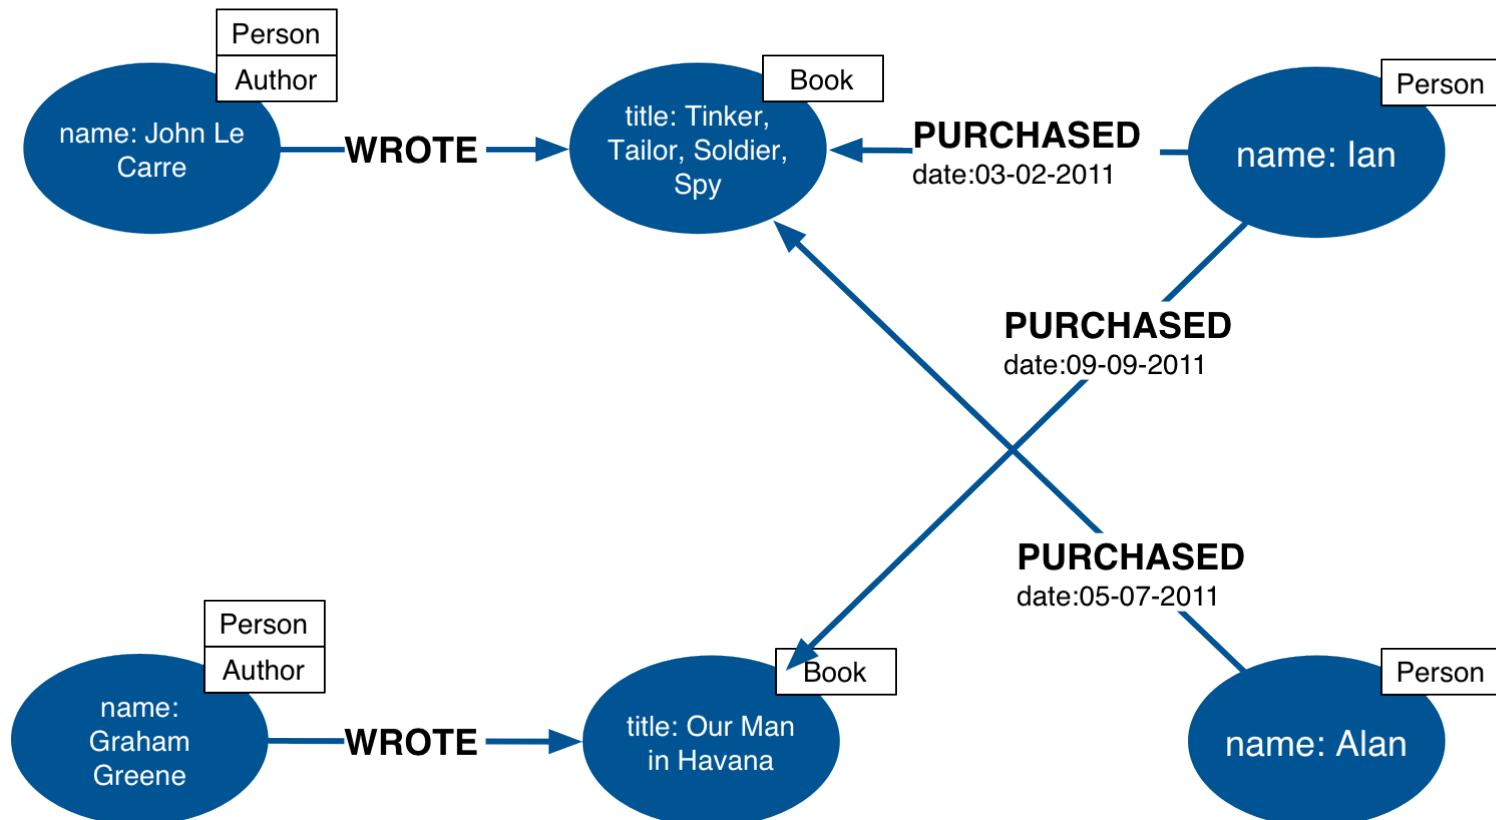
Neo4j is a Graph Database



Graph Model Building Blocks

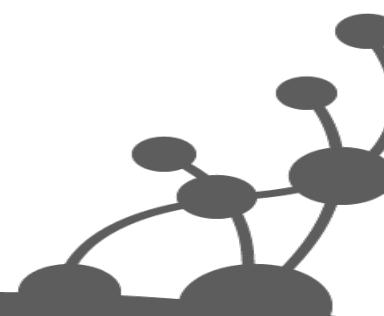


Labeled Property Graph Data Model

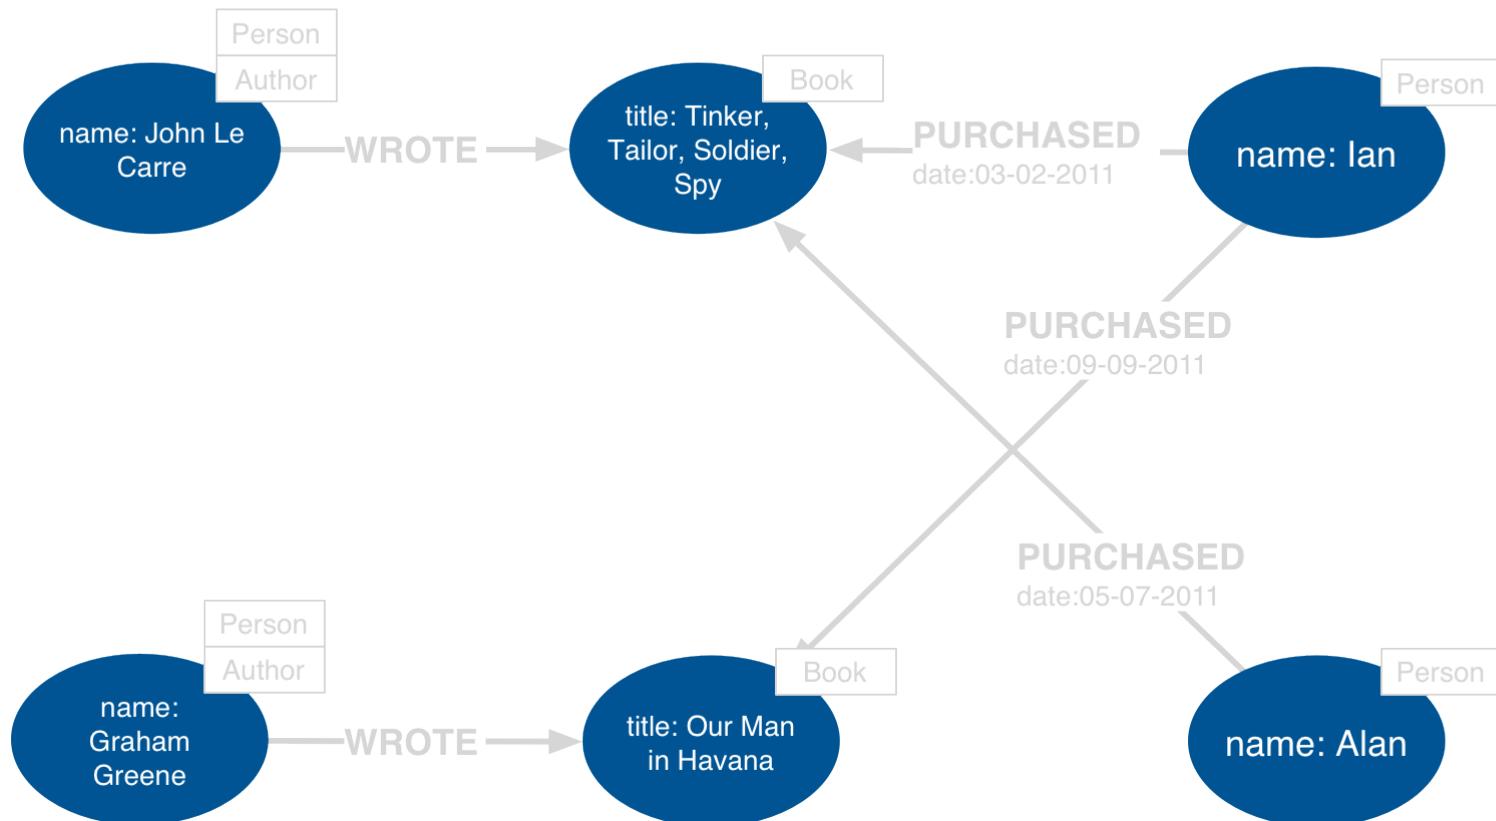


Four Building Blocks

- Nodes
- Relationships
- Properties
- Labels

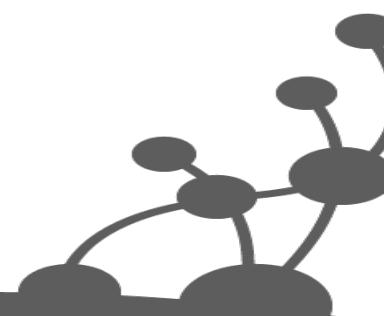


Nodes



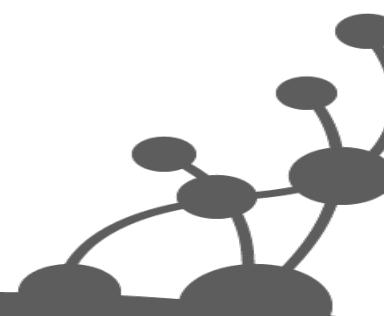
Nodes

- Used to represent *entities* and *complex value types* in your domain
- Can contain properties
 - Used to represent entity *attributes* and/or *metadata* (e.g. timestamps, version)
 - Key-value pairs
 - Java primitives
 - Arrays
 - null is not a valid value
 - Every node can have different properties

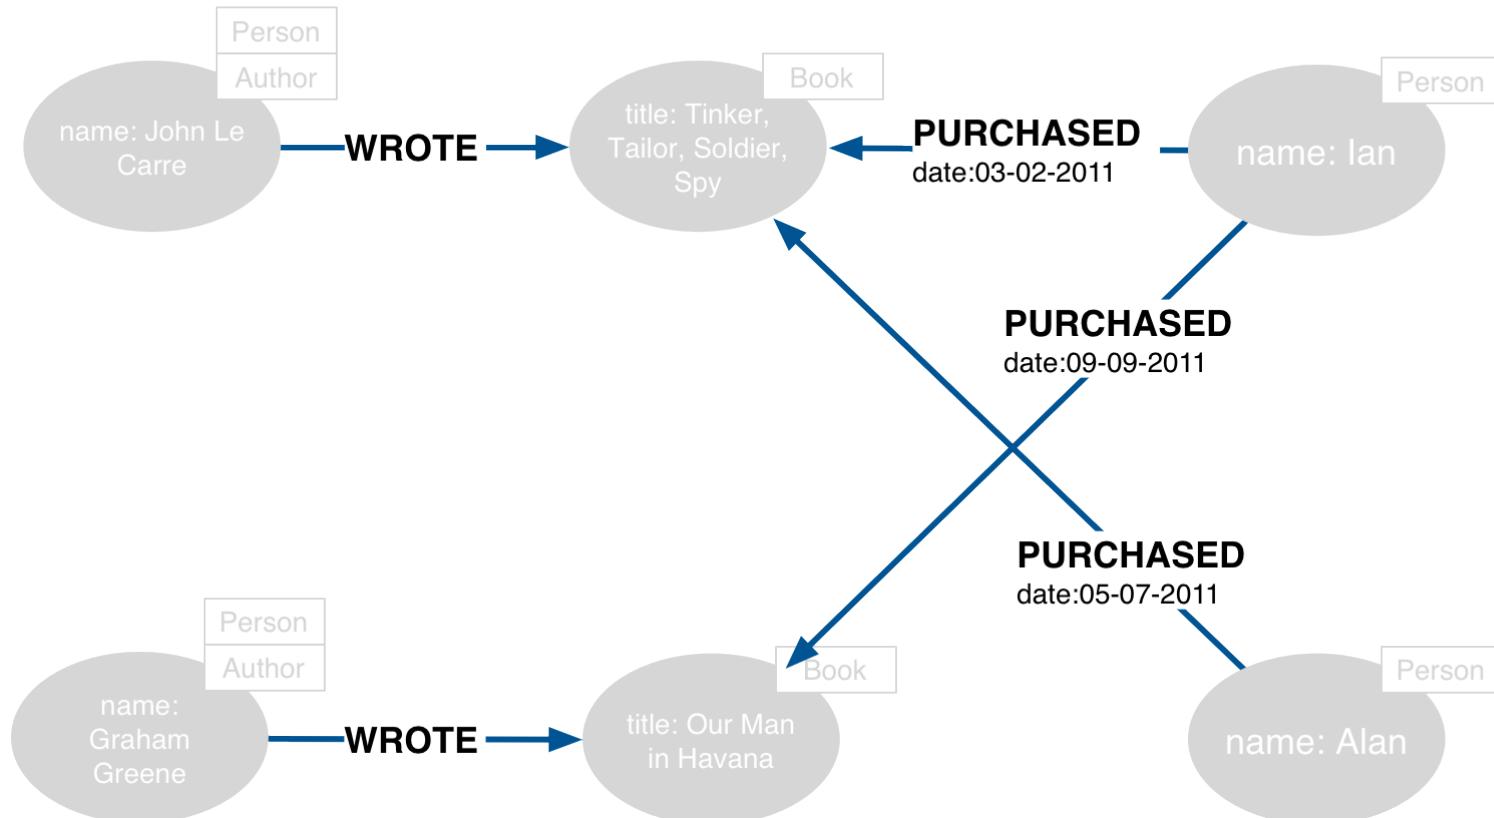


Entities and Value Types

- Entities
 - Have unique conceptual identity
 - Change attribute values, but identity remains the same
- Value types
 - No conceptual identity
 - Can substitute for each other if they have the same value
 - Simple: single value (e.g. colour, category)
 - Complex: multiple attributes (e.g. address)

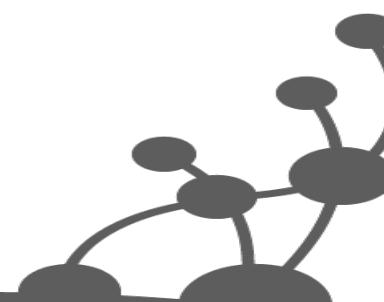


Relationships

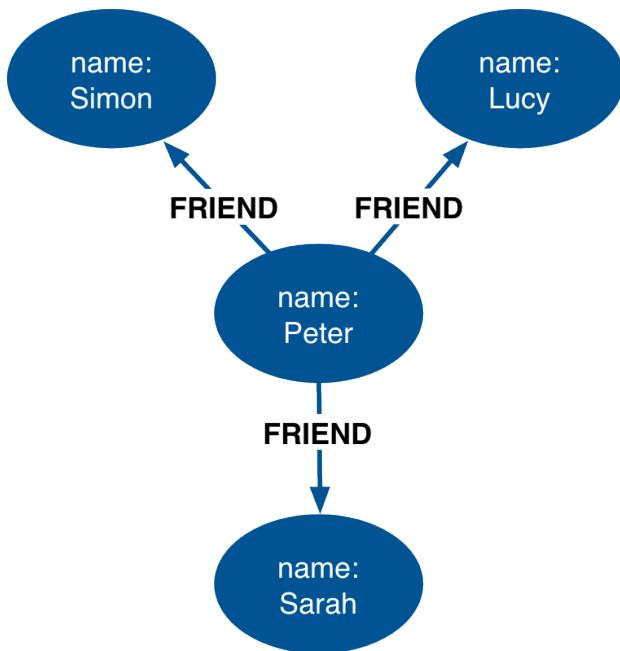


Relationships

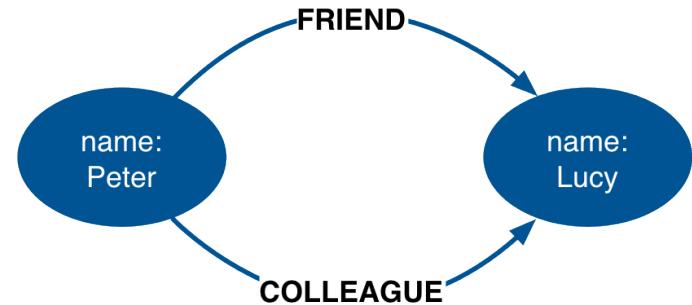
- Every relationship has a *name* and a *direction*
 - Add structure to the graph
 - Provide semantic context for nodes
- Can contain properties
 - Used to represent *quality* or *weight* of relationship, or *metadata*
- Every relationship must have a *start node* and *end node*
 - No dangling relationships



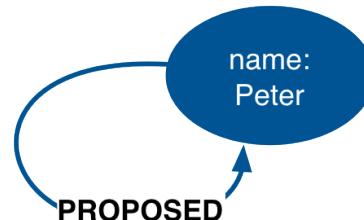
Relationships (continued)



Nodes can have more than one relationship



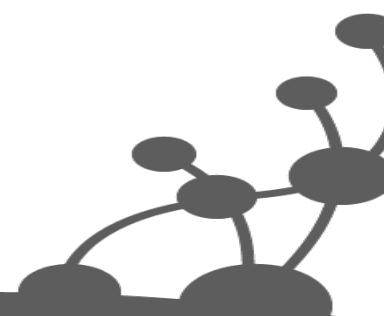
Nodes can be connected by more than one relationship



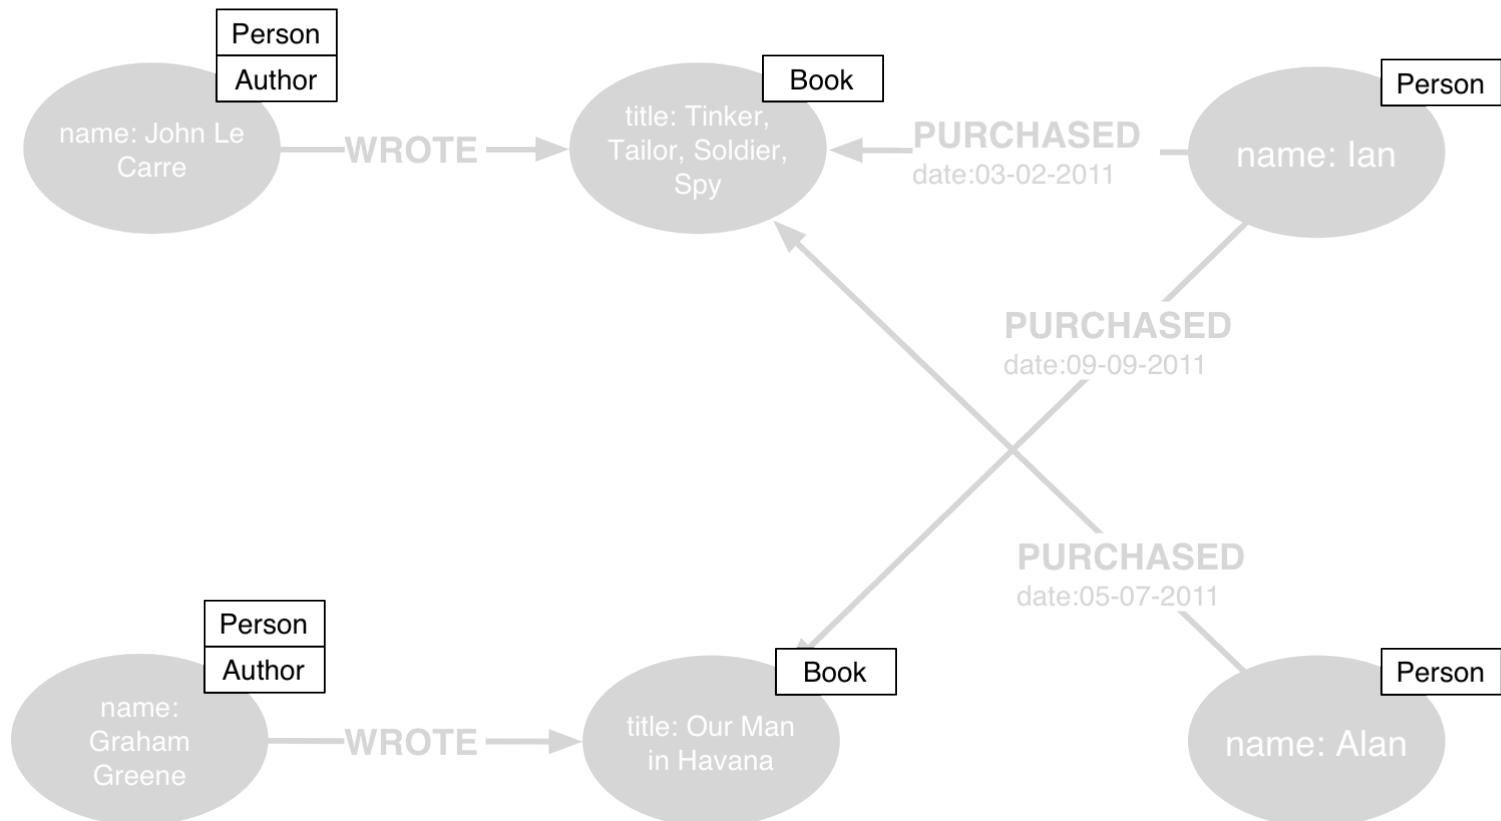
Self relationships are allowed

Variable Structure

- Relationships are defined with regard to node *instances*, not *classes* of nodes
 - Two nodes representing the same kind of “thing” can be connected in very different ways
 - Allows for structural variation in the domain
 - Contrast with relational schemas, where foreign key relationships apply to all rows in a table
 - No need to use *null* to represent the absence of a connection

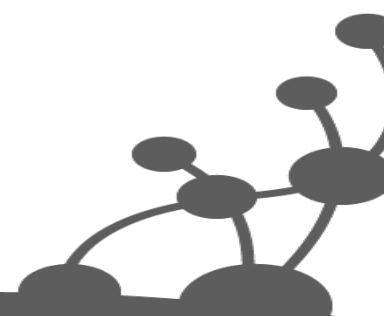


Labels



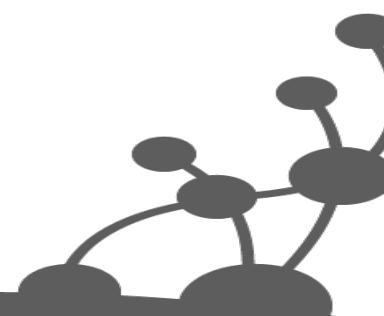
Labels

- Every node can have *zero or more* labels
- Used to represent *roles* (e.g. user, product, company)
 - Group nodes
 - Allow us to associate *indexes* and *constraints* with groups of nodes

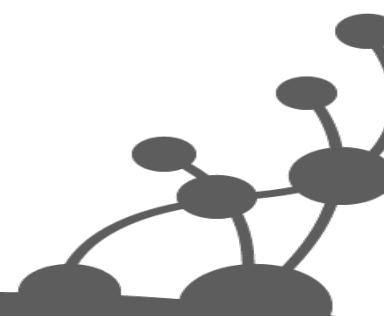


Four Building Blocks

- Nodes
 - Entities
- Relationships
 - Connect entities and structure domain
- Properties
 - Entity attributes, relationship qualities, and metadata
- Labels
 - Group nodes by role

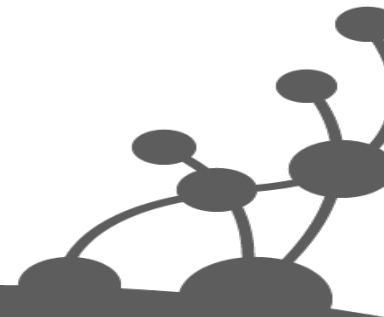


Cypher Query Language



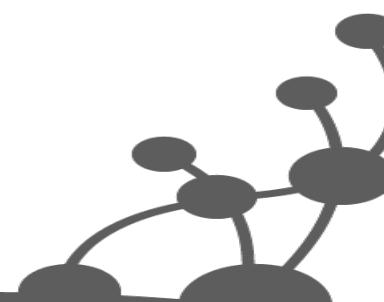
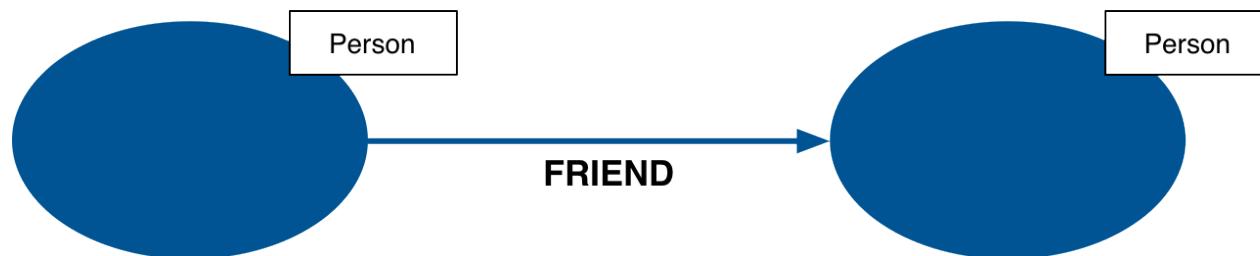
Nodes and Relationships

○ - -> ○



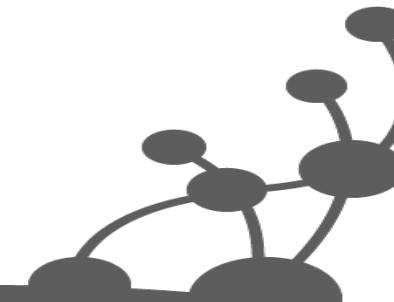
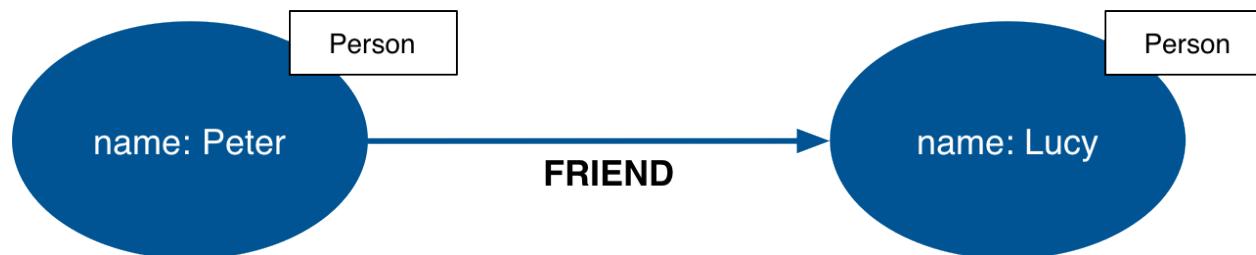
Labels and Relationship Types

(:Person)-[:FRIEND]->(:Person)



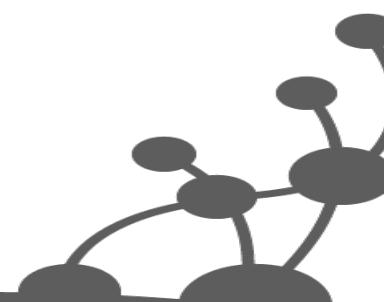
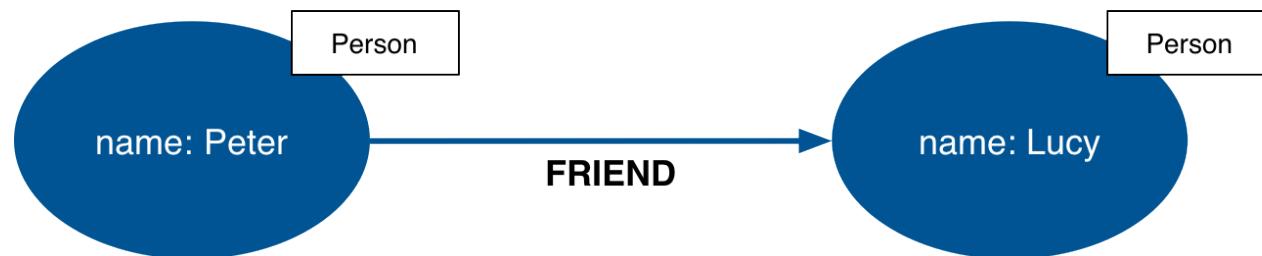
Properties

(:Person{name:'Peter'})-[:FRIEND]->(:Person{name:'Lucy'})



Identifiers

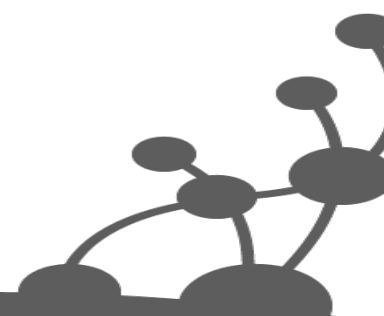
(p1:Person{name:'Peter'})-[r:FRIEND]-(p2:Person{name:'Lucy'})



Cypher

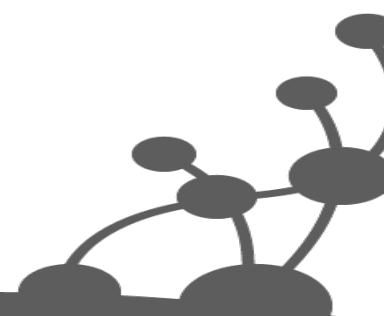
```
MATCH graph_pattern
RETURN results
```

<http://docs.neo4j.org/chunked/milestone/query-match.html>
<http://docs.neo4j.org/chunked/milestone/query-return.html>



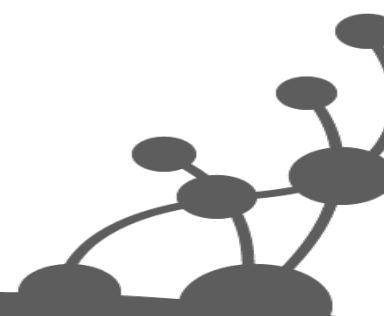
Example Query

```
MATCH (:Person{name:'Peter'})  
      -[:FRIEND]->(friends)  
RETURN friends
```



Find This Pattern

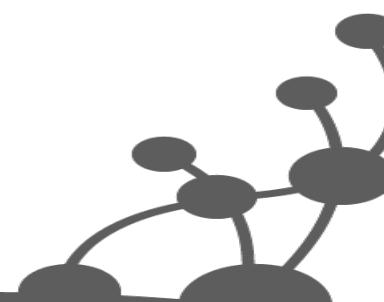
```
MATCH (:Person{name:'Peter'})  
      -[:FRIEND]->(friends)  
RETURN friends
```



Lookup Using Identifier + Label

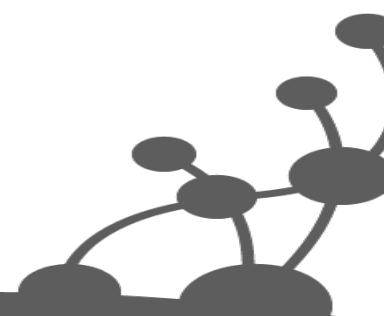
```
MATCH (:Person{name: 'Peter'})  
      -[:FRIEND]->(friends)  
RETURN friends
```

Search nodes labeled
'Person', matching on
'name' property



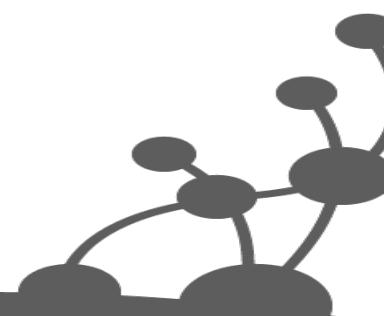
Return Nodes

```
MATCH (:Person{name:'Peter'})  
      -[:FRIEND]->(friends)  
RETURN friends
```



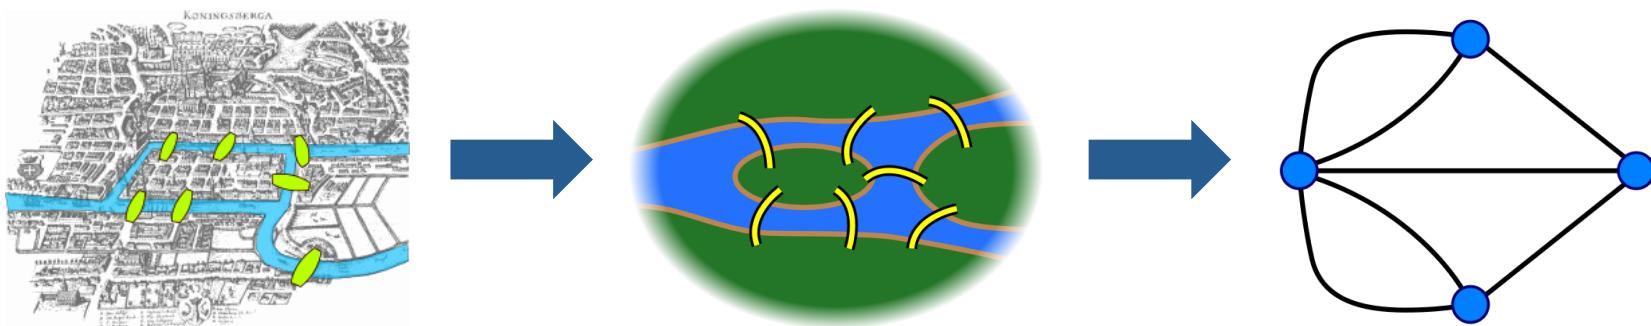
Exercise 1

Modeling Example



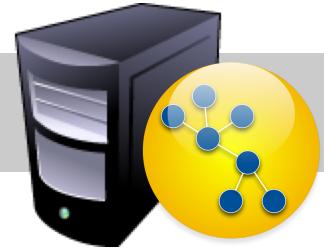
Models

Purposeful abstraction of a domain designed to satisfy particular application/end-user goals

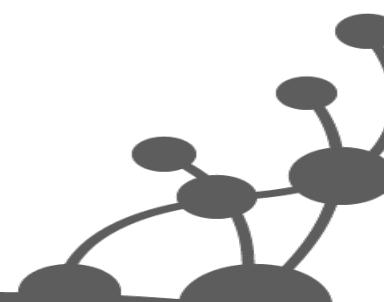


Images: en.wikipedia.org

Example Application



- Knowledge management
 - People, companies, skills
 - Cross organizational
- Find my professional social network
 - Exchange knowledge
 - Interest groups
 - Help
 - Staff projects

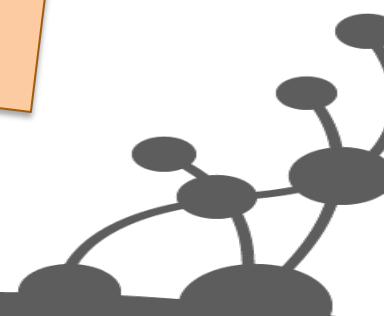


Application/End-User Goals

As an employee

**I want to know who in the company
has similar skills to me**

So that we can exchange knowledge



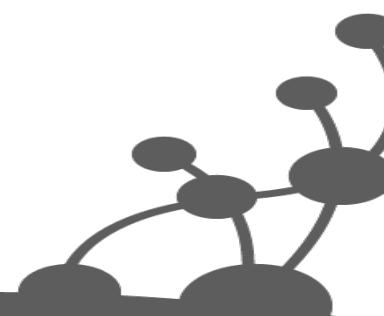
Questions To Ask of the Domain

As an employee

*I want to know who in the company
has similar skills to me*

So that we can exchange knowledge

Which people, who work for the same company
as me, have similar skills to me?



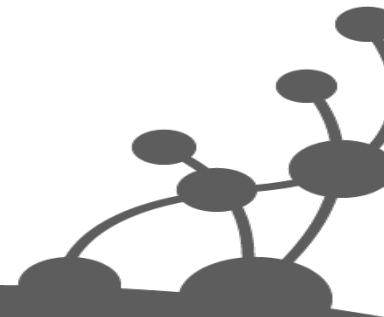
Identify Entities

Which **people**, who work for the same **company** as me, have similar **skills** to me?

Person

Company

Skill

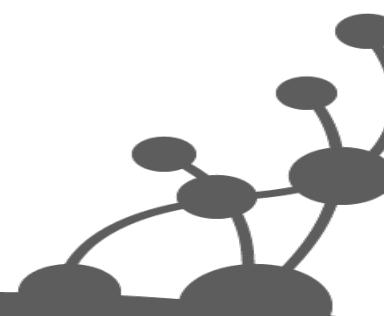


Identify Relationships Between Entities

Which people, who **work** for the same company
as me, **have** similar skills to me?

Person **WORKS_FOR** Company

Person **HAS_SKILL** Skill



Convert to Cypher Paths

Relationship

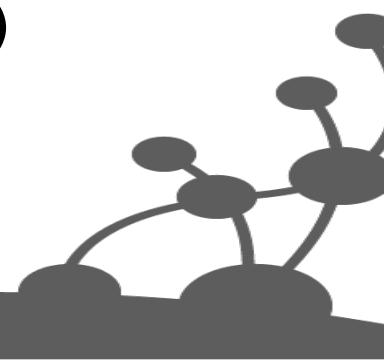
Person WORKS_FOR Company

Person HAS_SKILL Skill

Label



(:Person)-[:WORKS_FOR]->(:Company),
(:Person)-[:HAS_SKILL]->(:Skill)

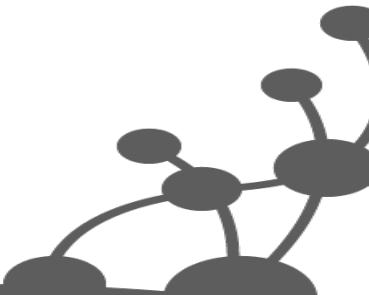
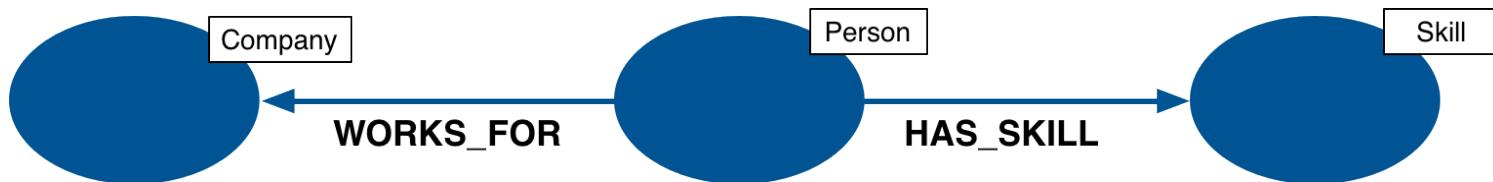


Consolidate Paths

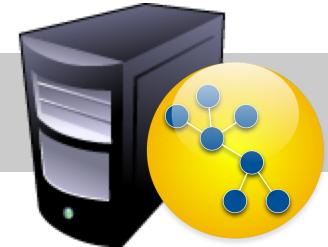
(:Person)-[:WORKS_FOR]->(:Company),
(:Person)-[:HAS_SKILL]->(:Skill)



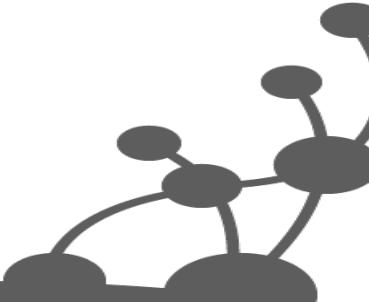
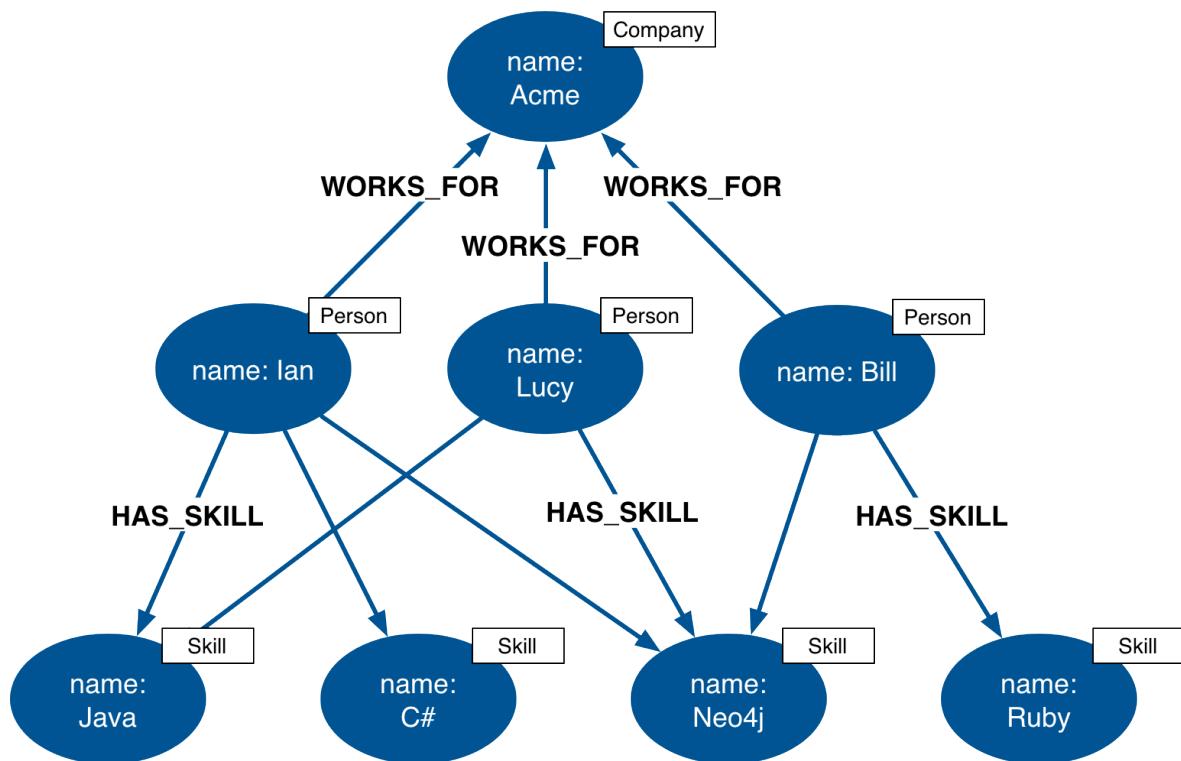
(:Company)<-[:WORKS_FOR]-(:Person)-[:HAS_SKILL]->(:Skill)



Candidate Data Model

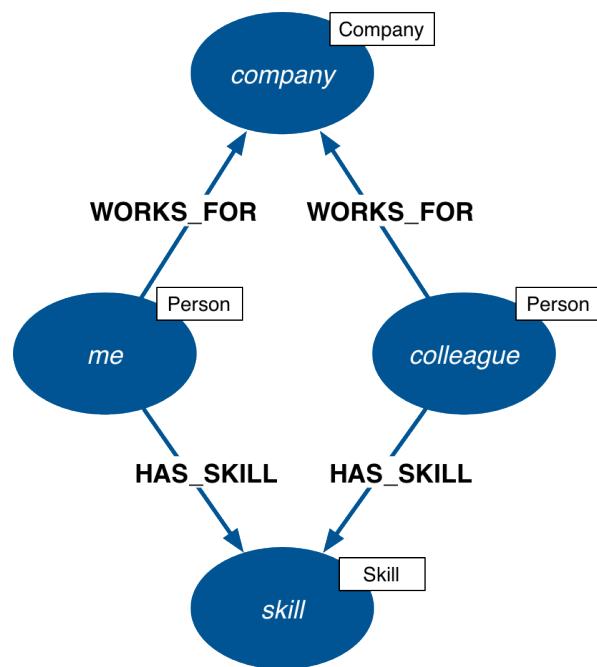


(:Company)<-[:WORKS_FOR]-(:Person)-[:HAS_SKILL]->(:Skill)

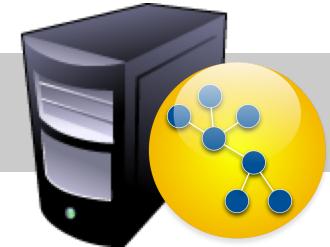


Express Question as Graph Pattern

Which people, who work for the same company as me, have similar skills to me?

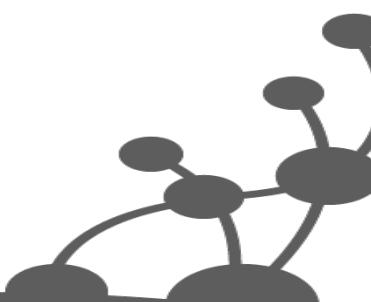
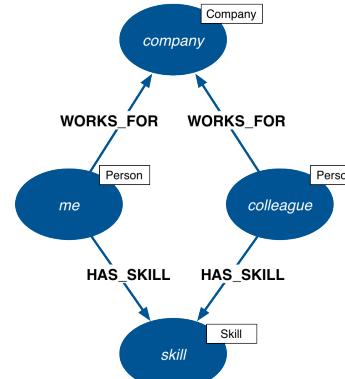


Cypher Query



Which people, who work for the same company as me, have similar skills to me?

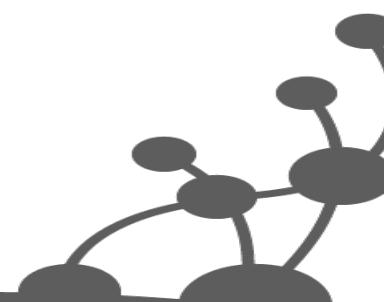
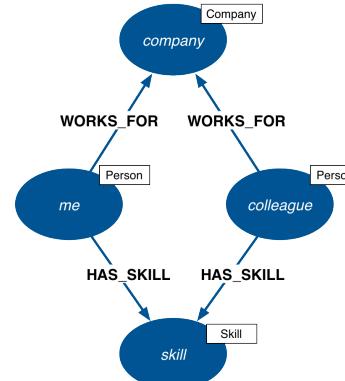
```
MATCH (company)<-[:WORKS_FOR]-(:Person{name:'Ian'})  
      -[:HAS_SKILL]->(skill),  
  (company)<-[:WORKS_FOR]-(colleague)-[:HAS_SKILL]->(skill)  
RETURN colleague.name AS name,  
       count(skill) AS score,  
       collect(skill.name) AS skills  
ORDER BY score DESC
```



Graph Pattern

Which people, who work for the same company as me, have similar skills to me?

```
MATCH (company)<-[:WORKS_FOR]-(:Person{name:'Ian'})  
      -[:HAS_SKILL]->(skill),  
  (company)<-[:WORKS_FOR]-(colleague)-[:HAS_SKILL]->(skill)  
RETURN colleague.name AS name,  
       count(skill) AS score,  
       collect(skill.name) AS skills  
ORDER BY score DESC
```

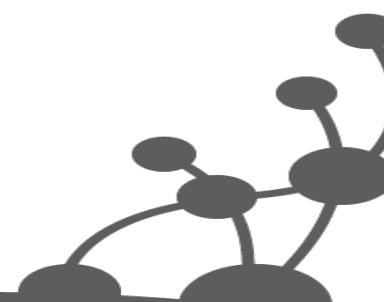
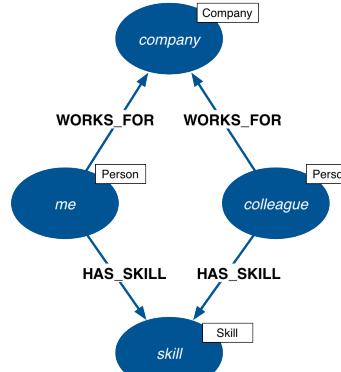


Anchor Pattern in Graph

Which people, who work for the same company as me, have similar skills to me?

```
MATCH (company)<-[:WORKS_FOR]-(:Person{name:'Ian'})  
      -[:HAS_SKILL]->(skill),  
  (company)<-[:WORKS_FOR]-(colleague)-[:HAS_SKILL]->(skill)  
RETURN colleague.name AS name,  
       count(skill) AS score,  
       collect(skill.name) AS skills  
ORDER BY score DESC
```

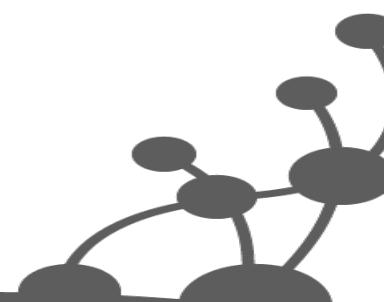
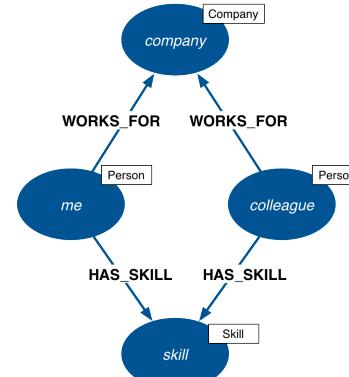
Search nodes labeled
'Person', matching on
'name' property



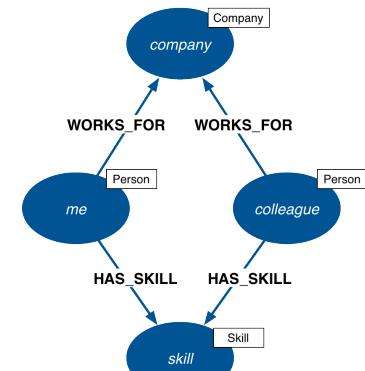
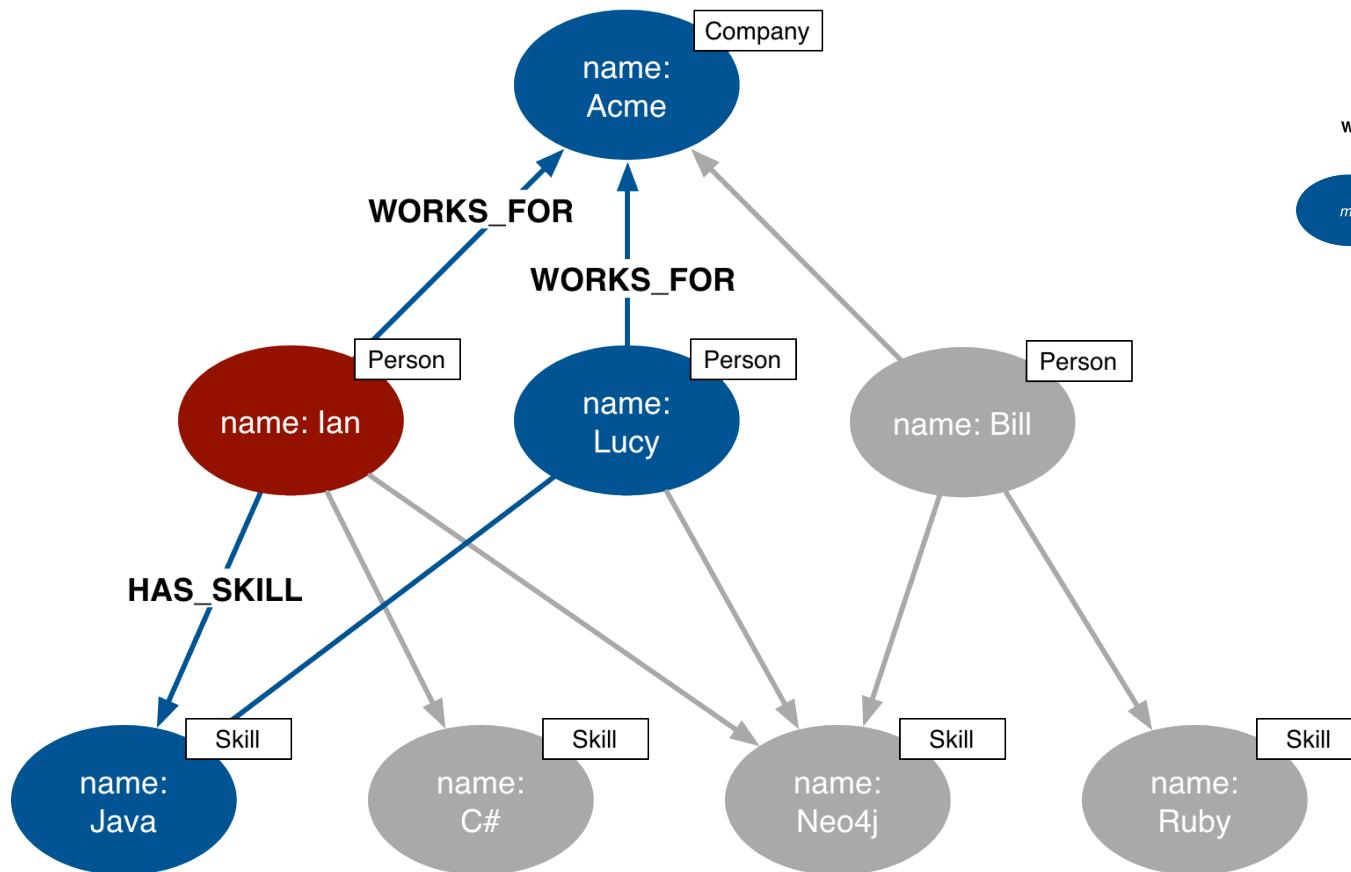
Create Projection of Results

Which people, who work for the same company as me, have similar skills to me?

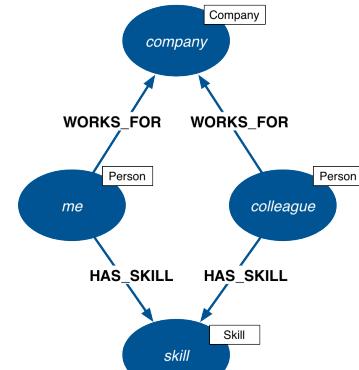
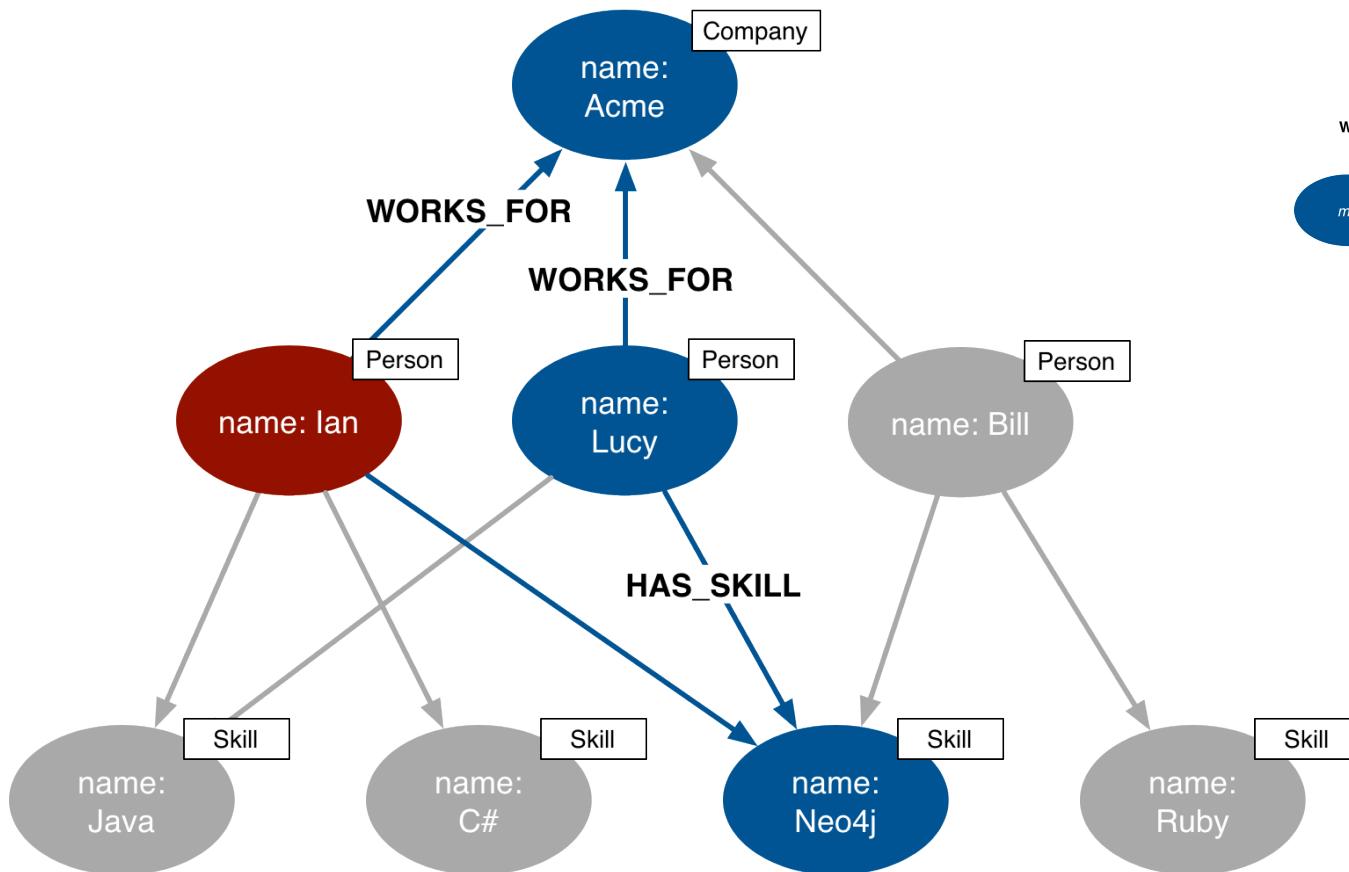
```
MATCH (company)<-[:WORKS_FOR]-(:Person{name:'Ian'})  
      -[:HAS_SKILL]->(skill),  
  (company)<-[:WORKS_FOR]-(colleague)-[:HAS_SKILL]->(skill)  
RETURN colleague.name AS name,  
       count(skill) AS score,  
       collect(skill.name) AS skills  
ORDER BY score DESC
```



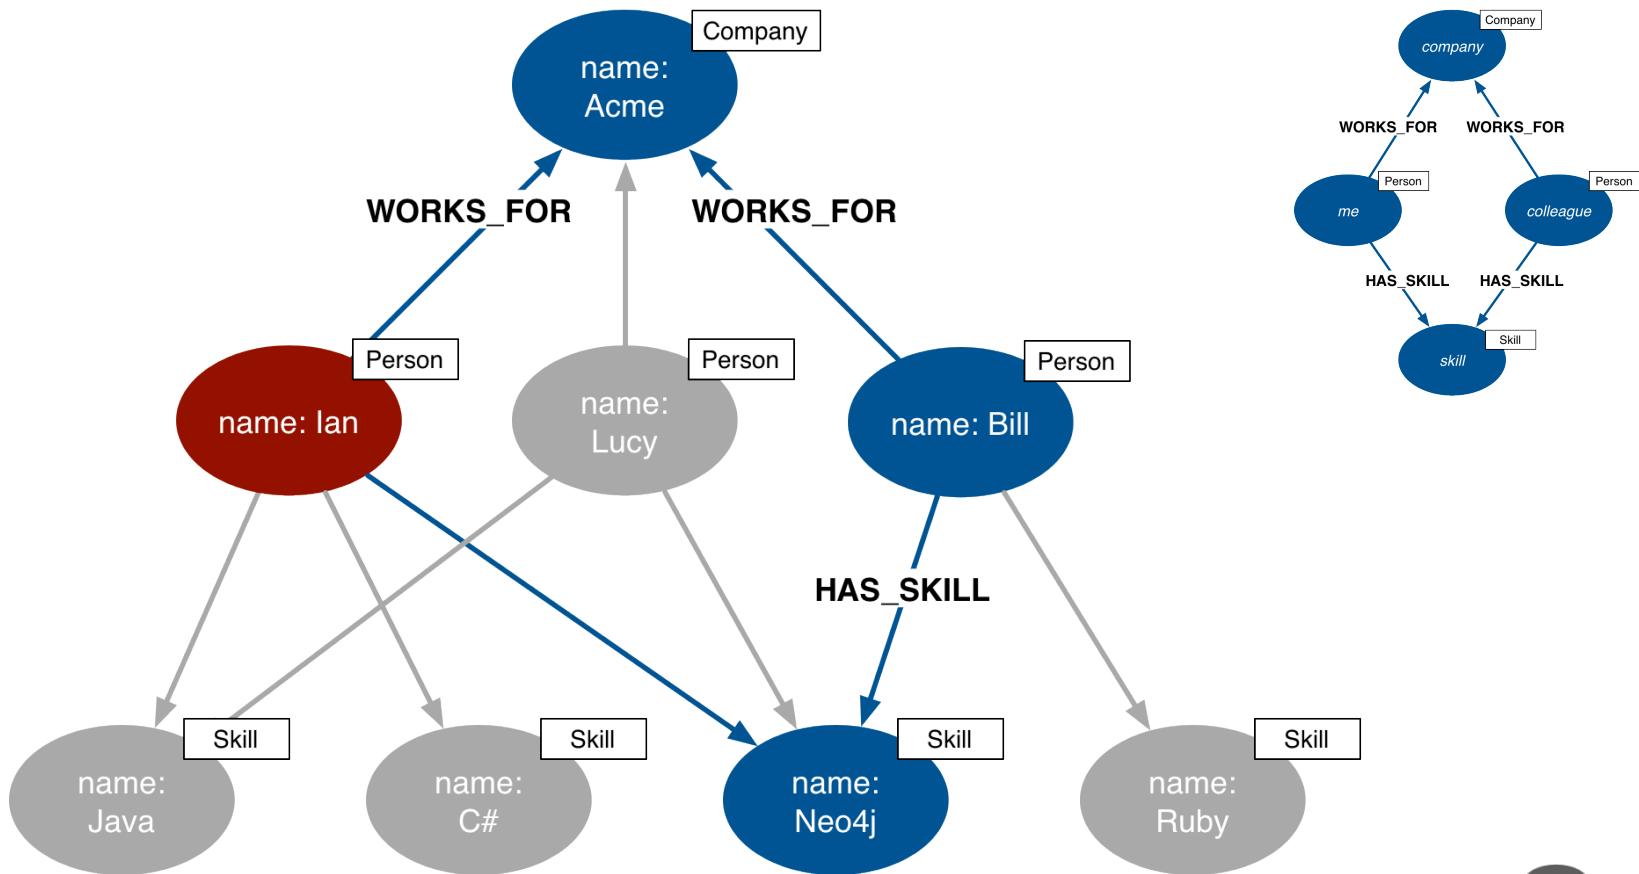
First Match



Second Match

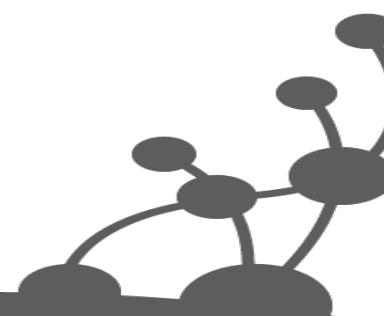


Third Match



Running the Query

```
+-----+  
| name | score | skills |  
+-----+  
| "Lucy" | 2      | ["Java", "Neo4j"] |  
| "Bill" | 1      | ["Neo4j"]       |  
+-----+  
2 rows
```



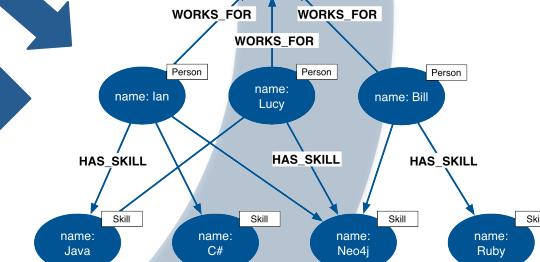
From User Story to Model and Query

*As an employee
I want to know who in the company
has similar skills to me
So that we can exchange knowledge*

Which people, who work for the same
company as me, have similar skills to me?

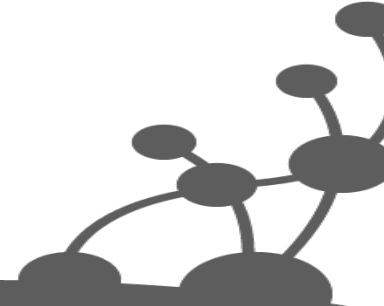
Person WORKS_FOR Company
Person HAS_SKILL Skill

```
MATCH (company)<-[ :WORKS_FOR ]-(me:Person)-[ :HAS_SKILL ]->(skill),  
(company)<-[ :WORKS_FOR ]-(colleague)-[ :HAS_SKILL ]->(skill)  
WHERE me.name = {name}  
RETURN colleague.name AS name,  
count(skill) AS score,  
collect(skill.name) AS skills  
ORDER BY score DESC
```

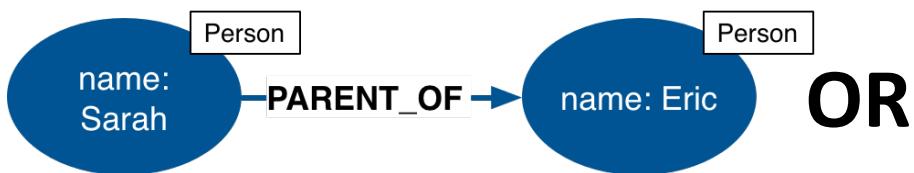
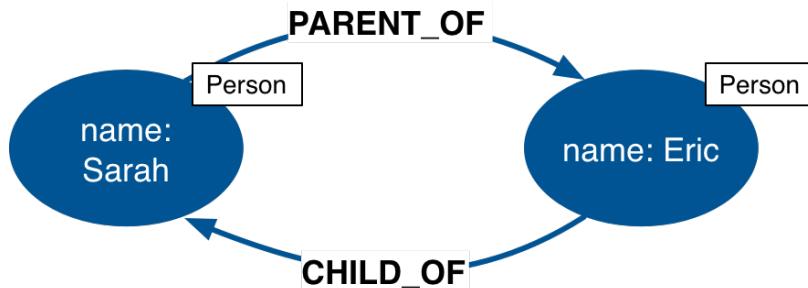
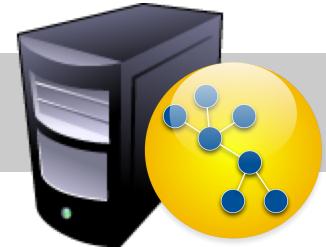


(:Company)<-[:WORKS_FOR]-(:Person)-[:HAS_SKILL]->(:Skill)

Modeling Guidelines



Symmetric Relationships



OR



Infer Symmetric Relationship

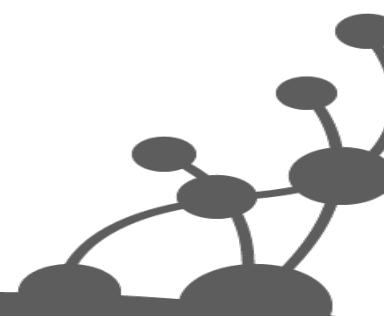


Find child:

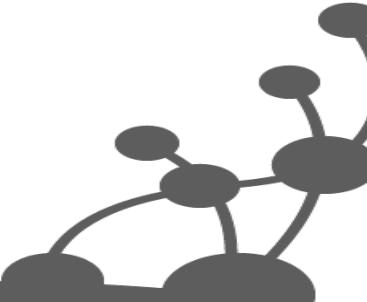
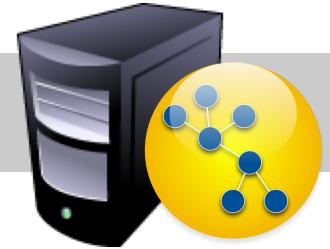
```
MATCH (parent{name:'Sarah'})  
      -[:PARENT_OF]->(child)  
RETURN child
```

Find parent:

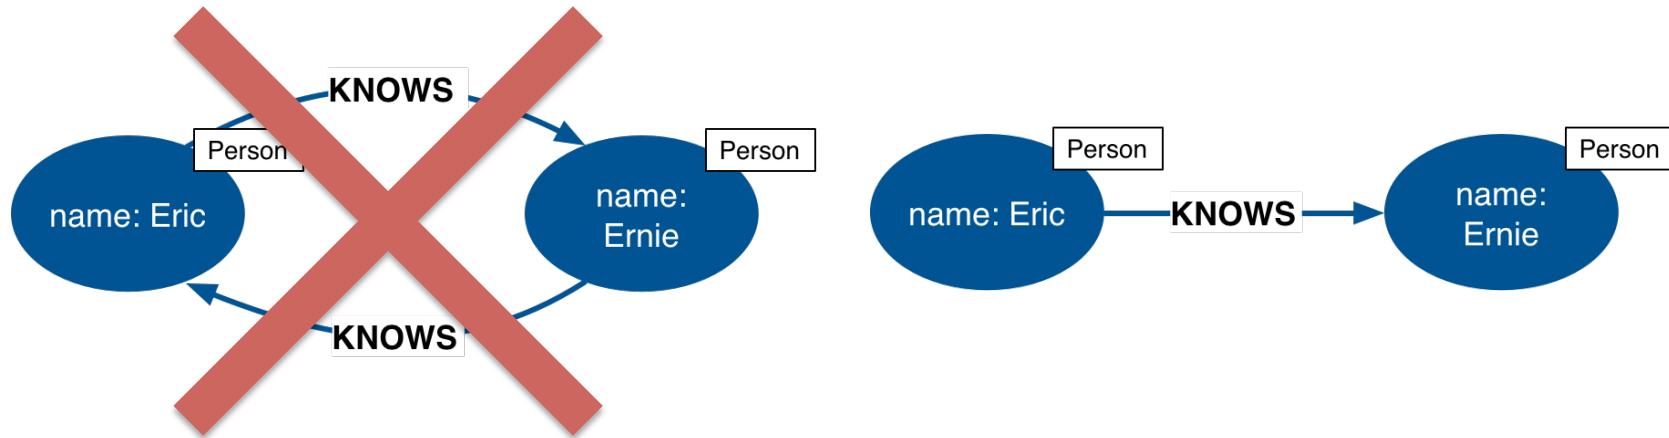
```
MATCH (parent)-[:PARENT_OF]->  
      (child{name:'Eric'})  
RETURN parent
```



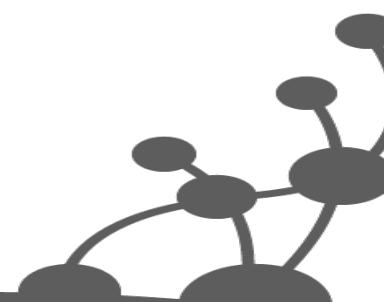
Bi-Directional Relationships



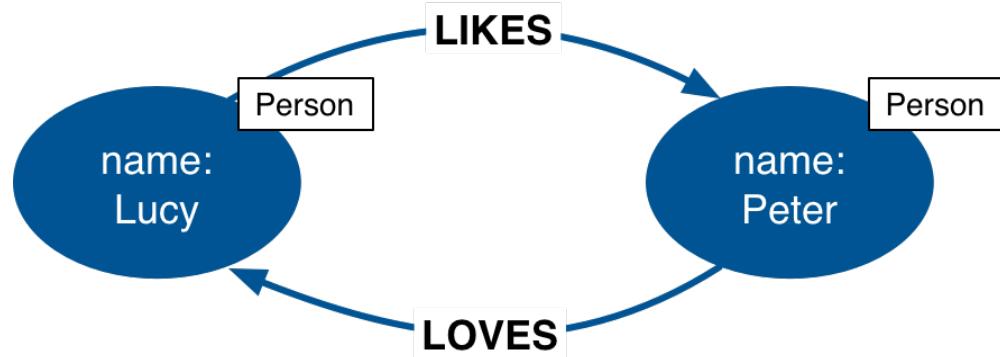
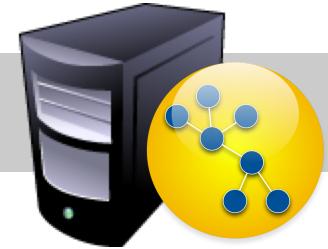
Use Single Relationship and Ignore Relationship Direction in Queries



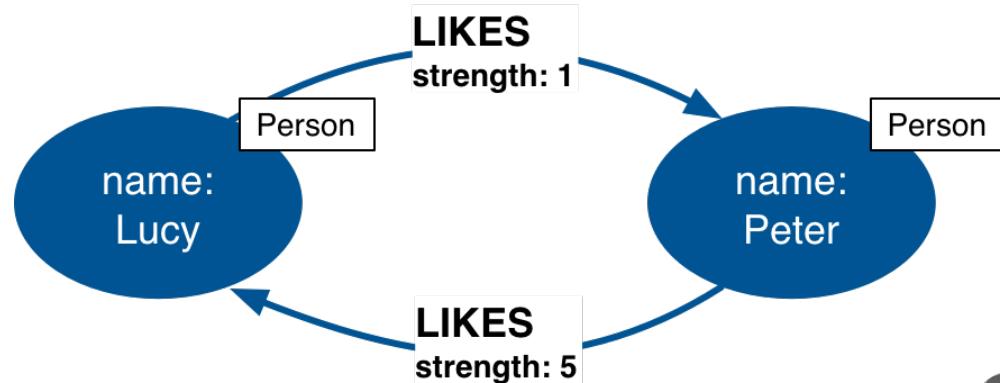
```
MATCH (p1{name:'Eric'})  
  -[:KNOWS]-(p2)  
RETURN p2
```



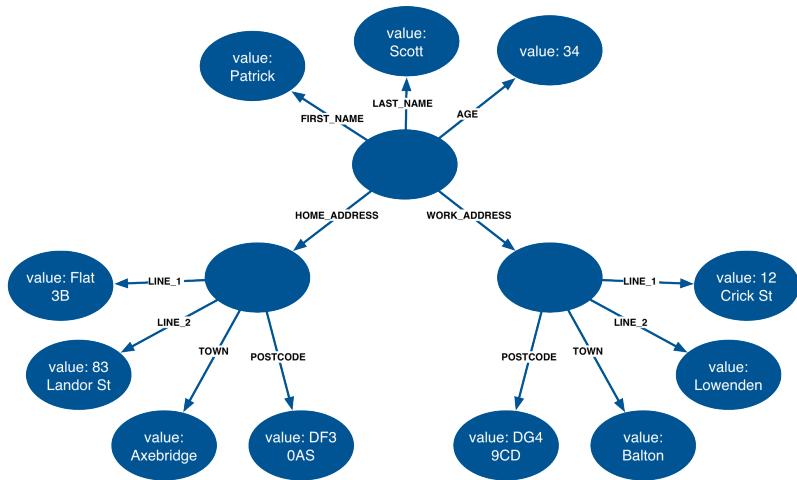
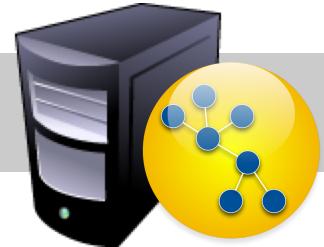
Qualified Bi-Directional Relationships



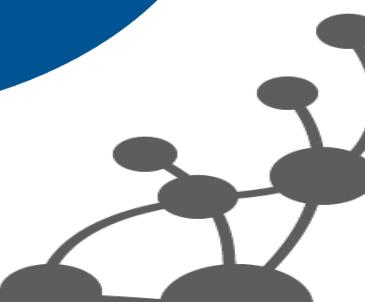
OR



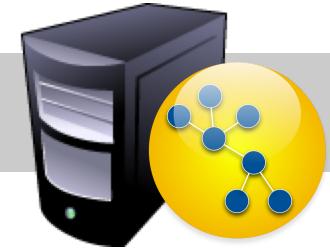
Properties Versus Relationships



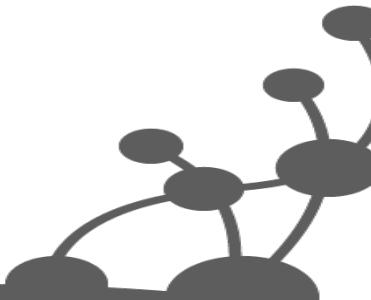
first-name: Patrick
last-name: Scott
age: 34
home-address: Flat 3B,
83 Landor St,
Axebridge,
DF3 0AS
work-address: Acme Ltd,
12 Crick St,
Balton,
DG4 9CD



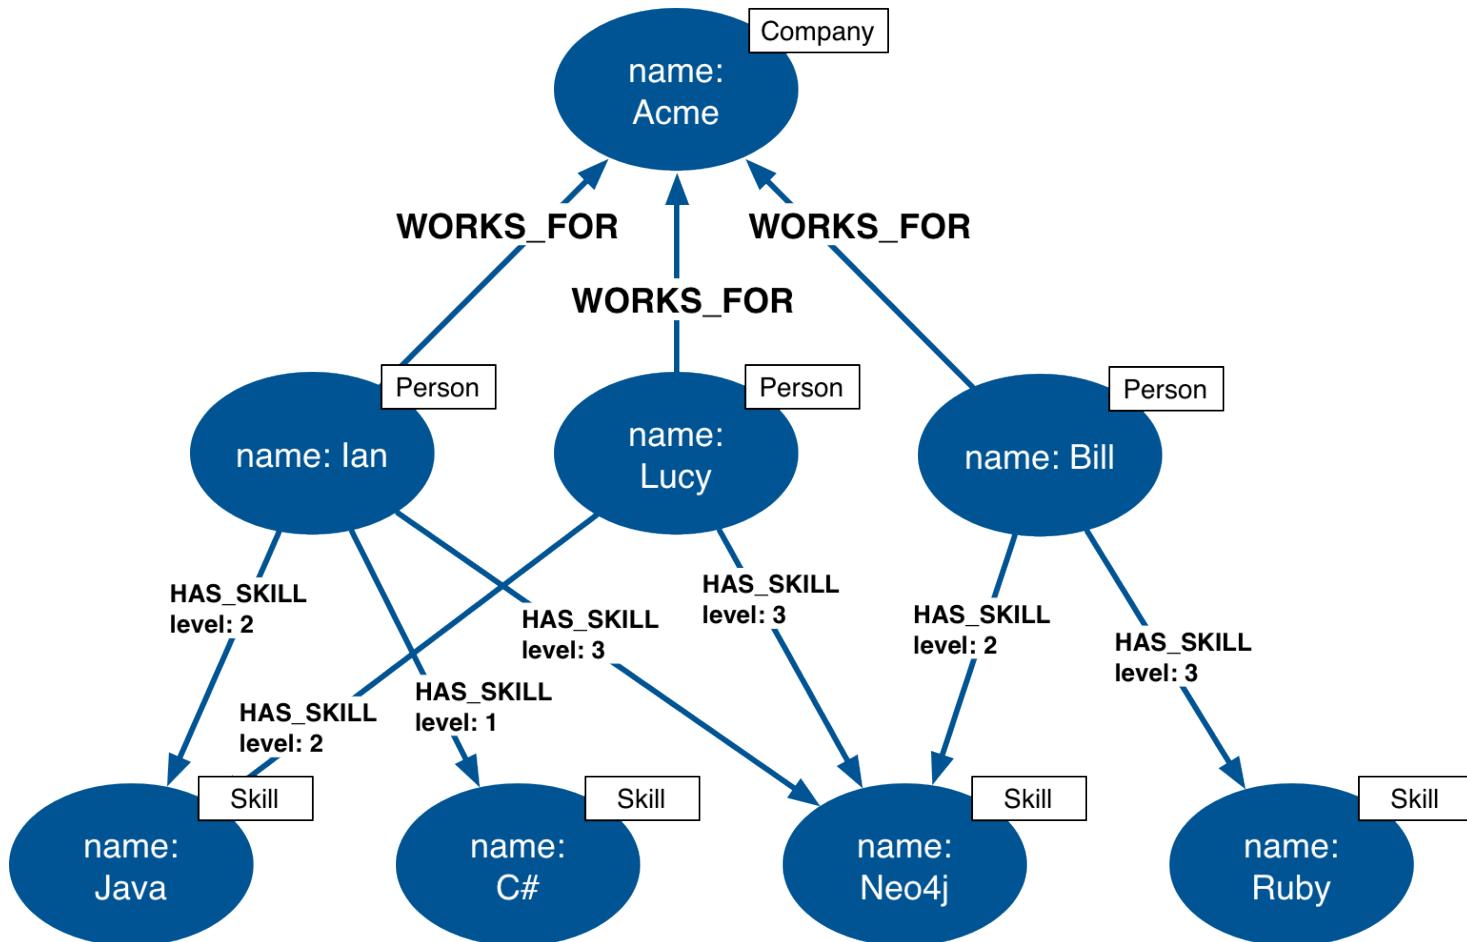
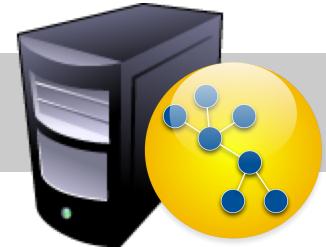
Use Relationships When...



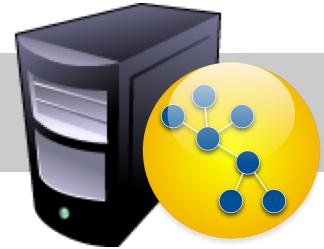
- You need to specify the weight, strength, or some other quality of the *relationship*
- AND/OR the attribute value comprises a complex value type (e.g. address)
- Examples:
 - Find all my colleagues who are *level 2 or above* (relationship quality) in a *skill* (attribute value) we have in common
 - Find all recent orders delivered to the same *delivery address* (complex value type)



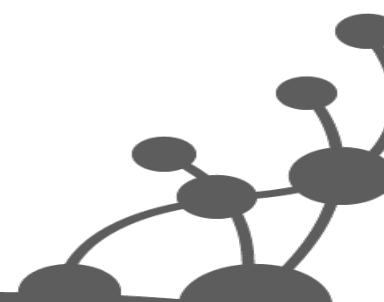
Example: Find Expert Colleagues



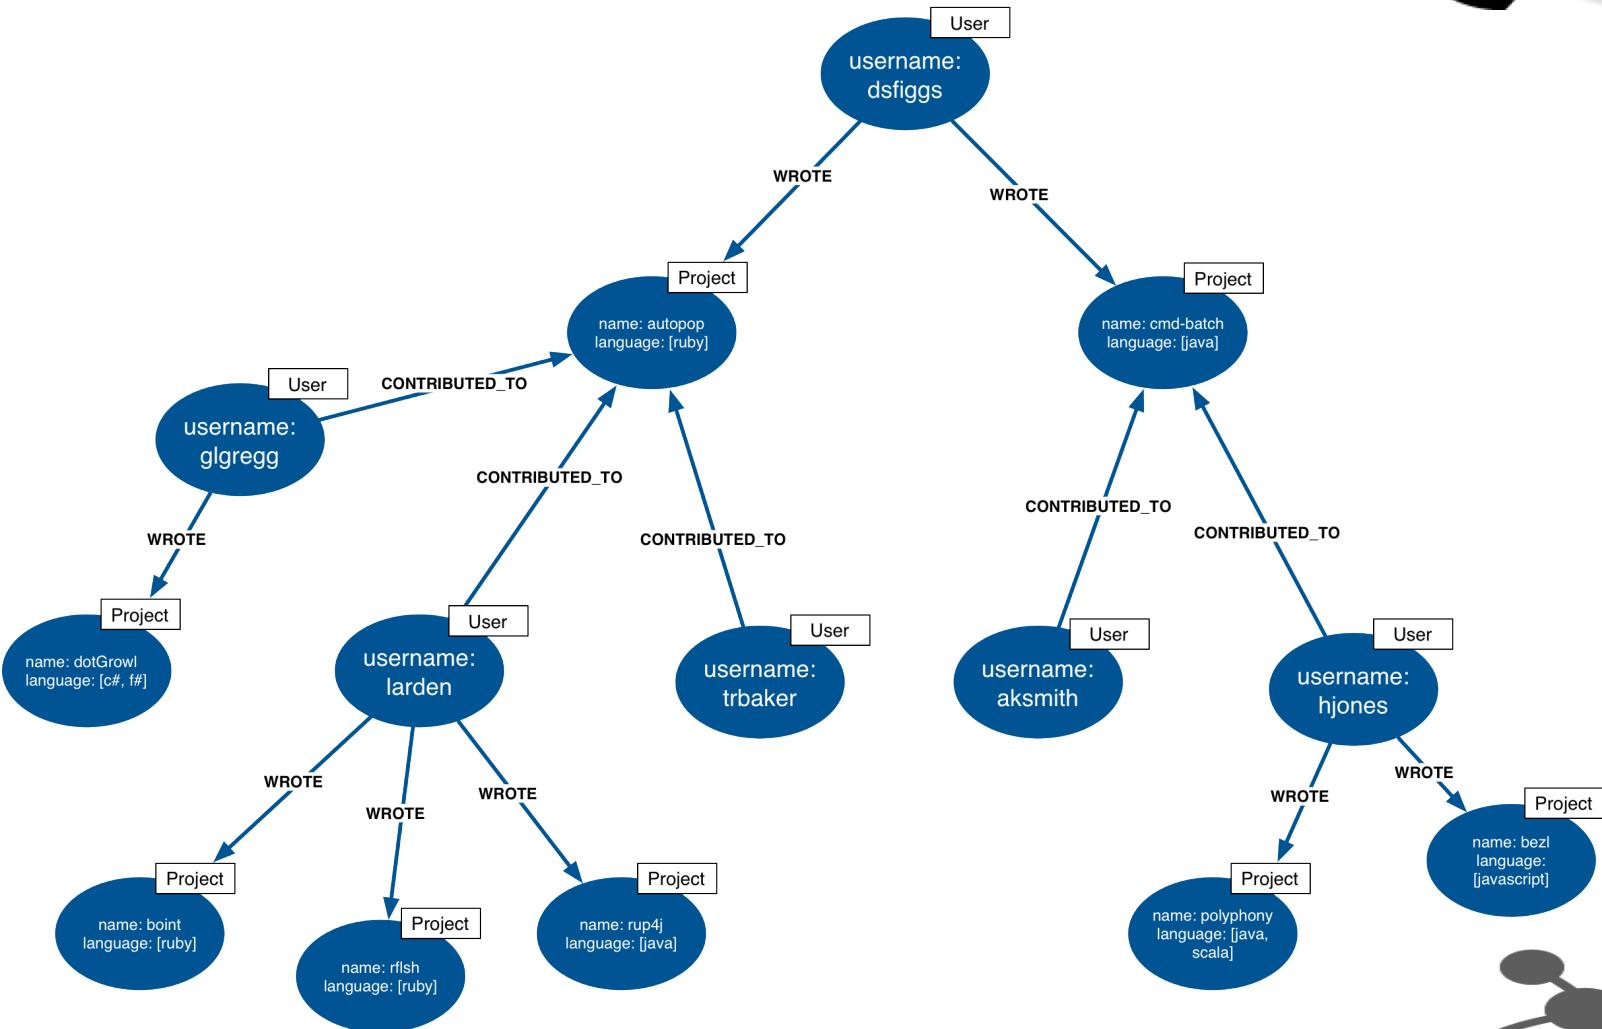
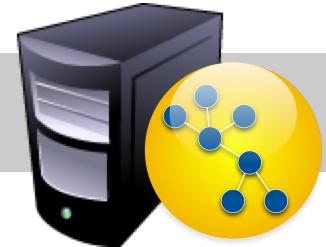
Use Properties When...



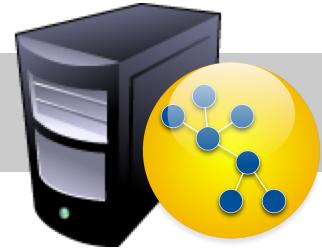
- There's no need to qualify the relationship
- AND the attribute value comprises a simple value type (e.g. colour)
- Examples:
 - Find those projects written by contributors to my projects that use the same *language* (attribute value) as my projects



Example: Similar By Language



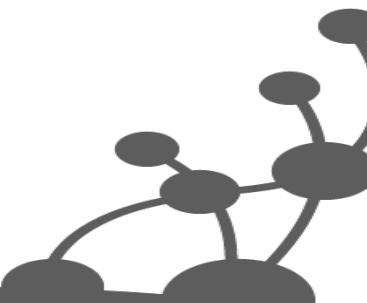
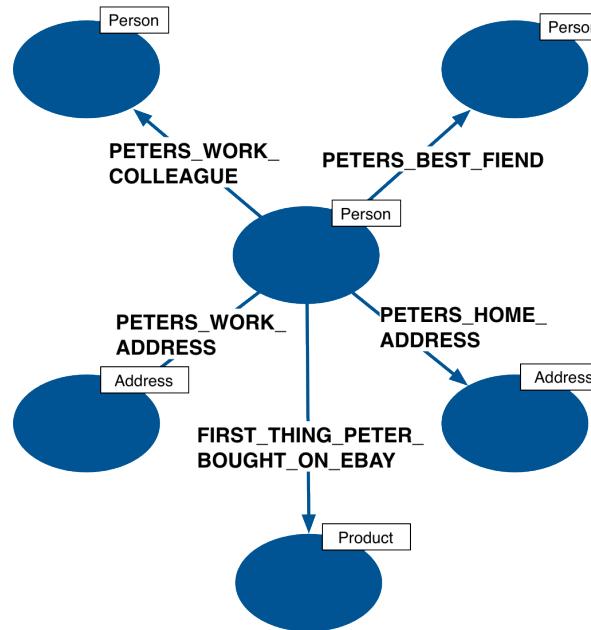
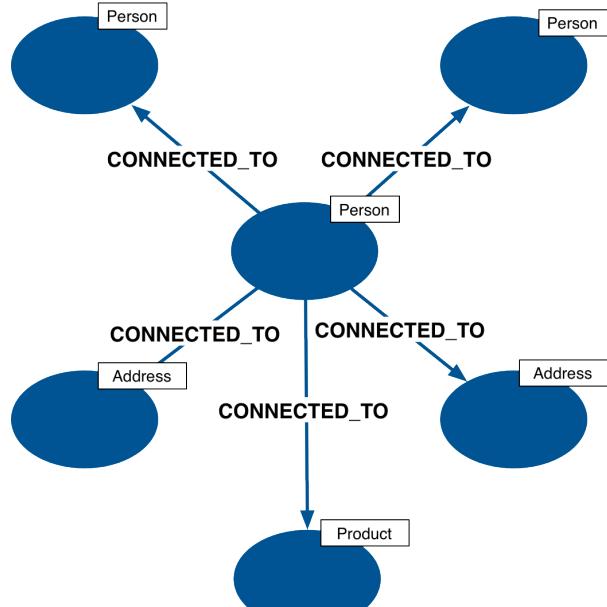
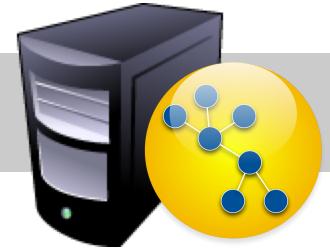
If Performance is Critical...



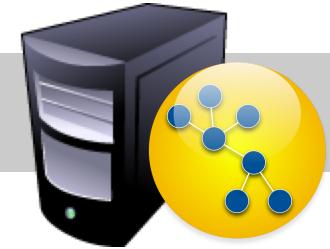
- Small property lookup on a node will be quicker than traversing a relationship
 - But traversing a relationship is still faster than a SQL join...
- However, *many small properties* on a node, or a lookup on a *large string* or *large array* property will impact performance
 - Always performance test against a representative dataset



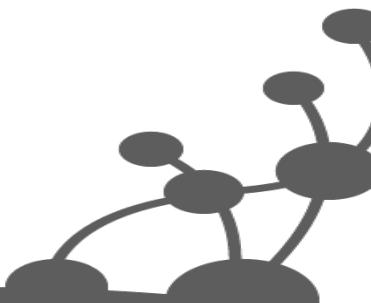
Relationship Granularity



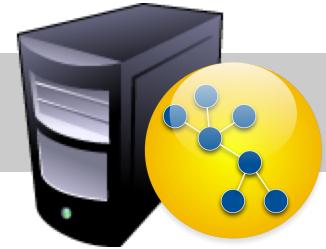
Align With Use Cases



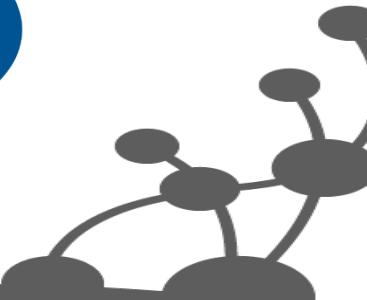
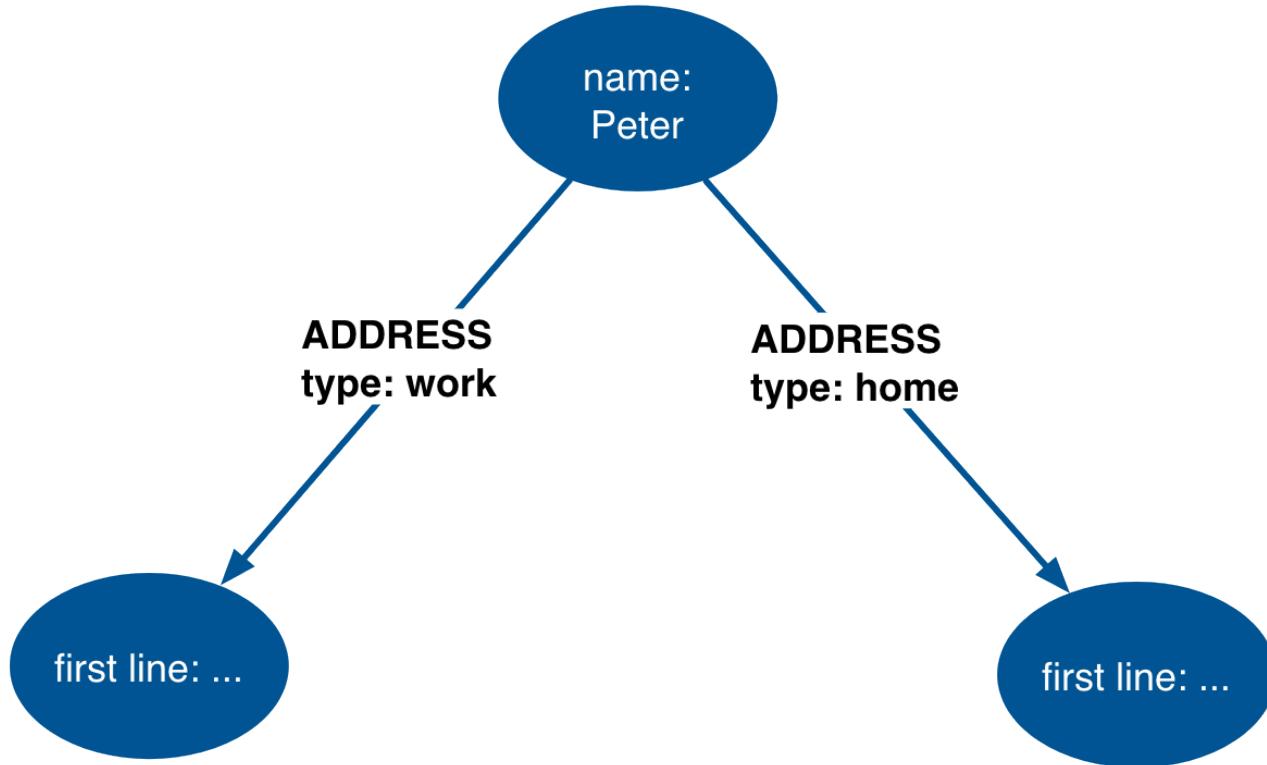
- Relationships are the “royal road” into the graph
- When querying, well-named relationships help discover only what is *absolutely necessary*
 - And eliminate unnecessary portions of the graph from consideration



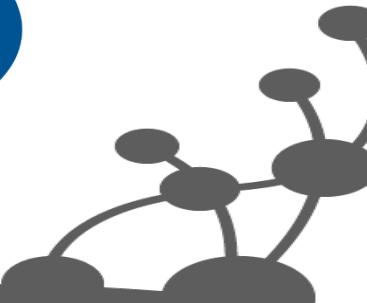
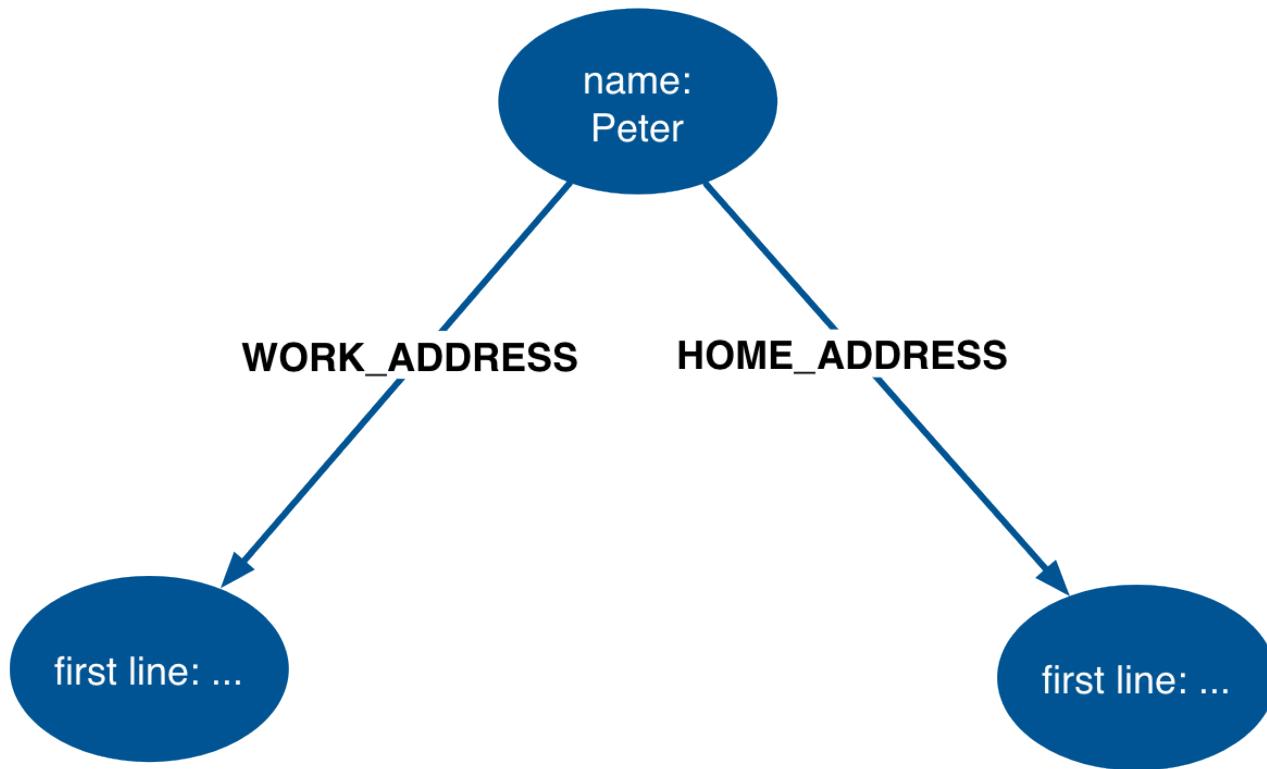
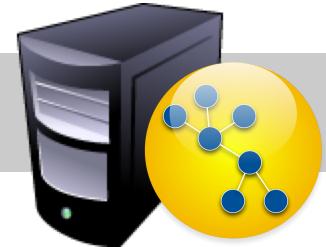
General Relationships



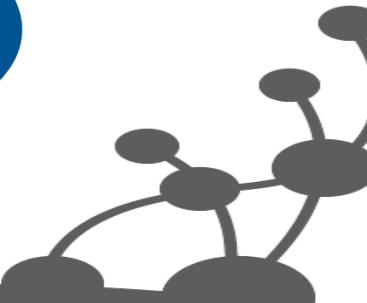
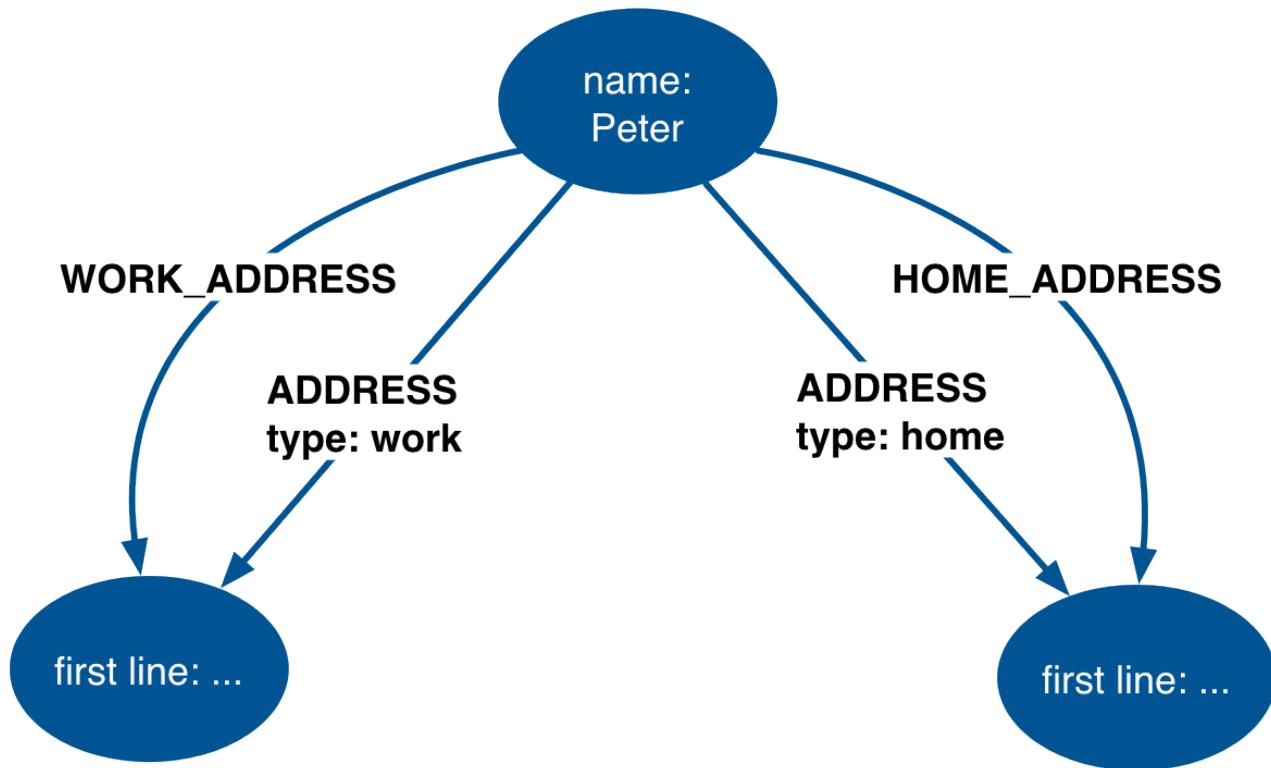
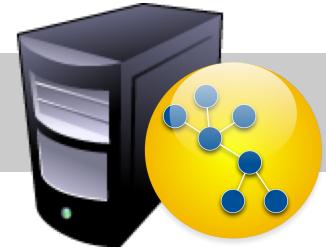
- Qualified by property



Specific Relationships

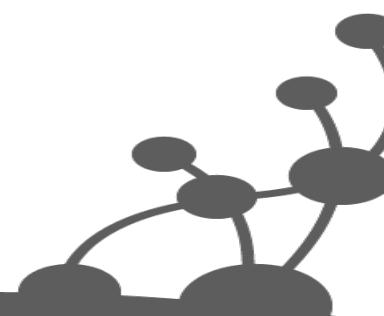


Best of Both Worlds



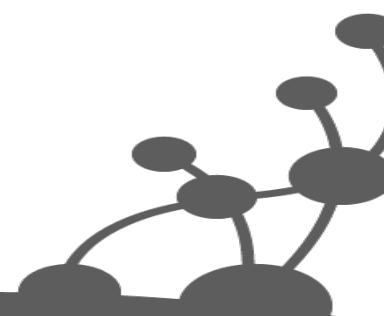
Exercise 2

Creating Data



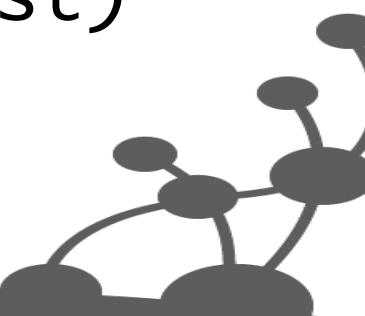
Exercise 2 - Create Some data

- Clean the database
- Execute *create-1.txt*
- View the results
 - MATCH (n) RETURN n



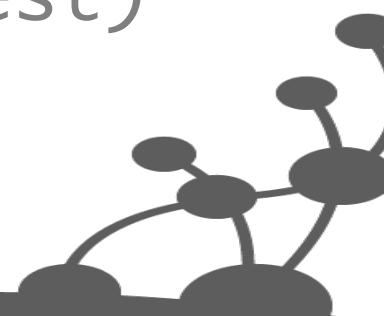
create-1.txt

```
CREATE
(ben:Person{username: 'ben'}),
(acme:Company{name: 'Acme, Inc'}),
(neo4j:Skill{name: 'Neo4j'}),
(rest:Skill{name: 'REST'}),
(ben)-[:WORKS_FOR]->(acme),
(ben)-[:HAS_SKILL{level:1}]->(neo4j),
(ben)-[:HAS_SKILL{level:3}]->(rest)
```



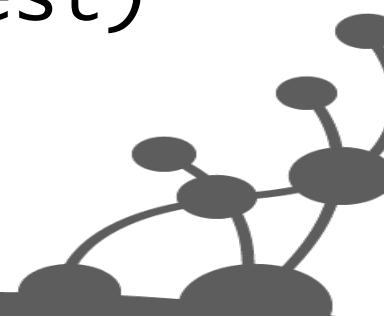
Create Nodes

```
CREATE  
(ben:Person{username: 'ben'}),  
(acme:Company{name: 'Acme, Inc'}),  
(neo4j:Skill{name: 'Neo4j'}),  
(rest:Skill{name: 'REST'}),  
(ben)-[:WORKS_FOR]->(acme),  
(ben)-[:HAS_SKILL{level:1}]->(neo4j),  
(ben)-[:HAS_SKILL{level:3}]->(rest)
```



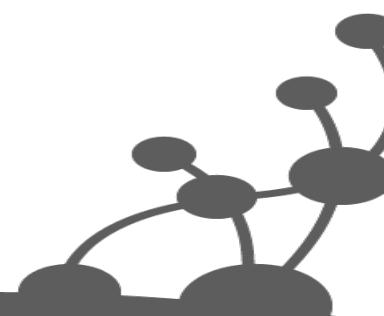
Connect Nodes

```
CREATE  
(ben:Person{username: 'ben'}),  
(acme:Company{name: 'Acme, Inc'}),  
(neo4j:Skill{name: 'Neo4j'}),  
(rest:Skill{name: 'REST'}),  
(ben)-[:WORKS_FOR]->(acme),  
(ben)-[:HAS_SKILL{level:1}]->(neo4j),  
(ben)-[:HAS_SKILL{level:3}]->(rest)
```



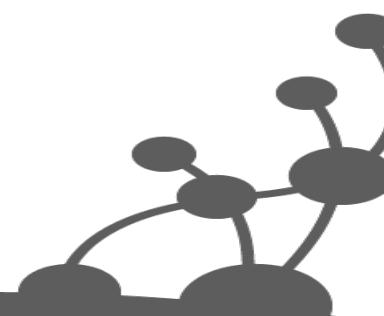
Create Some More Data

- Create more people
 - Same skills (Neo4j and REST)
 - Same company (Acme)
- View the results
 - `MATCH (n) RETURN n`



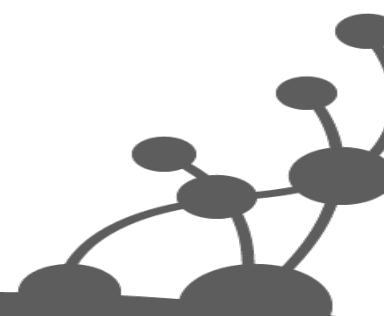
Your Turn

- Clean the database
- Execute *create-2.txt*, *create-3.txt* and *create-4.txt*
 - After each operation, view the results
- What happens if you add or remove properties when specifying unique nodes and relationships?



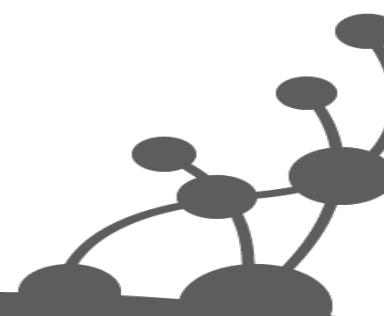
Creating Unique Nodes and Relationships

```
MERGE (c:Company{name:'Acme'})  
MERGE (p:Person{username:'ian'})  
MERGE (s1:Skill{name:'Java'})  
MERGE (s2:Skill{name:'C#'})  
MERGE (s3:Skill{name:'Neo4j'})  
MERGE (c)<-[:WORKS_FOR]-(p)  
MERGE (p)-[r1:HAS_SKILL]->(s1)  
MERGE (p)-[r2:HAS_SKILL]->(s2)  
MERGE (p)-[r3:HAS_SKILL]->(s3)  
SET r1.level = 2  
SET r2.level = 2  
SET r3.level = 3  
RETURN c, p, s1, s2, s3
```



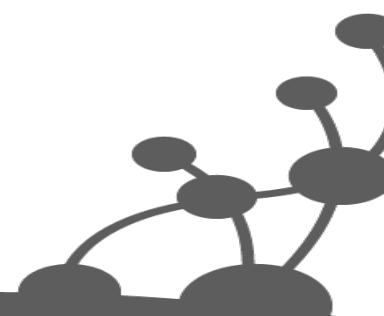
Create Unique Nodes

```
MERGE (c:Company{name:'Acme'})  
MERGE (p:Person{username:'ian'})  
MERGE (s1:Skill{name:'Java'})  
MERGE (s2:Skill{name:'C#'})  
MERGE (s3:Skill{name:'Neo4j'})  
MERGE (c)<-[:WORKS_FOR]-(p)  
MERGE (p)-[r1:HAS_SKILL]->(s1)  
MERGE (p)-[r2:HAS_SKILL]->(s2)  
MERGE (p)-[r3:HAS_SKILL]->(s3)  
SET r1.level = 2  
SET r2.level = 2  
SET r3.level = 3  
RETURN c, p, s1, s2, s3
```



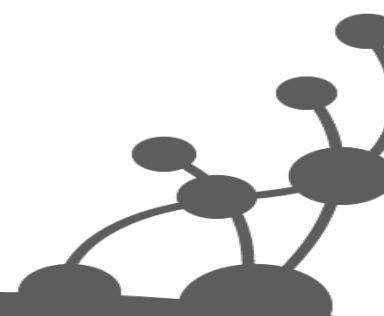
Create Unique Relationships

```
MERGE (c:Company{name:'Acme'})  
MERGE (p:Person{username:'ian'})  
MERGE (s1:Skill{name:'Java'})  
MERGE (s2:Skill{name:'C#'})  
MERGE (s3:Skill{name:'Neo4j'})  
MERGE (c)<-[:WORKS_FOR]-(p)  
MERGE (p)-[r1:HAS_SKILL]->(s1)  
MERGE (p)-[r2:HAS_SKILL]->(s2)  
MERGE (p)-[r3:HAS_SKILL]->(s3)  
SET r1.level = 2  
SET r2.level = 2  
SET r3.level = 3  
RETURN c, p, s1, s2, s3
```



Set Relationship Properties

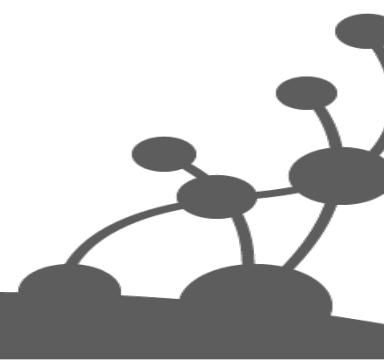
```
MERGE (c:Company{name:'Acme'})  
MERGE (p:Person{username:'ian'})  
MERGE (s1:Skill{name:'Java'})  
MERGE (s2:Skill{name:'C#'})  
MERGE (s3:Skill{name:'Neo4j'})  
MERGE (c)<-[:WORKS_FOR]-(p)  
MERGE (p)-[r1:HAS_SKILL]->(s1)  
MERGE (p)-[r2:HAS_SKILL]->(s2)  
MERGE (p)-[r3:HAS_SKILL]->(s3)  
SET r1.level = 2  
SET r2.level = 2  
SET r3.level = 3  
RETURN c, p, s1, s2, s3
```



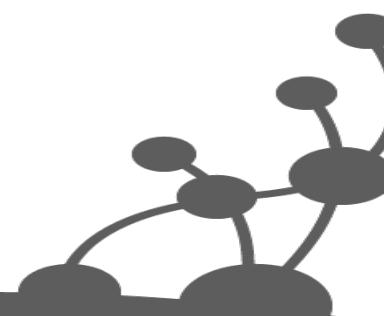
MERGE

MERGE ensures that a pattern exists in the graph.
Either the pattern already exists, or it needs to
be created.

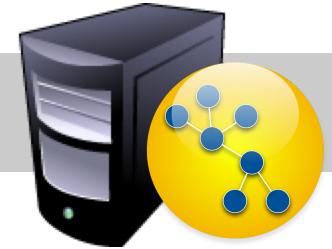
<http://docs.neo4j.org/chunked/milestone/query-merge.html>



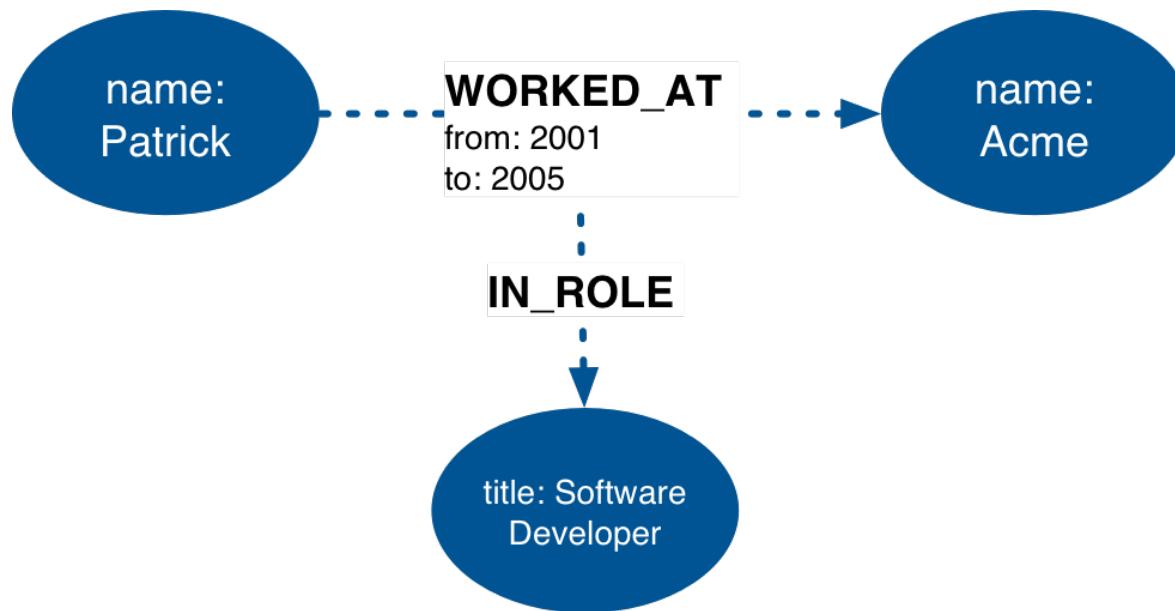
Common Graph Structures



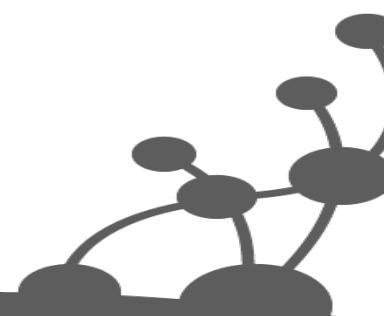
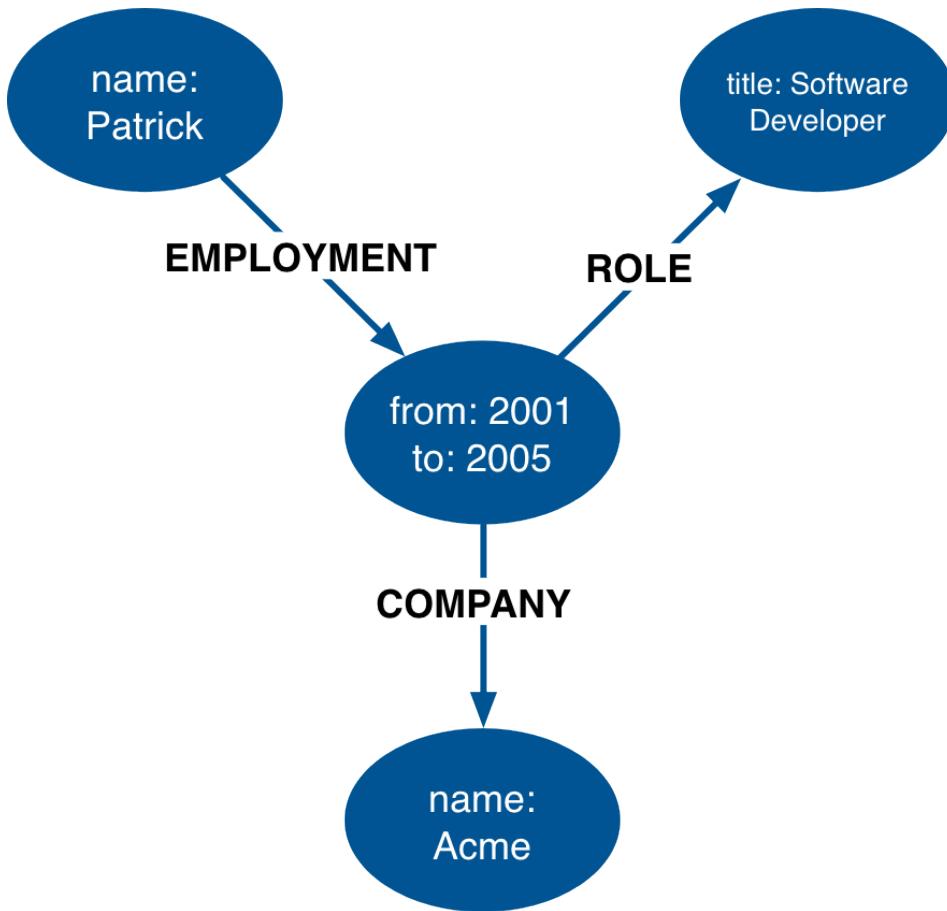
Intermediate Nodes



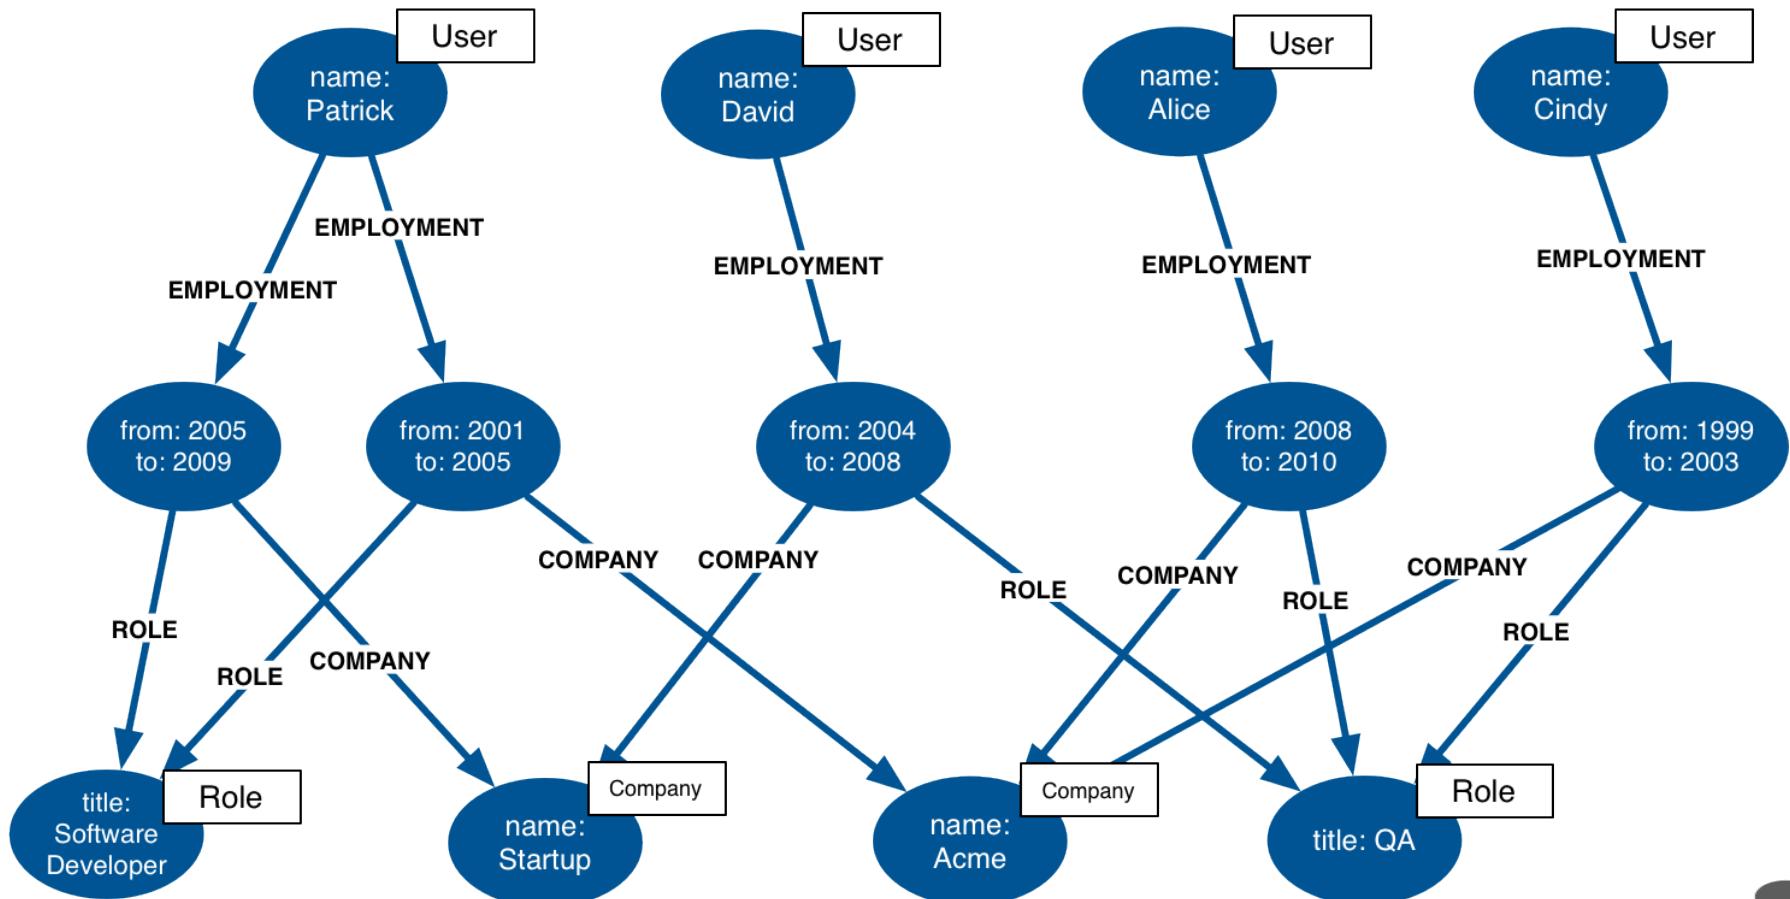
- Connect more than 2 nodes in a single context
 - Hyperedges (n-ary relationships)
- Relate something to a relationship



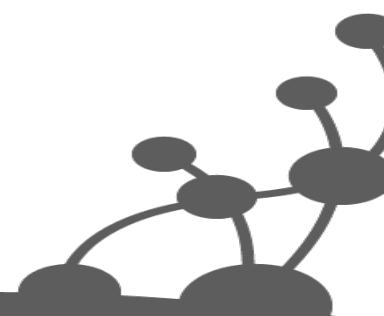
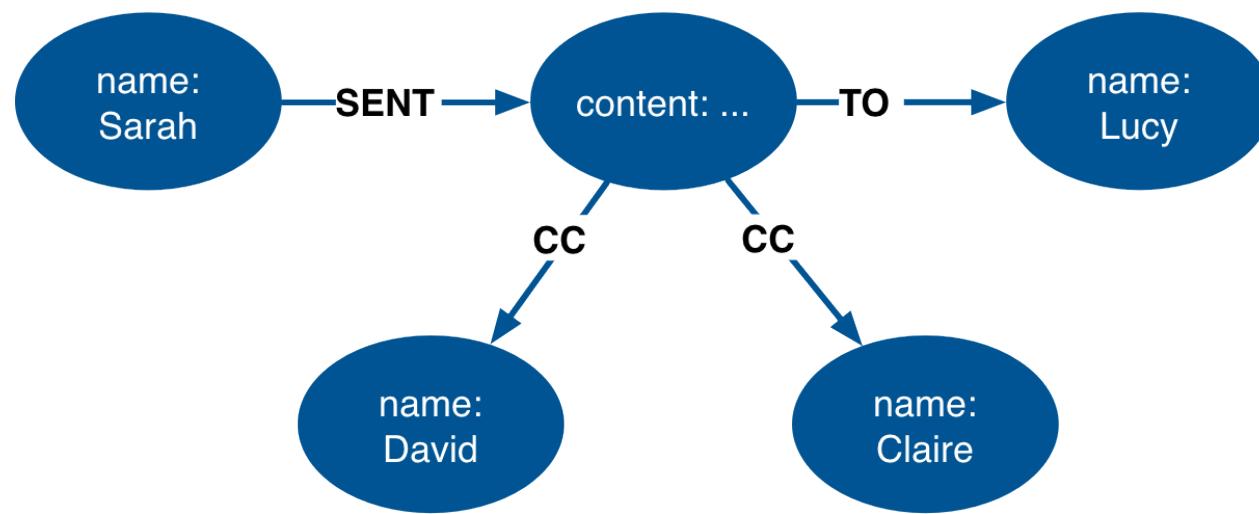
Rich Context, Multiple Dimensions



Dimensions Shared Between Contexts

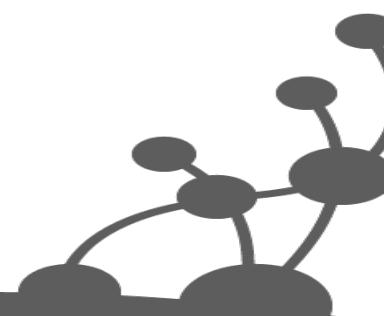


Multiple Parties

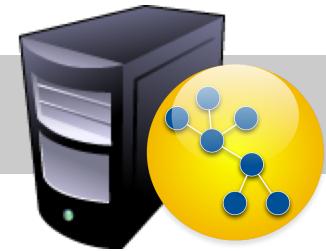


Considerations

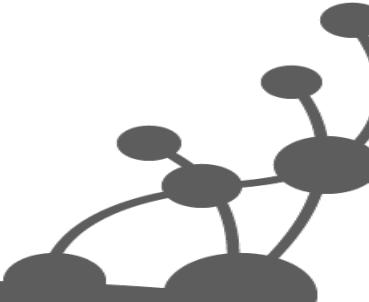
- An intermediate node provides flexibility
 - It allows more than two nodes to be connected in a single context
- But it can be overkill, and will have an impact on performance



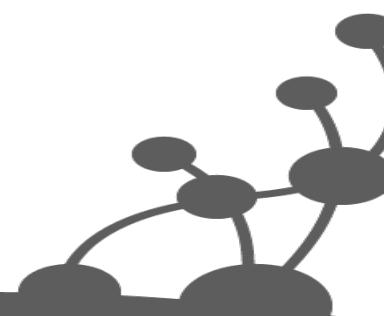
Linked Lists



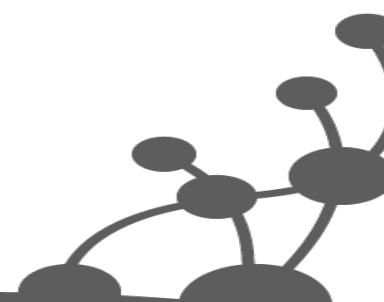
- Entities are linked in a sequence
- You need to traverse the sequence
- You may need to identify the beginning or end (first/last, earliest/latest, etc.)
- Examples
 - Event stream
 - Episodes of a TV series
 - Job history



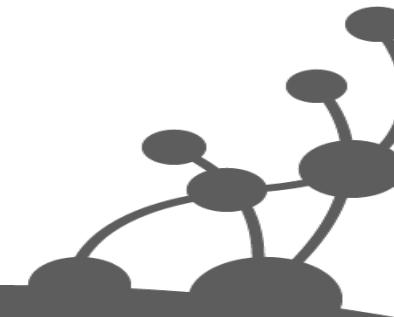
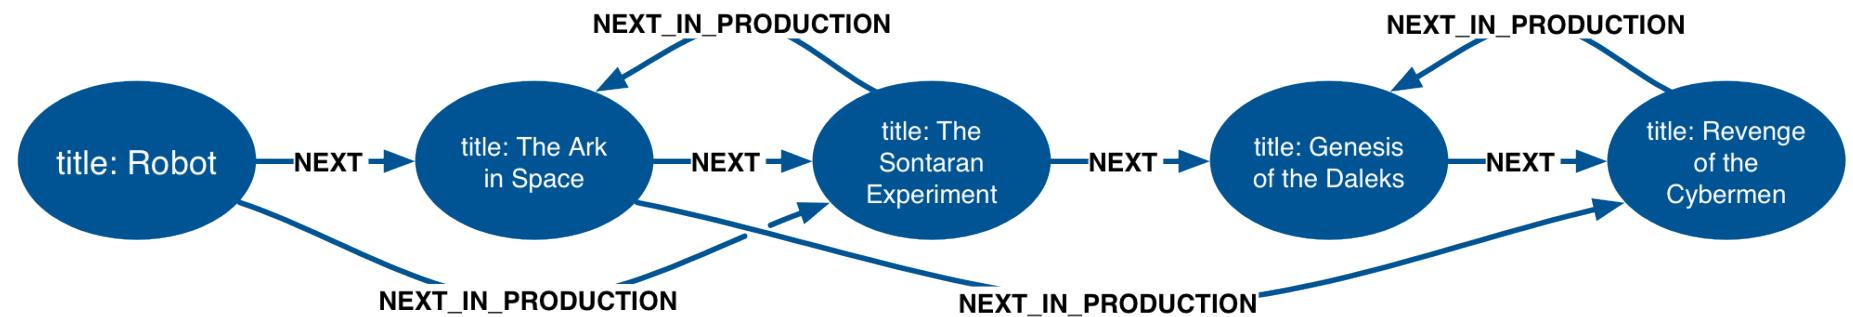
Linked List



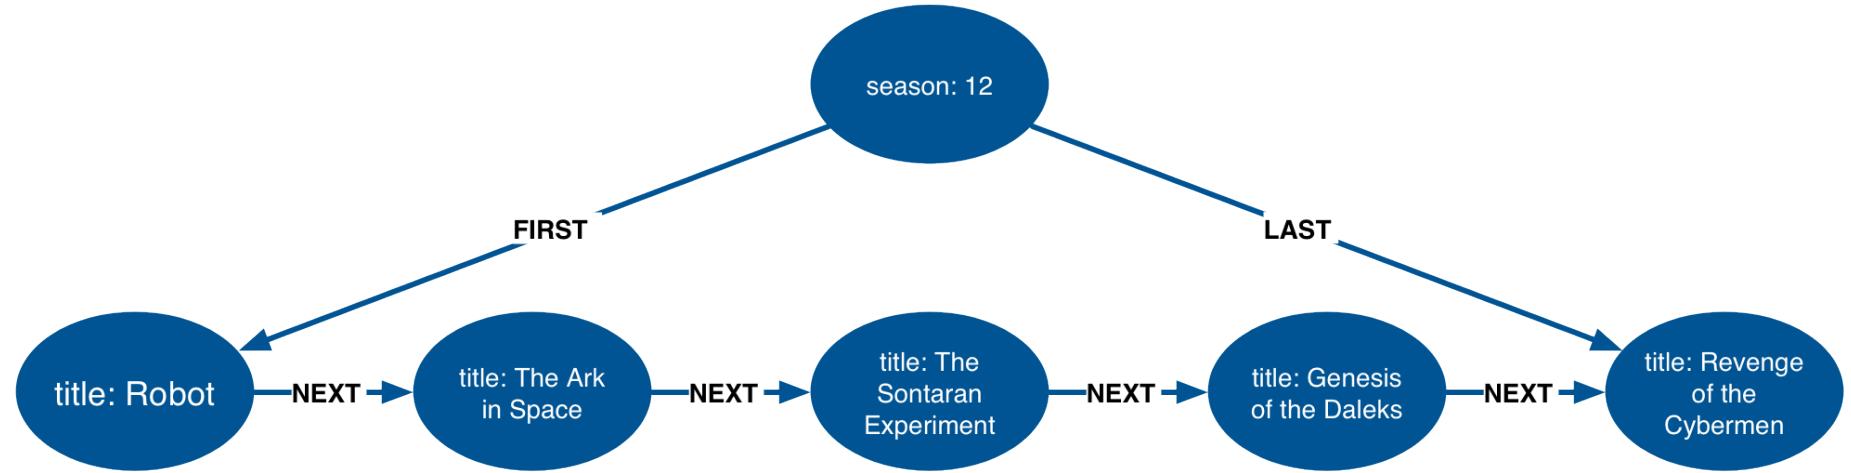
Doubly Linked List



Interleaved Linked Lists

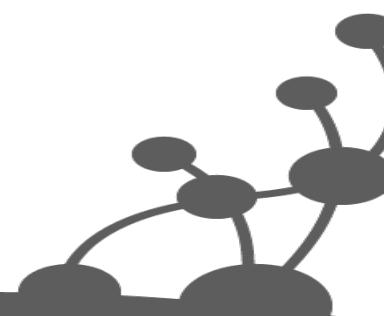


Pointers to Head and Tail



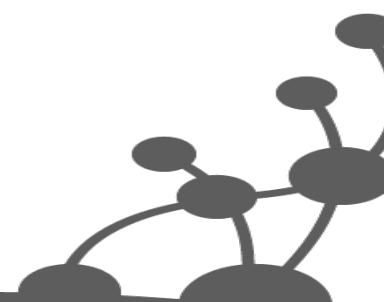
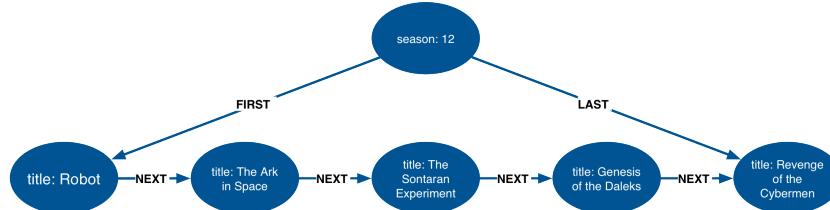
Exercise 3

Linked List Example



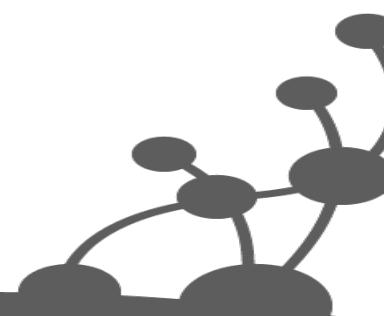
Season 12 of Doctor Who

- Add stories as they are broadcast
 - Maintain pointer to FIRST and LAST stories broadcast
- Find all stories broadcast so far
- Find latest story broadcast so far



Your Turn

- Clean the database
- Execute *setup.txt*
 - Creates root season node
- Execute *add-node.txt*
 - Adds *Robot*
- Modify the query to add more stories in broadcast order
- At each stage, view the results
 - MATCH (n) RETURN n



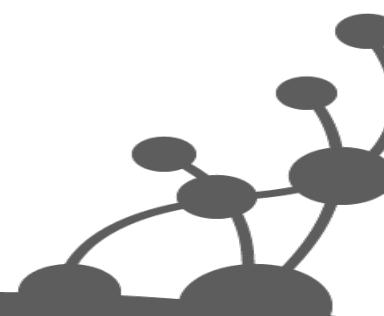
Add Story to Season

```
MERGE (season:Season{season:12})
MERGE (season)-[:LAST]->(newStory:Story{title:'Robot'})
WITH season, newStory

// Determine whether first story already exists
WITH season, newStory,
CASE WHEN NOT ((season)-[:FIRST]->()) THEN [1] ELSE []
END
AS firstExists

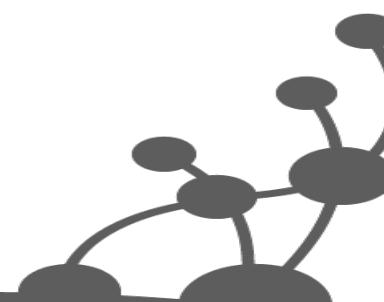
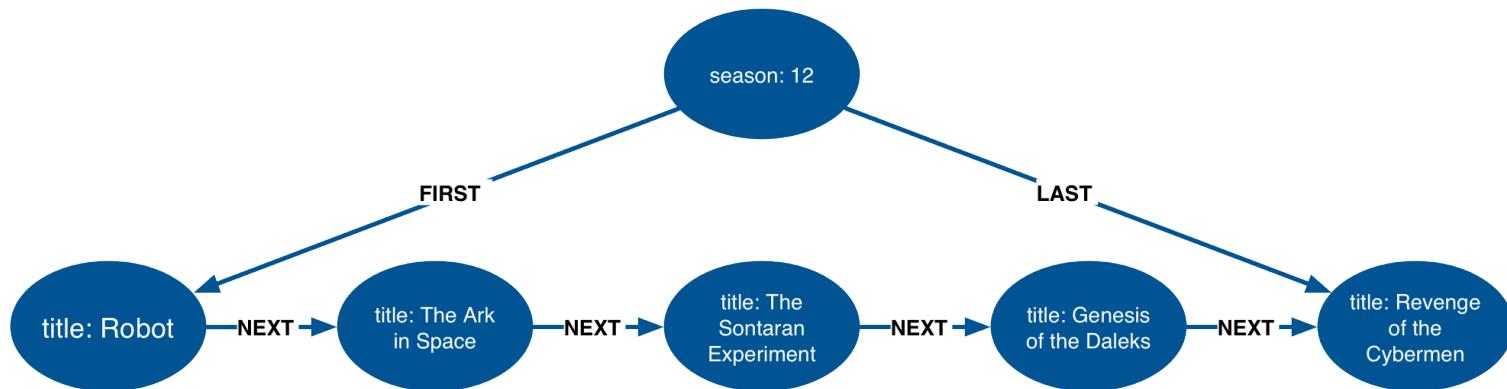
// Create FIRST rel newStory is first story
FOREACH (i IN firstExists | MERGE (season)-[:FIRST]->(newStory))
WITH season, newStory

// Delete old LAST relationship
MATCH (newStory)<-[:LAST]-(season)-[oldRel:LAST]->(oldLast)
DELETE oldRel
MERGE (oldLast)-[:NEXT]->(newStory)
```



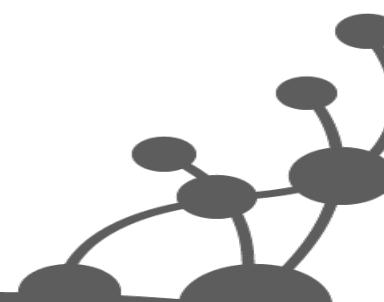
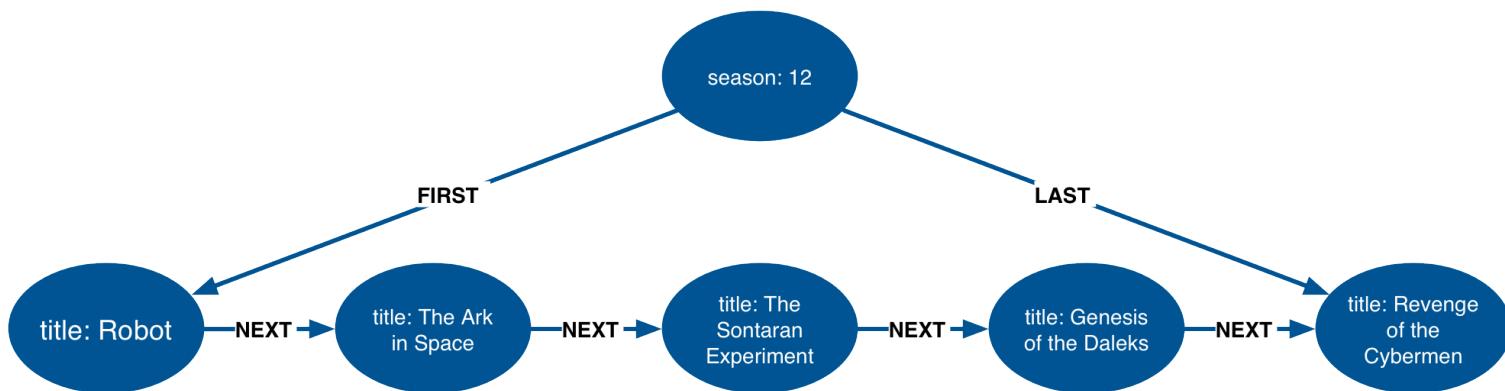
Query-1 - Find All Stories Broadcast So Far

```
MATCH (season:Season)-[:FIRST]->(firstStory)  
      -[:NEXT*0..]->(nextStory)  
RETURN nextStory.title AS nextStory
```

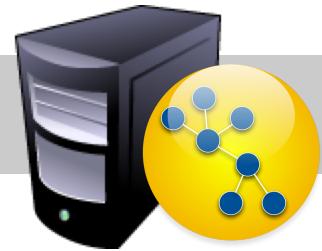


Query-2 - Find Last Story to be Broadcast

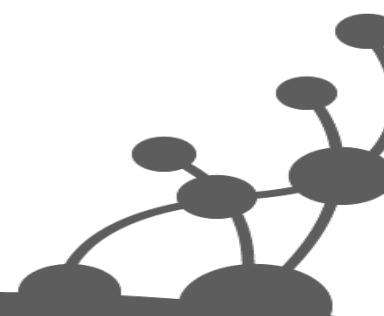
```
MATCH (season:Season)-[:LAST]->(lastStory)  
RETURN lastStory.title AS lastStory
```



In-Graph Indexes

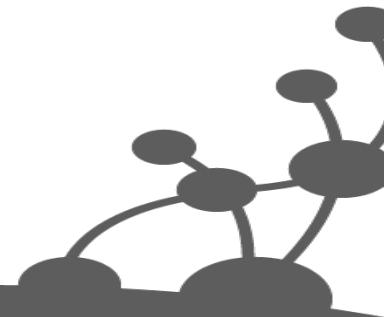


- Indexes are graphs:
 - B-tree (binary search)
 - R-tree (spatial access, multi-dimensional information)

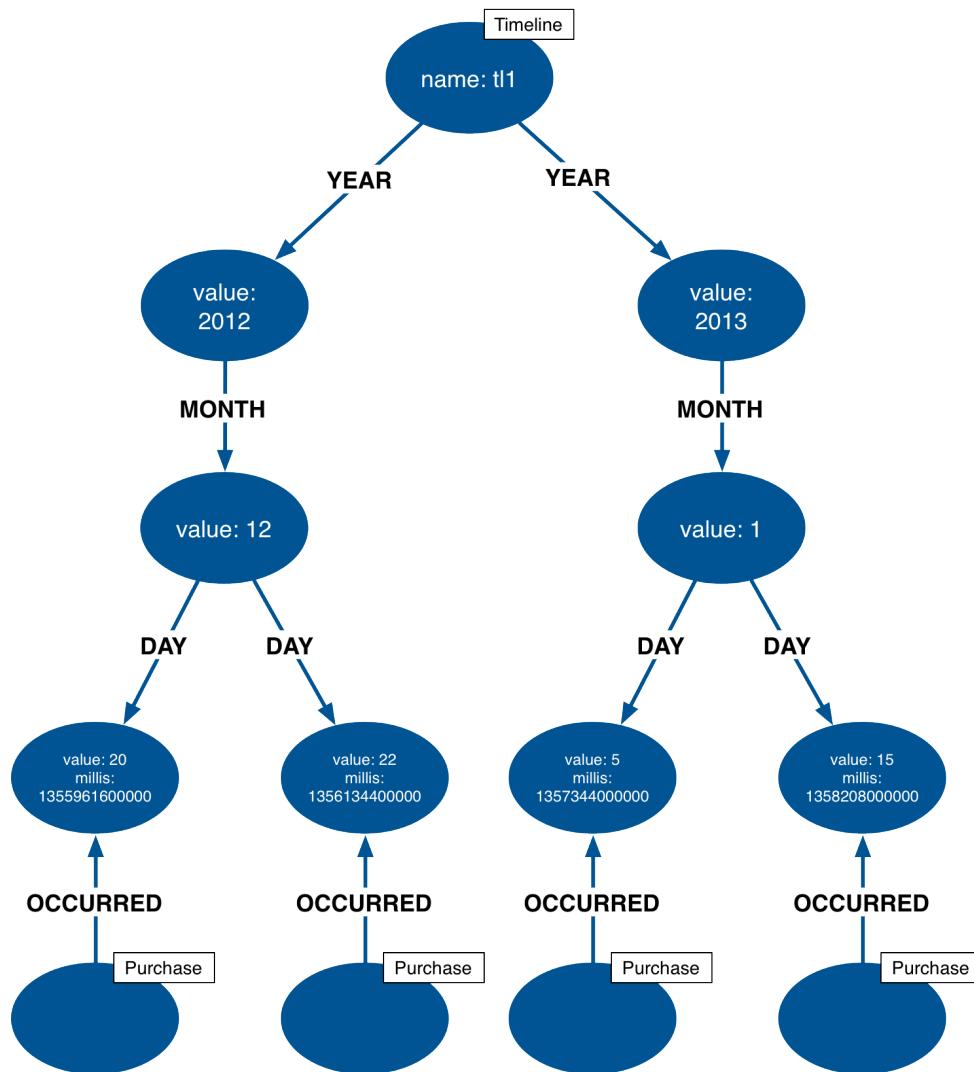


Timeline Tree

- Discrete events
 - No natural relationships to other events
- You need to find events at differing levels of granularity
 - Between two days
 - Between two months
 - Between two minutes

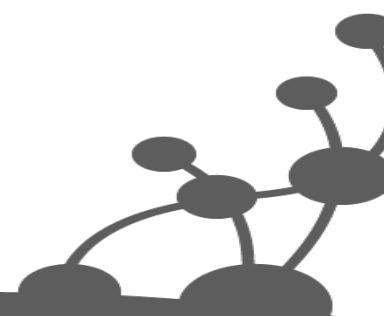


Example Timeline Tree



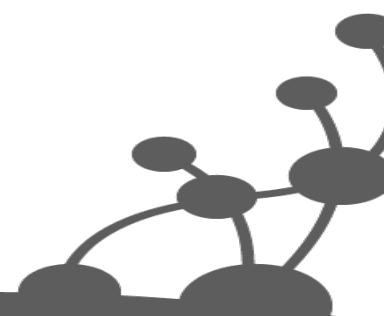
Exercise 4

Timeline Tree Example



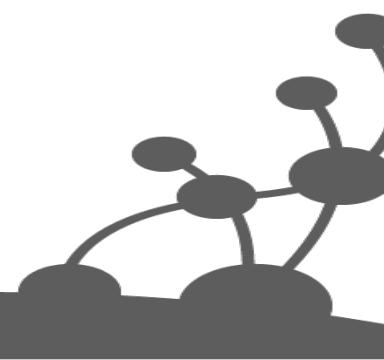
Your Turn

- Clean the database
- Execute *create-1.txt* to *create-6.txt* in order
- At each stage, view the results
 - MATCH (n) RETURN n



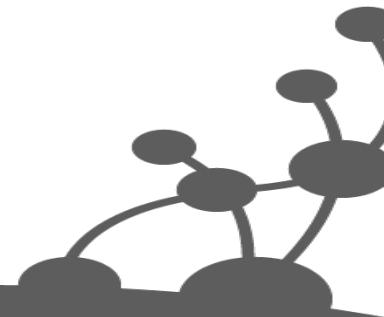
Lazily Insert Date Elements

```
MERGE (timeline:Timeline{name:'timeline-1'})  
MERGE (timeline)-[:YEAR]->(year{value:2007})  
MERGE (year)-[:MONTH]->(month{value:1})  
MERGE (month)-[:DAY]->(day{value:14})  
MERGE (day)<-[:OCCURRED]-  
(n:Purchase{name:'purchase-1'})
```



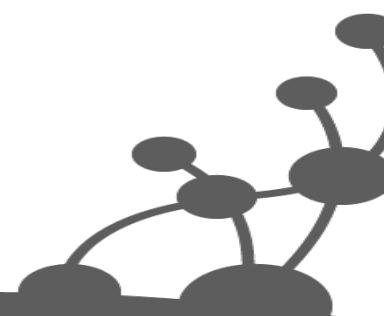
Create or Insert Root

```
MERGE (timeline:Timeline{name:'timeline-1'})  
MERGE (timeline)-[:YEAR]->(year{value:2007})  
MERGE (year)-[:MONTH]->(month{value:1})  
MERGE (month)-[:DAY]->(day{value:14})  
MERGE (day)<-[:OCCURRED]-  
(n:Purchase{name:'purchase-1'})
```



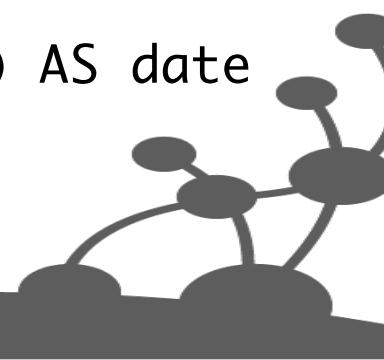
The Add ‘Locally Unique’ Nodes

```
MERGE (timeline:Timeline{name:'timeline-1'})  
MERGE (timeline)-[:YEAR]->(year{value:2007})  
MERGE (year)-[:MONTH]->(month{value:1})  
MERGE (month)-[:DAY]->(day{value:14})  
MERGE (day)<-[:OCCURRED]-  
(n:Purchase{name:'purchase-1'})
```

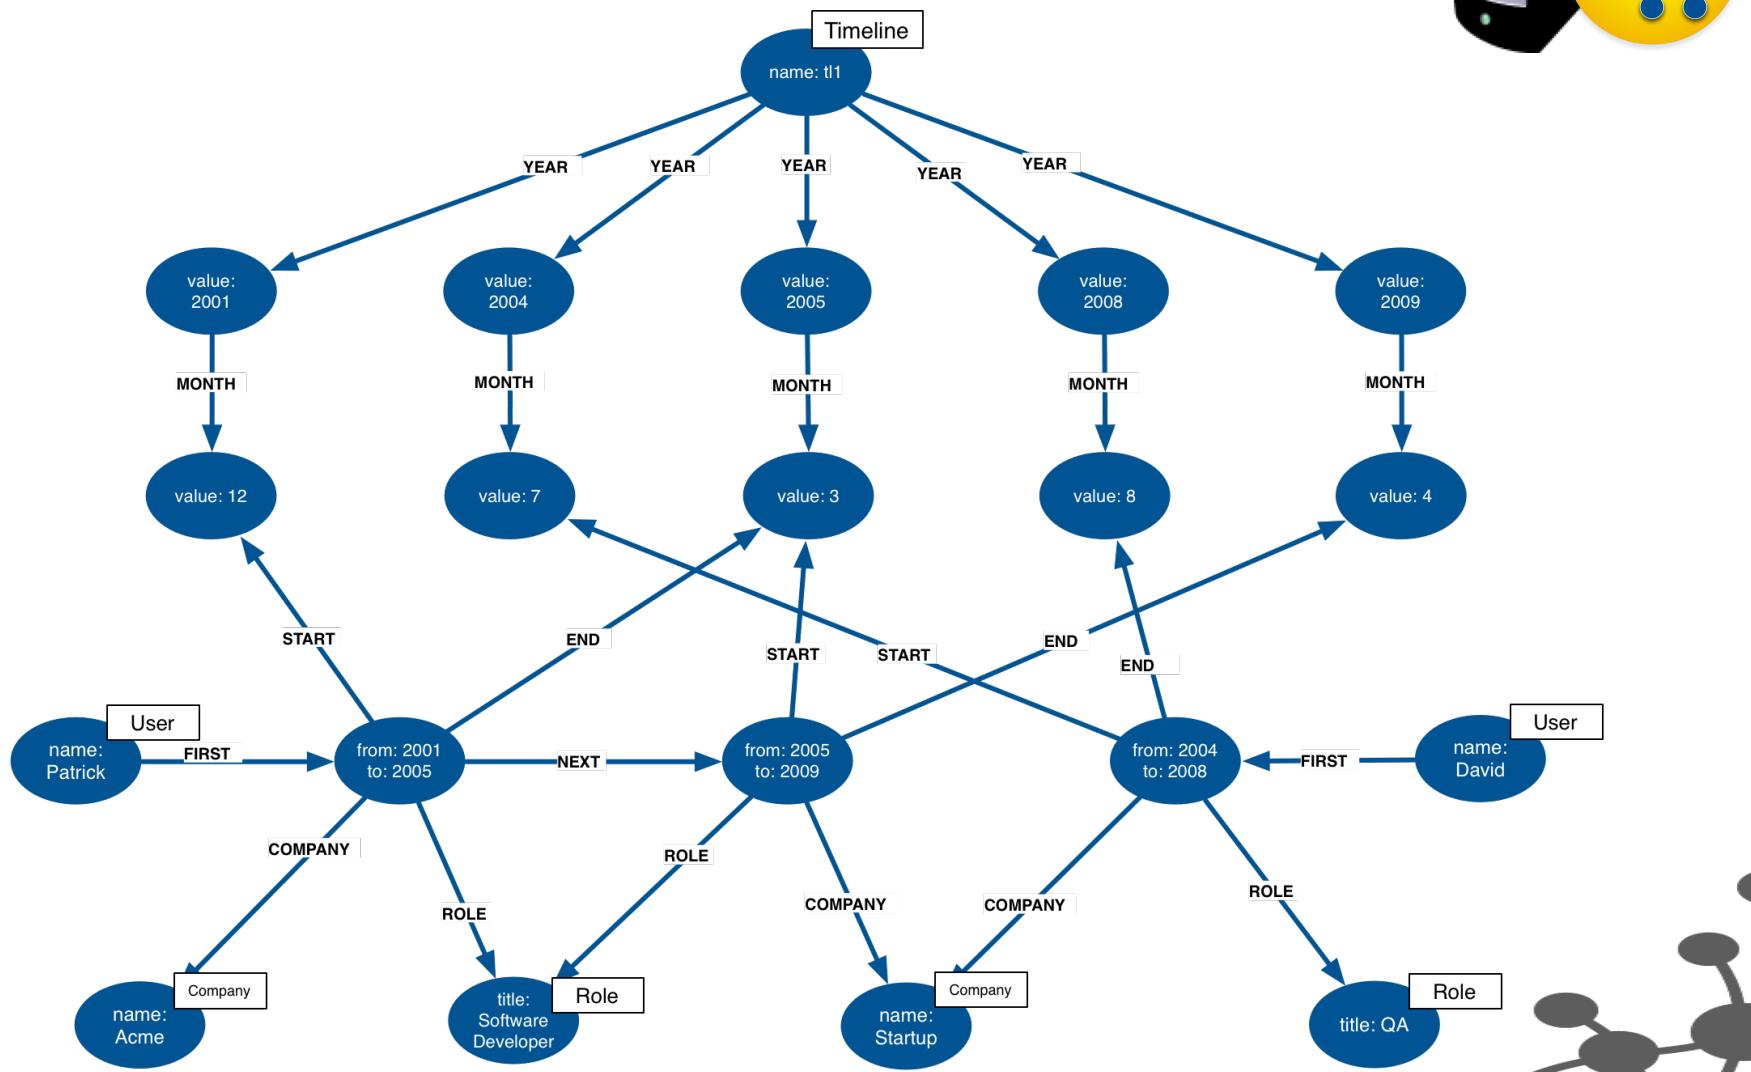


Query-1 - Get All Events Between Two Dates

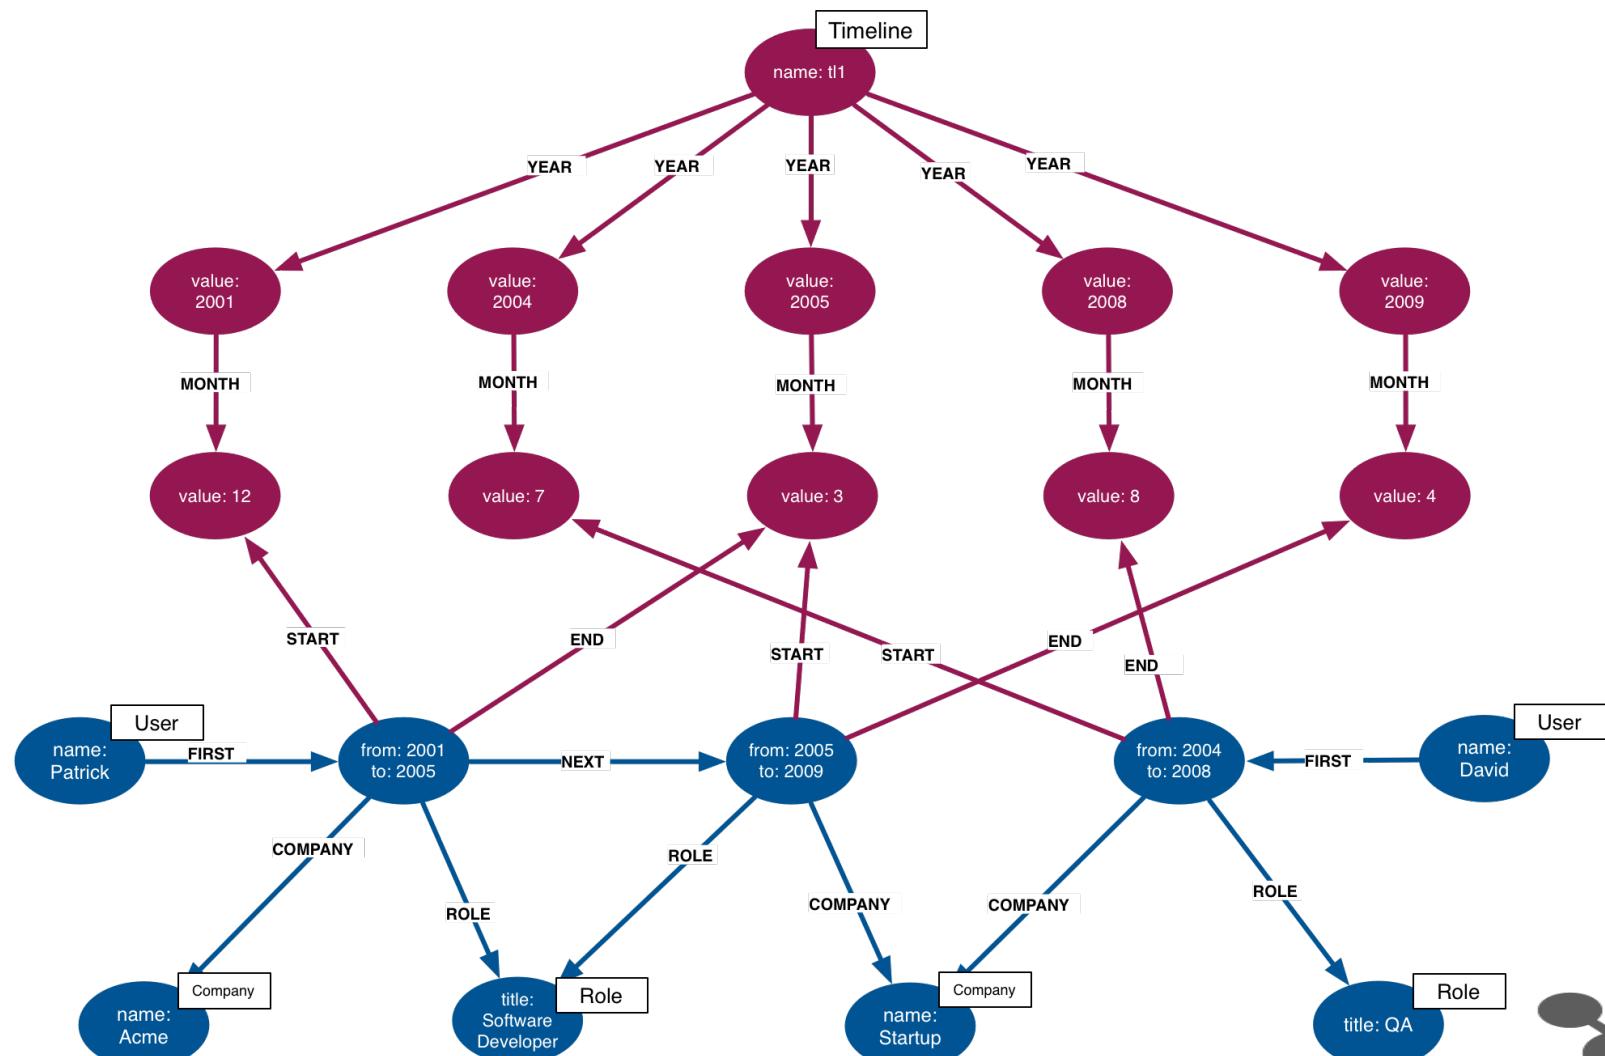
```
MATCH (timeline:Timeline{name:'timeline-1'})  
  -[:YEAR]->(y)  
  -[:MONTH]->(m)  
  -[:DAY]->(d)<-[:OCCURRED]-(n)  
WHERE (y.value > {startYear} AND y.value < {endYear})  
OR ({startYear} = {endYear})  
OR (y.value = {startYear}  
    AND ((m.value > {startMonth})  
        OR (m.value = {startMonth}  
            AND d.value >= {startDay})))  
OR (y.value = {endYear}  
    AND ((m.value < {endMonth})  
        OR (m.value = {endMonth}  
            AND d.value <= {endDay})))  
RETURN n.name,  
       (d.value + " - " + m.value + " - " + y.value) AS date
```



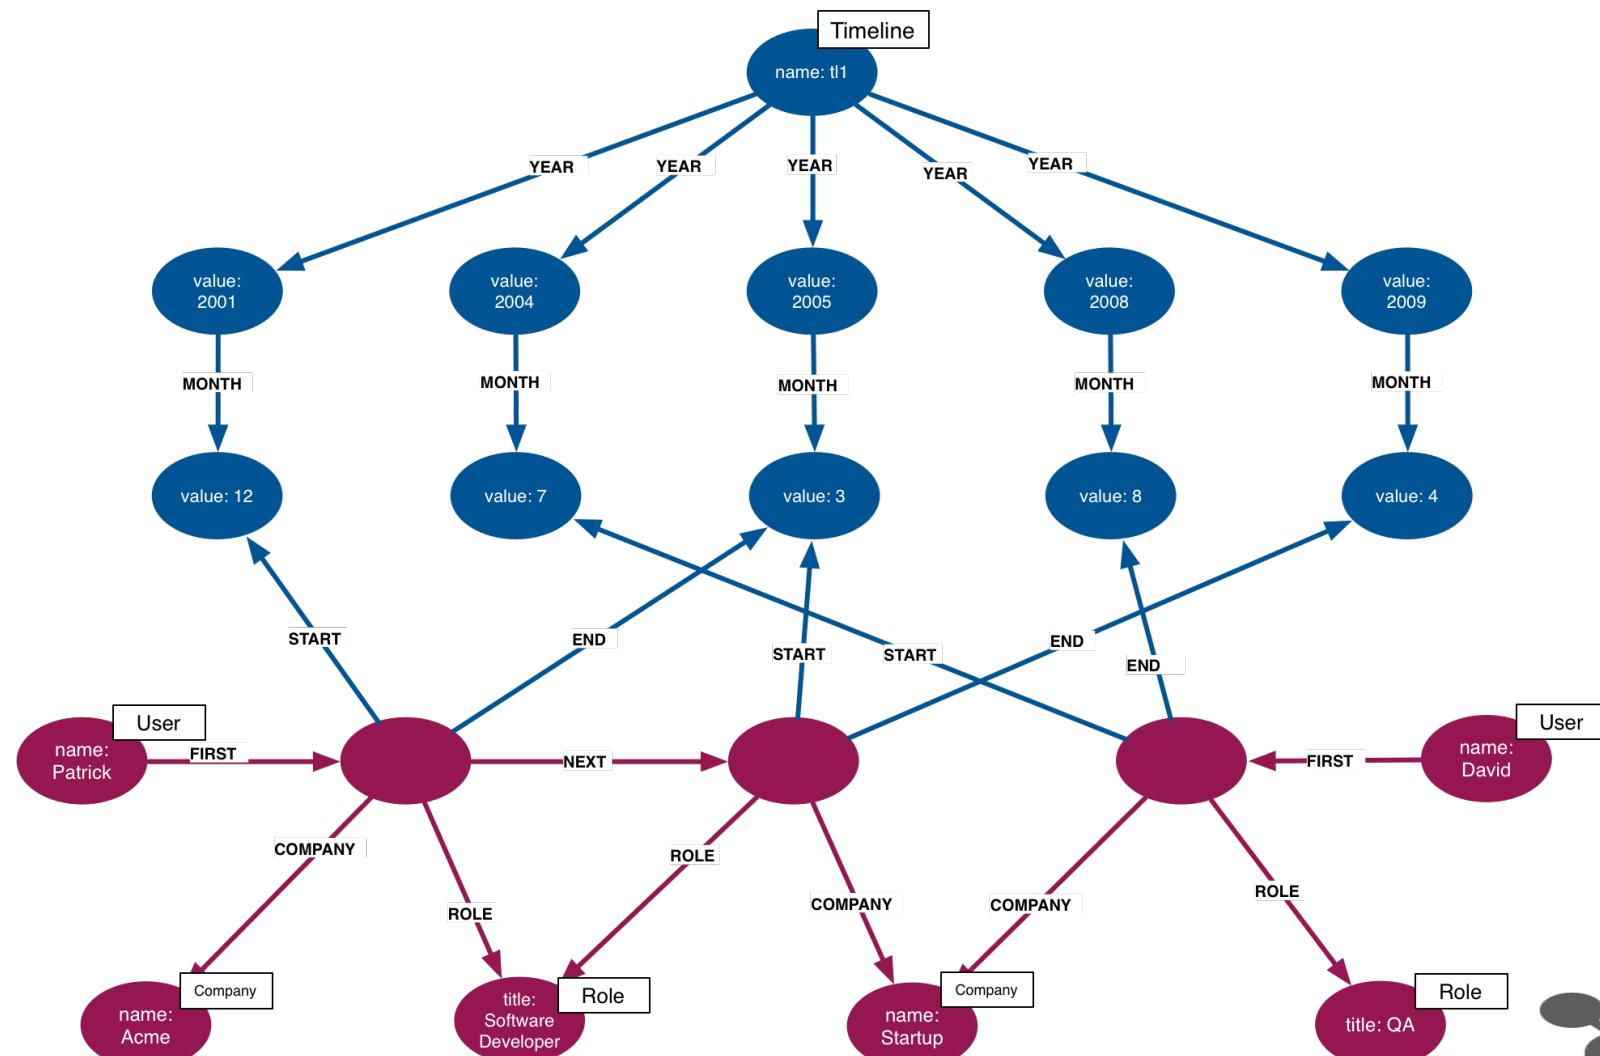
Composing Structures



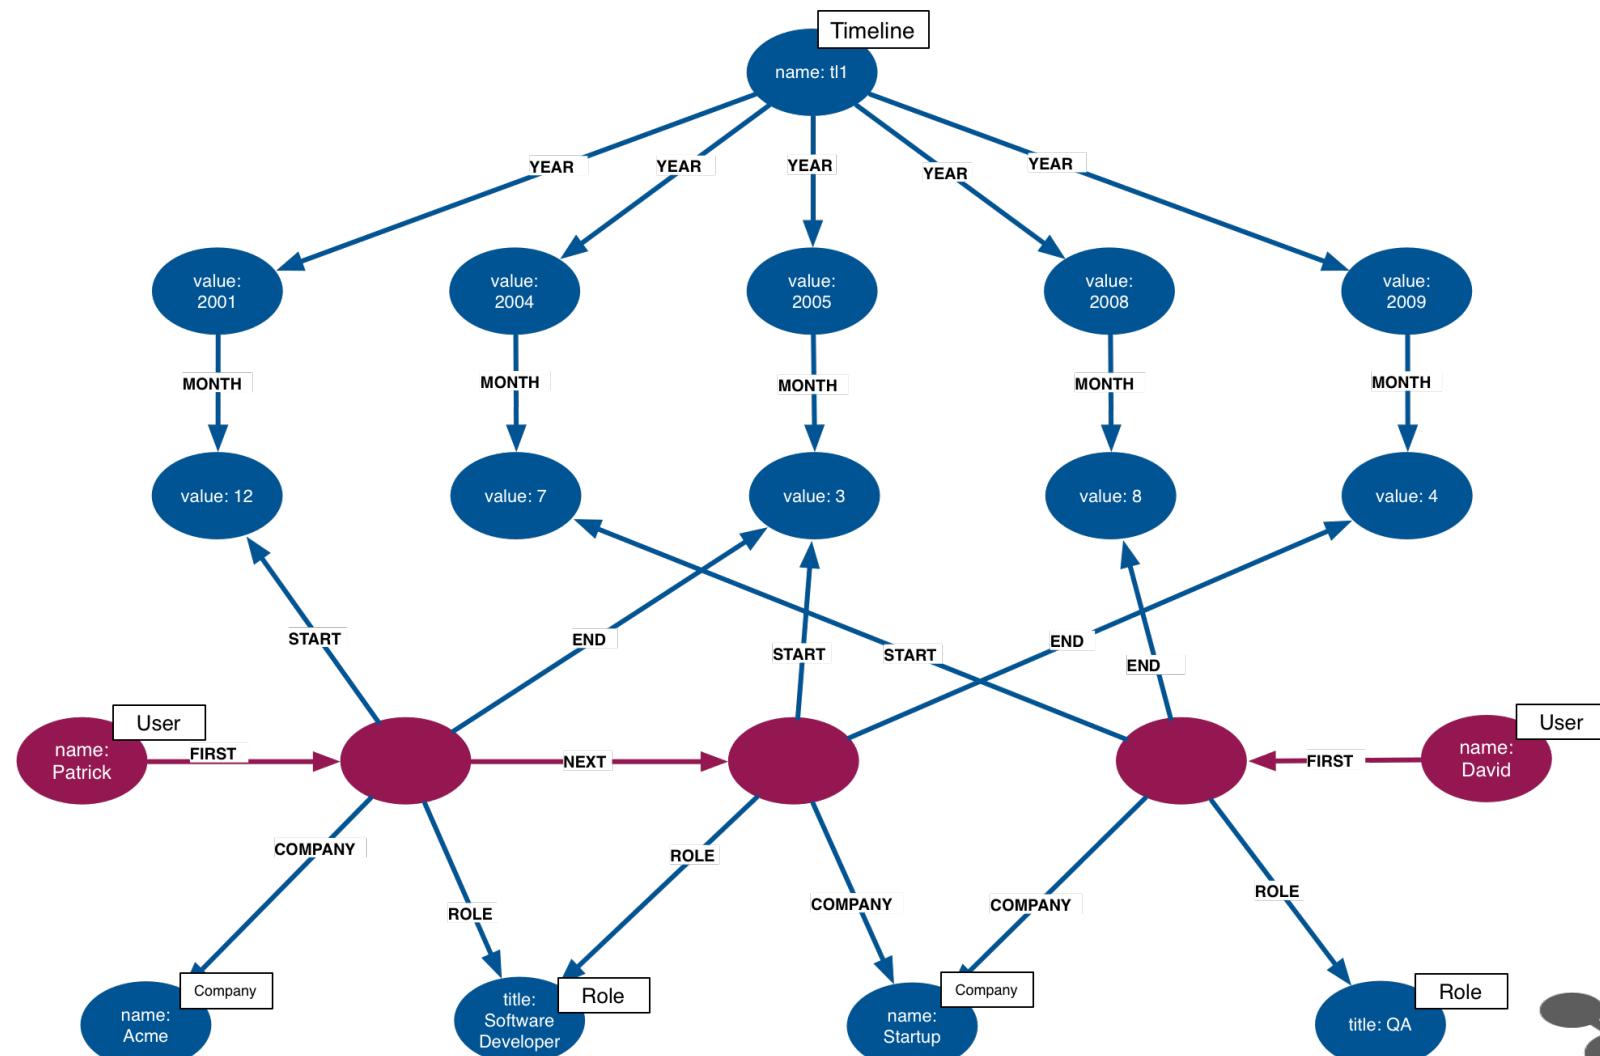
Timeline Tree



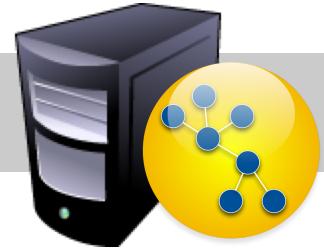
Intermediate Nodes



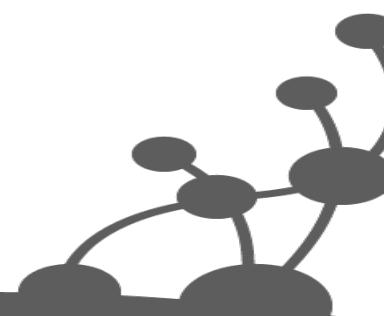
Linked Lists



Versioning Graphs

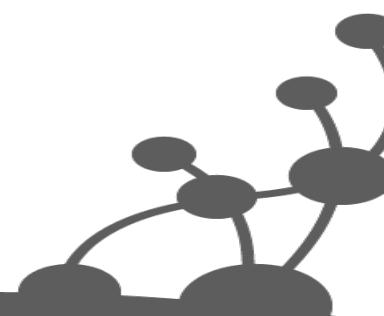


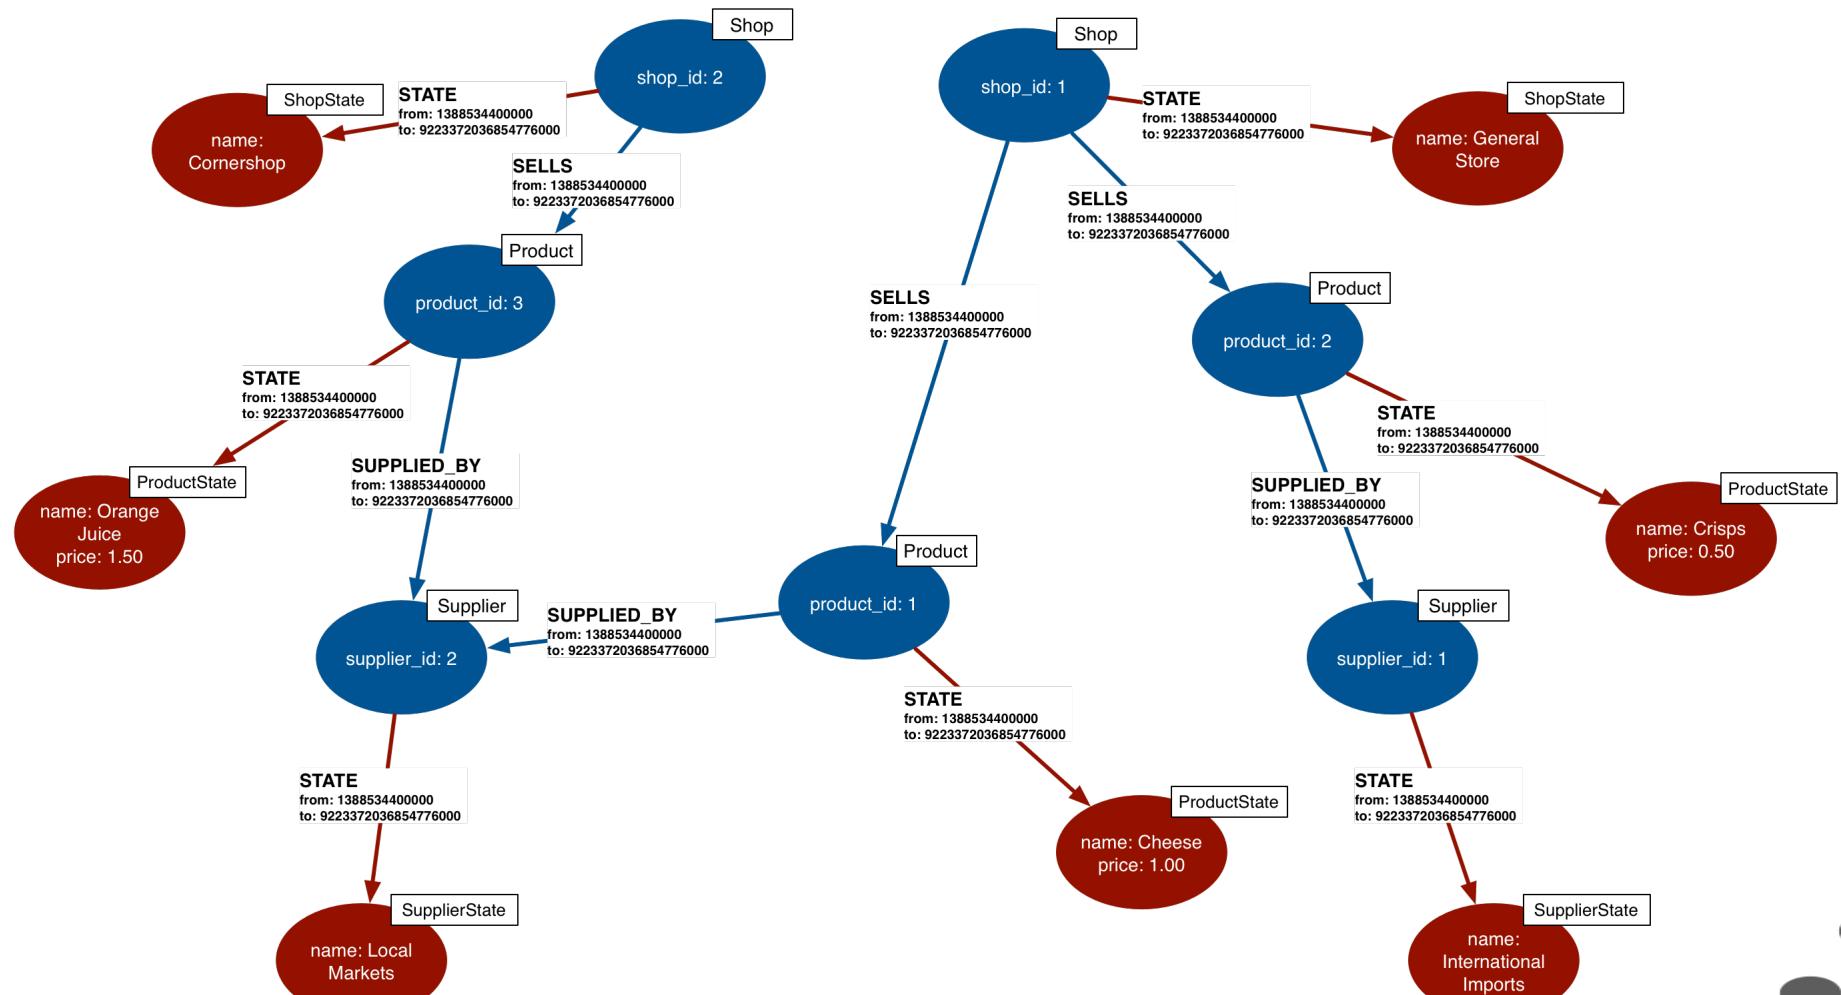
- Time-based
 - Universal versioning schema
 - Discrete, continuous sequence
 - Millis since the epoch



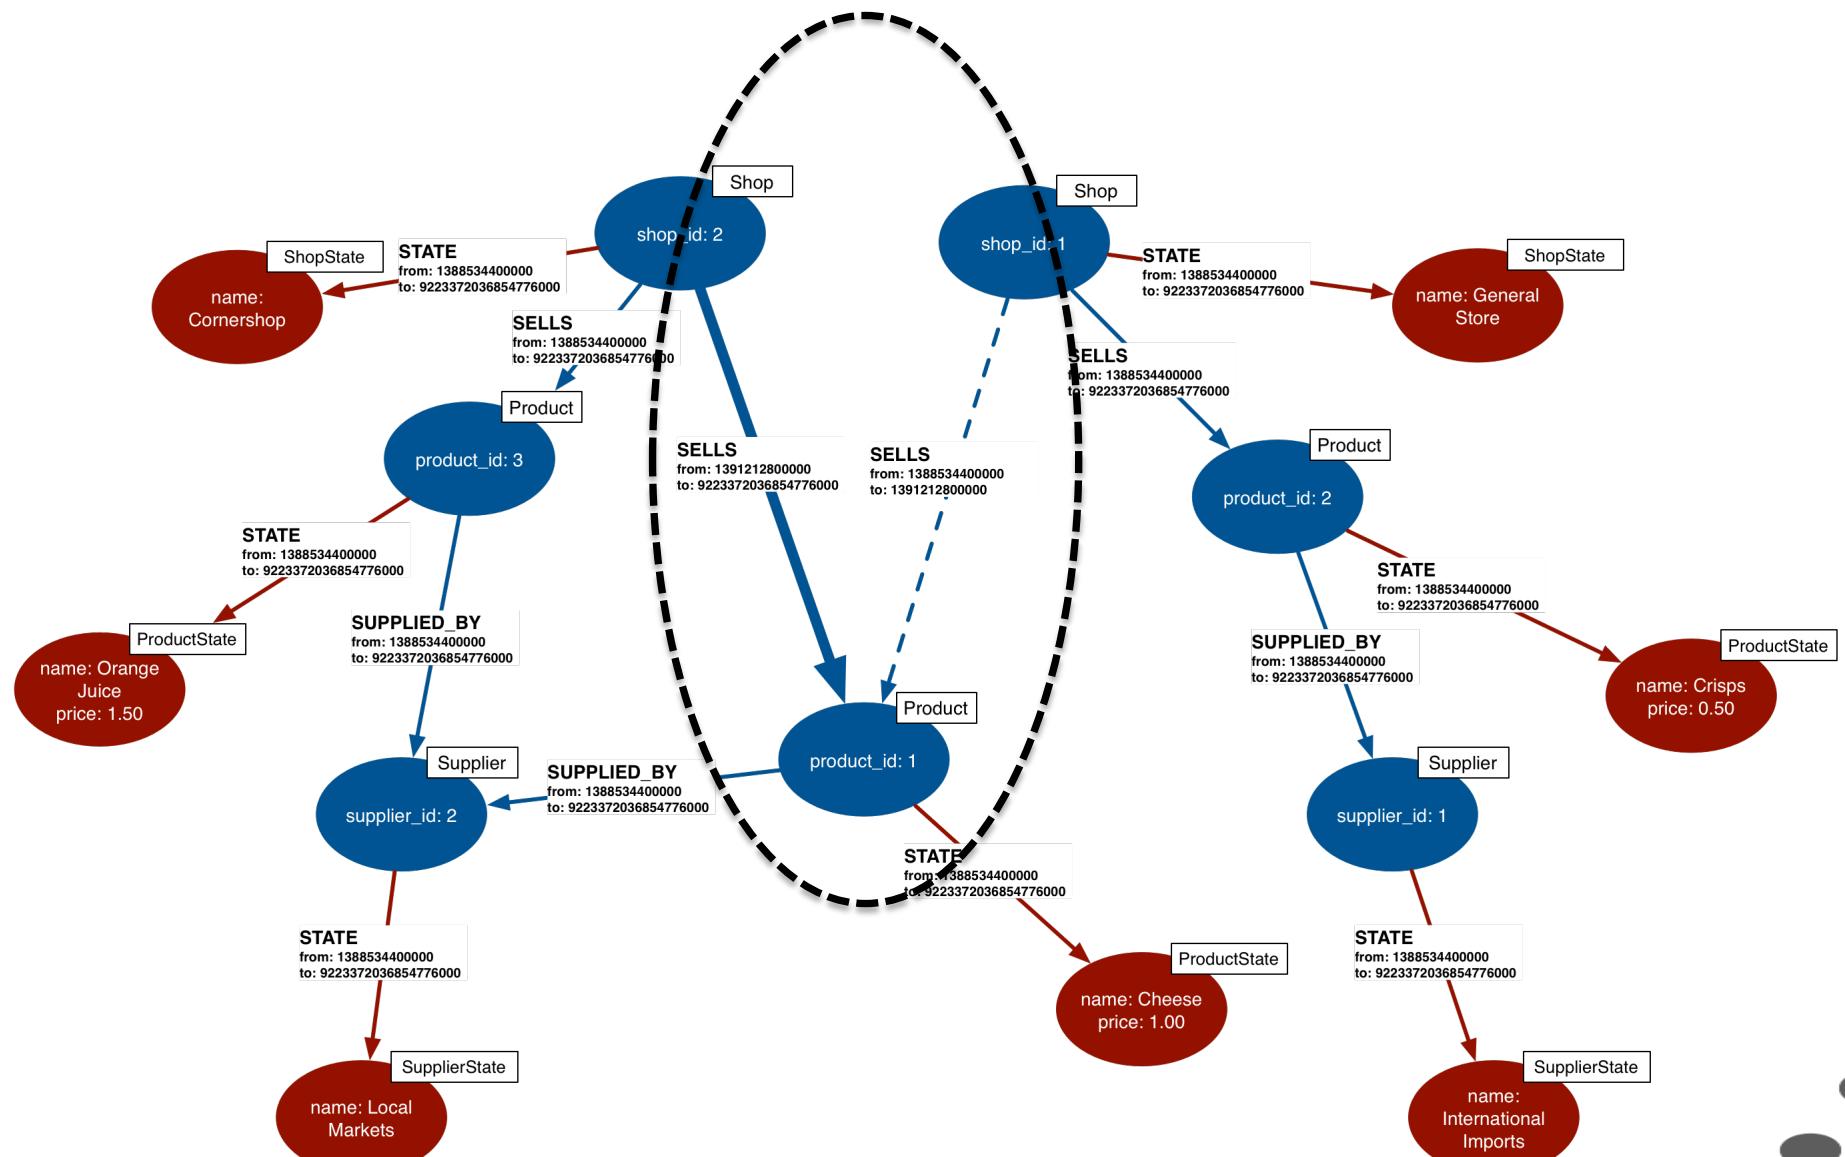
Separate Structure from State

- Structure
 - Identity nodes
 - Placeholders
 - Timestamped identity relationships
 - i.e. normal domain relationships
- State
 - State nodes
 - Snapshot of entity state
 - Timestamped state relationships





1 Jan 2014 = 1388534400000



1 Feb 2014 = 1391212800000

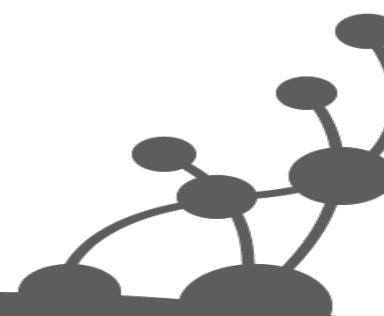
Find Current Structural Relationship

```
MATCH (p:Product{product_id:1})<-[r:SELLS]-(:Shop)
WHERE r.to = 9223372036854775807
MATCH (s:Shop{shop_id:2})
SET r.to = 1391212800000
CREATE (s)
-[:SELLS{from:1391212800000,to:9223372036854775807}]->(p)
```

9223372036854775807 = End of Time = EOT

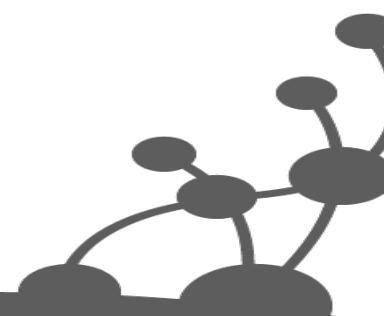
Archive Structural Relationship

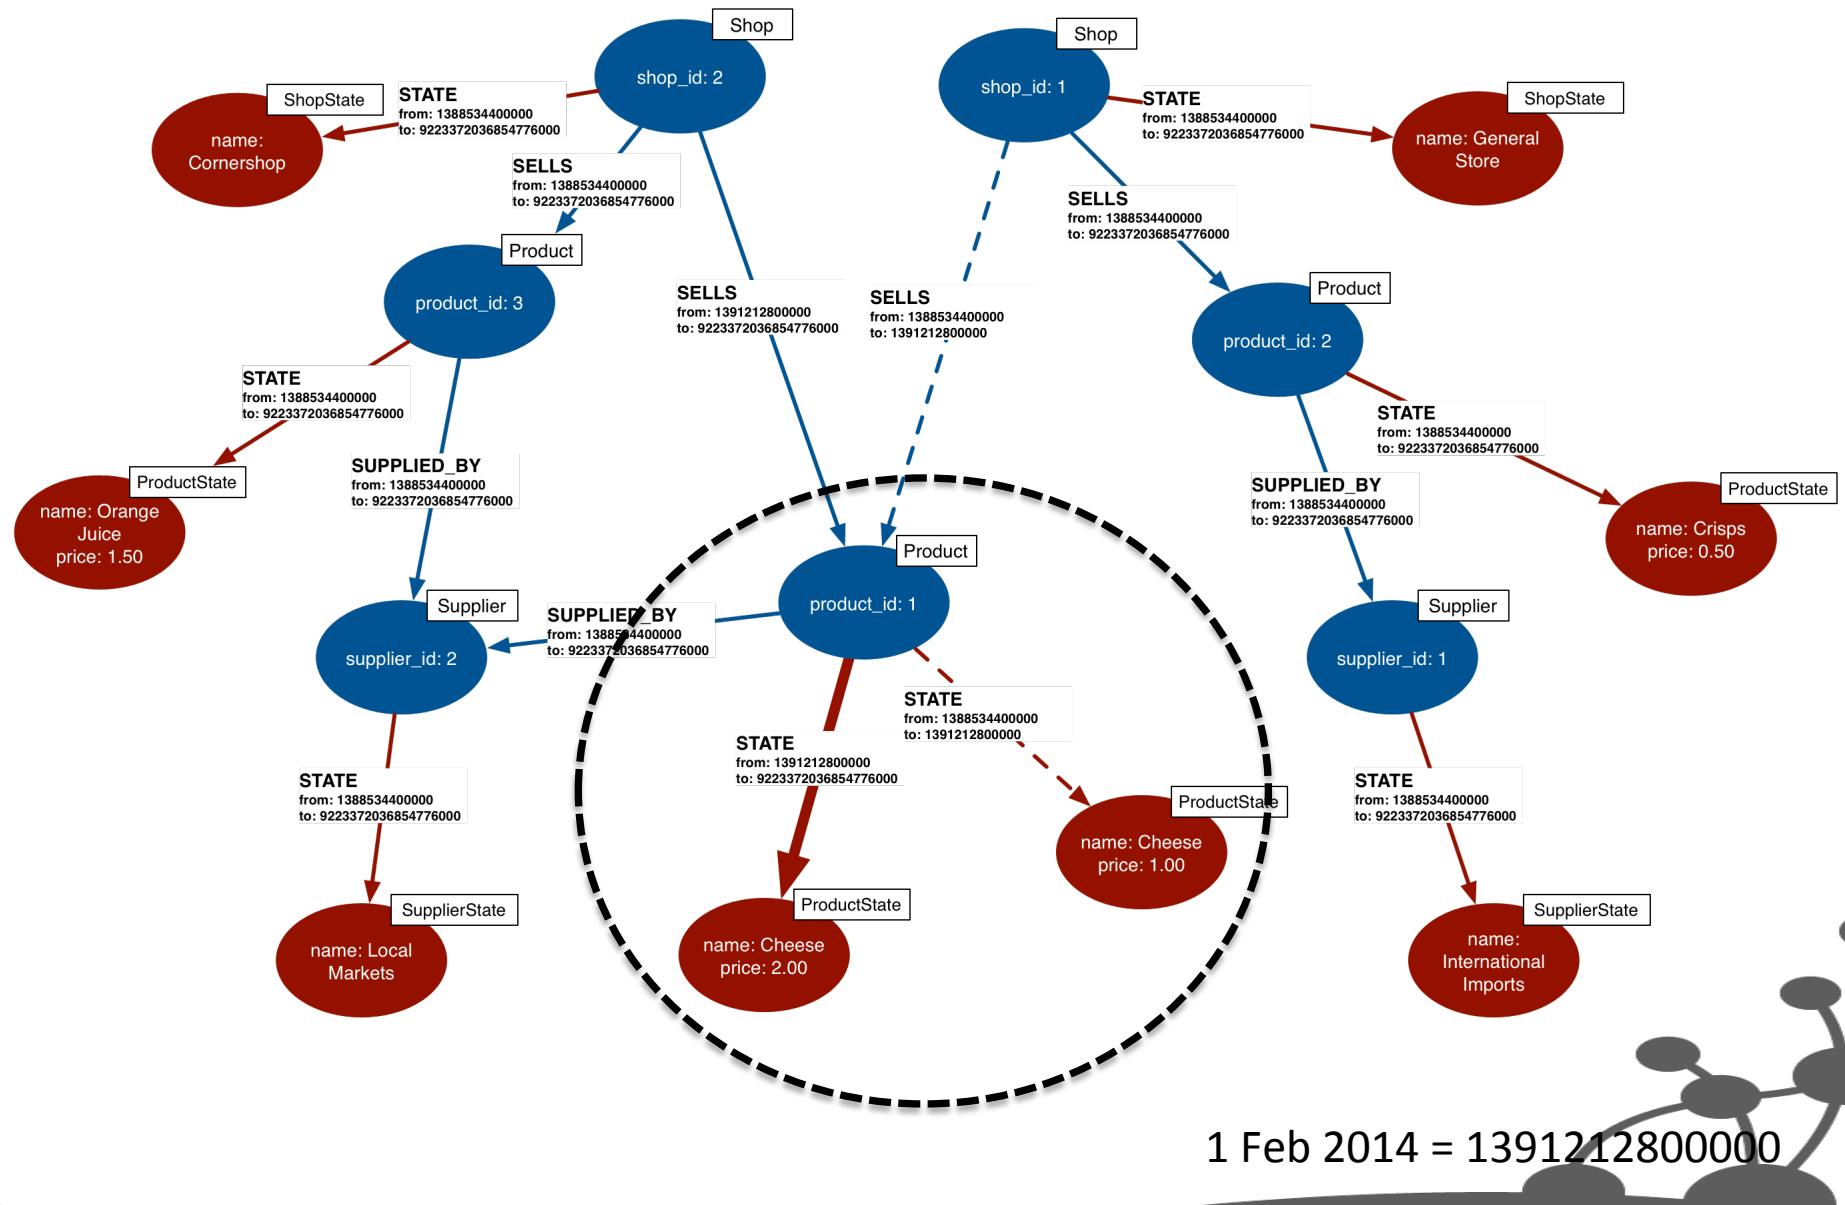
```
MATCH (p:Product{product_id:1})<-[r:SELLS]-(:Shop)
WHERE r.to = 9223372036854775807
MATCH (s:Shop{shop_id:2})
SET r.to = 1391212800000
CREATE (s)
    -[:SELLS{from:1391212800000,to:9223372036854775807}]->(p)
```



Create New Structural Relationships

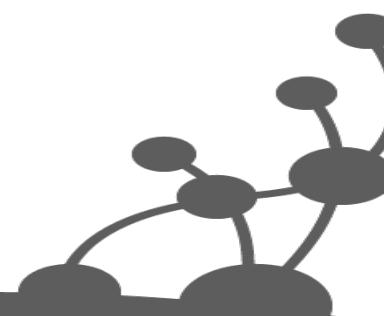
```
MATCH (p:Product{product_id:1})<-[r:SELLS]-(:Shop)
WHERE r.to = 9223372036854775807
MATCH (s:Shop{shop_id:2})
SET r.to = 139121280000
CREATE (s)
-[:SELLS{from:139121280000,to:9223372036854775807}]->(p)
```





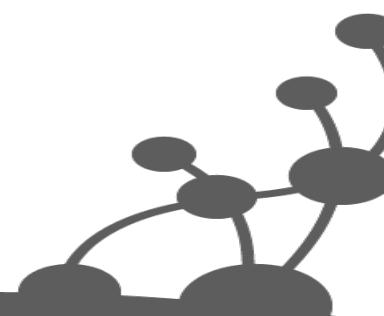
All Products Sold by Shop 1 on 5 January 2014

```
MATCH (s:Shop{shop_id:1})-[r1:SELLS]->(p:Product)
WHERE (r1.from <= 1388880000000 AND r1.to > 1388880000000)
MATCH (p)-[r2:STATE]->(ps:ProductState)
WHERE (r2.from <= 1388880000000 AND r2.to > 1388880000000)
RETURN p.product_id AS productId,
       ps.name AS product,
       ps.price AS price
ORDER BY price DESC
```



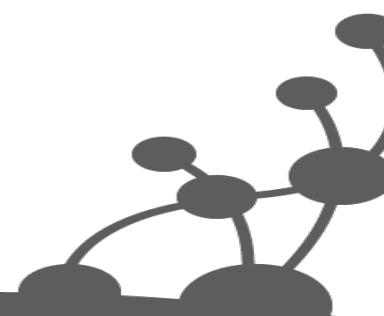
Find Structure

```
MATCH (s:Shop{shop_id:1})-[r1:SELLS]->(p:Product)
WHERE (r1.from <= 1388880000000 AND r1.to > 1388880000000)
MATCH (p)-[r2:STATE]->(ps:ProductState)
WHERE (r2.from <= 1388880000000 AND r2.to > 1388880000000)
RETURN p.product_id AS productId,
       ps.name AS product,
       ps.price AS price
ORDER BY price DESC
```



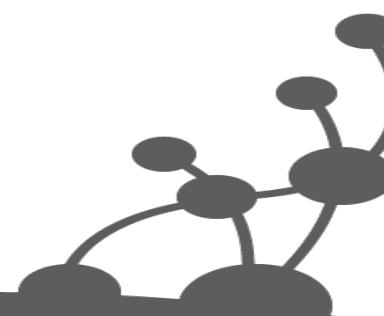
Find State

```
MATCH (s:Shop{shop_id:1})-[r1:SELLS]->(p:Product)
WHERE (r1.from <= 1388880000000 AND r1.to > 1388880000000)
MATCH (p)-[r2:STATE]->(ps:ProductState)
WHERE (r2.from <= 1388880000000 AND r2.to > 1388880000000)
RETURN p.product_id AS productId,
       ps.name AS product,
       ps.price AS price
ORDER BY price DESC
```

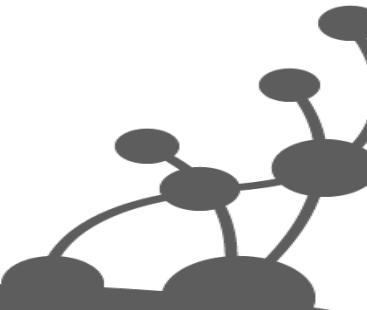


Return Results

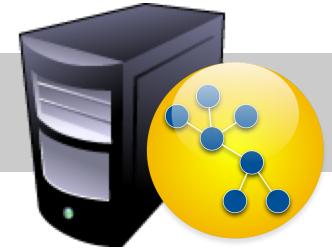
```
MATCH (s:Shop{shop_id:1})-[r1:SELLS]->(p:Product)
WHERE (r1.from <= 1388880000000 AND r1.to > 1388880000000)
MATCH (p)-[r2:STATE]->(ps:ProductState)
WHERE (r2.from <= 1388880000000 AND r2.to > 1388880000000)
RETURN p.product_id AS productId,
       ps.name AS product,
       ps.price AS price
ORDER BY price DESC
```



Evolving a Graph Model



Refactoring

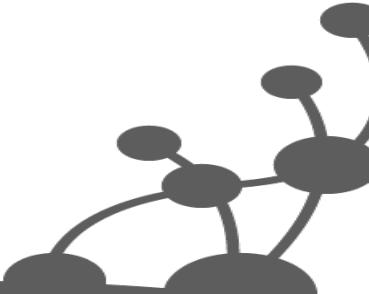


Definition

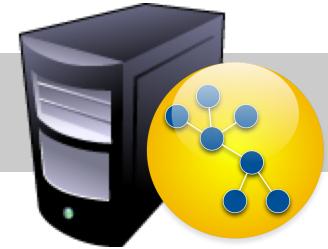
- Restructure graph without changing informational semantics

Reasons

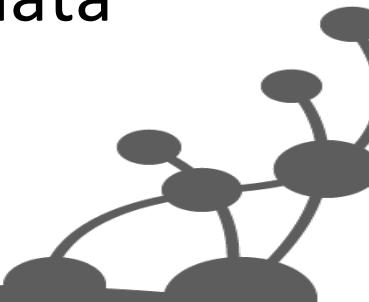
- Improve design
- Enhance performance
- Accommodate new functionality
- Enable iterative and incremental development of data model



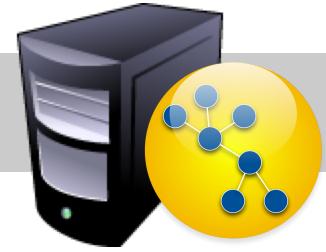
Data Migrations



- Execute in repeatable order
- Backup database
- Execute in batches
 - Unbounded results will generate large transactions and may trigger Out of Memory exceptions
- Apply migrations to test data to ensure existing functionality doesn't break
- Ensure application can accommodate old and new structures if performing against live data



Extract Node from Property

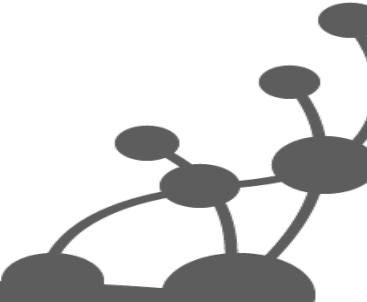


Problem

- You've modeled an attribute as a property with a simple value, but now need to:
 - Qualify the attribute semantics AND/OR
 - Introduce a complex value AND/OR
 - Reify the relationship represented by the value

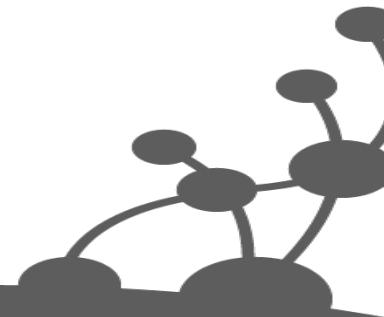
Solution

- Create a new node per unique property value
- Connect existing nodes to the new property nodes
- Remove the old property

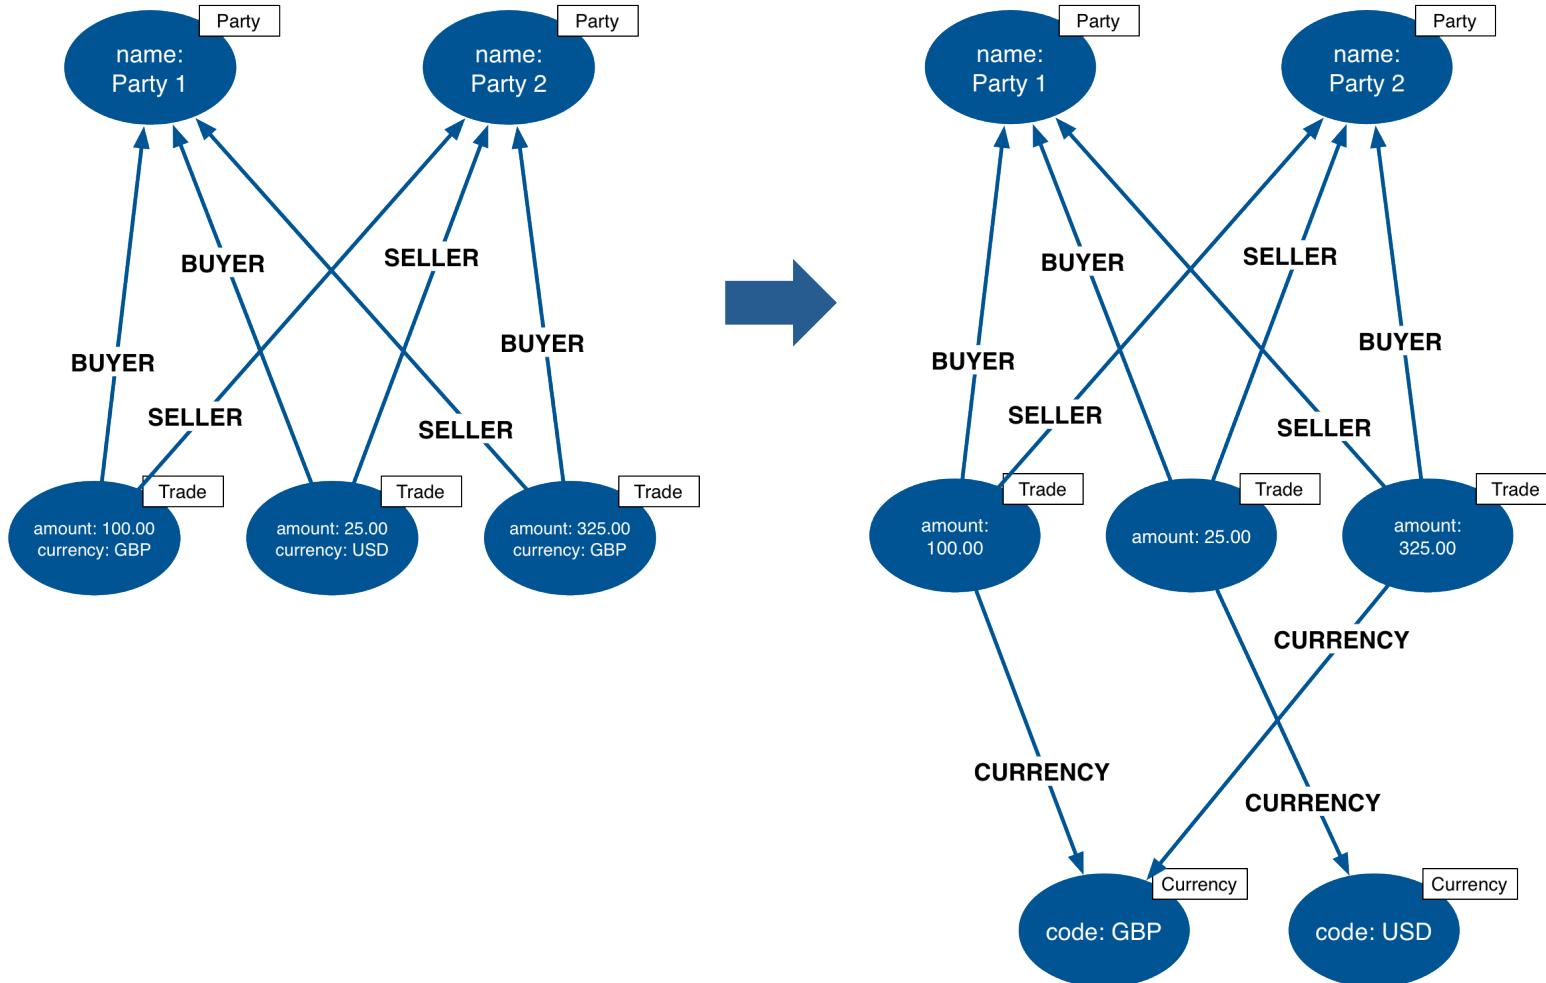


Exercise 5

Extract Node From Property

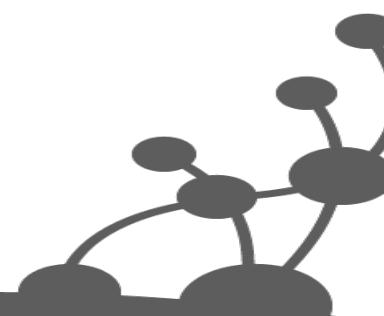


Example: (n.currency) to (:Currency)



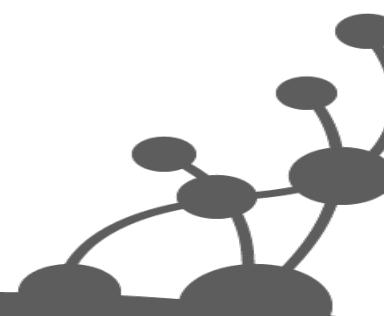
Your Turn

- Clean the database
- Execute *setup.txt*
- View the results
 - MATCH (n) RETURN n
- Execute *update-1.txt* repeatedly, until *numberRemoved* is zero
 - At each stage, view the results



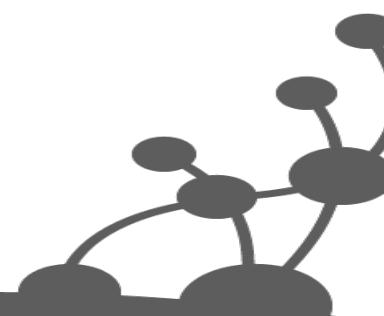
Extract Node From Property

```
MATCH (t:Trade) WHERE has(t.currency)
WITH t LIMIT {batchSize}
MERGE (c:Currency{code:t.currency})
MERGE (t)-[:CURRENCY]->(c)
REMOVE t.currency
RETURN count(t) AS numberRemoved
```



Select Batch of Nodes With Property

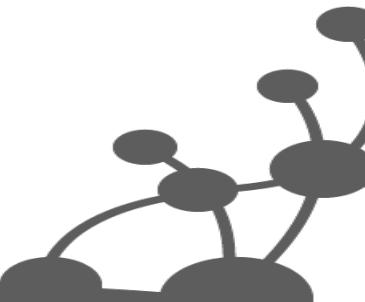
```
MATCH (t:Trade) WHERE has(t.currency)
WITH t LIMIT {batchSize}
MERGE (c:Currency{code:t.currency})
MERGE (t)-[:CURRENCY]->(c)
REMOVE t.currency
RETURN count(t) AS numberRemoved
```



Create Property Node

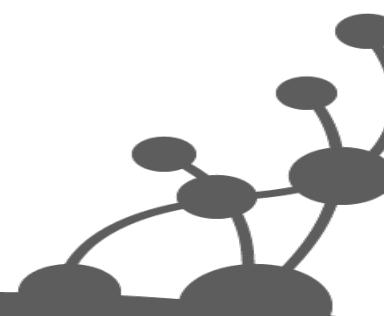
```
MATCH (t:Trade) WHERE has(t.currency)  
WITH t LIMIT {batchSize}  
MERGE (c:Currency{code:t.currency})  
MERGE (t)-[:CURRENCY]->(c)  
REMOVE t.currency  
RETURN count(t) AS numberRemoved
```

Copy property value
from existing node



Relate Existing Node to Property Node

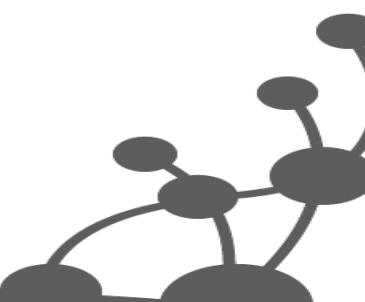
```
MATCH (t:Trade) WHERE has(t.currency)
WITH t LIMIT {batchSize}
MERGE (c:Currency{code:t.currency})
MERGE (t)-[:CURRENCY]->(c)
REMOVE t.currency
RETURN count(t) AS numberRemoved
```



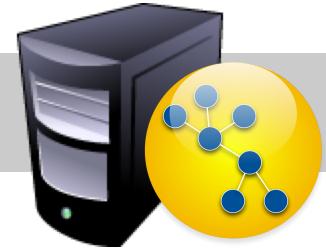
Remove Old Property

```
MATCH (t:Trade) WHERE has(t.currency)
WITH t LIMIT {batchSize}
MERGE (c:Currency{code:t.currency})
MERGE (t)-[:CURRENCY]->(c)
REMOVE t.currency
RETURN count(t) AS numberRemoved
```

Repeat until
numberRemoved
is zero



Extract Node from Array Property

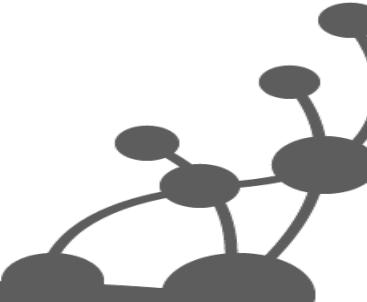


Problem

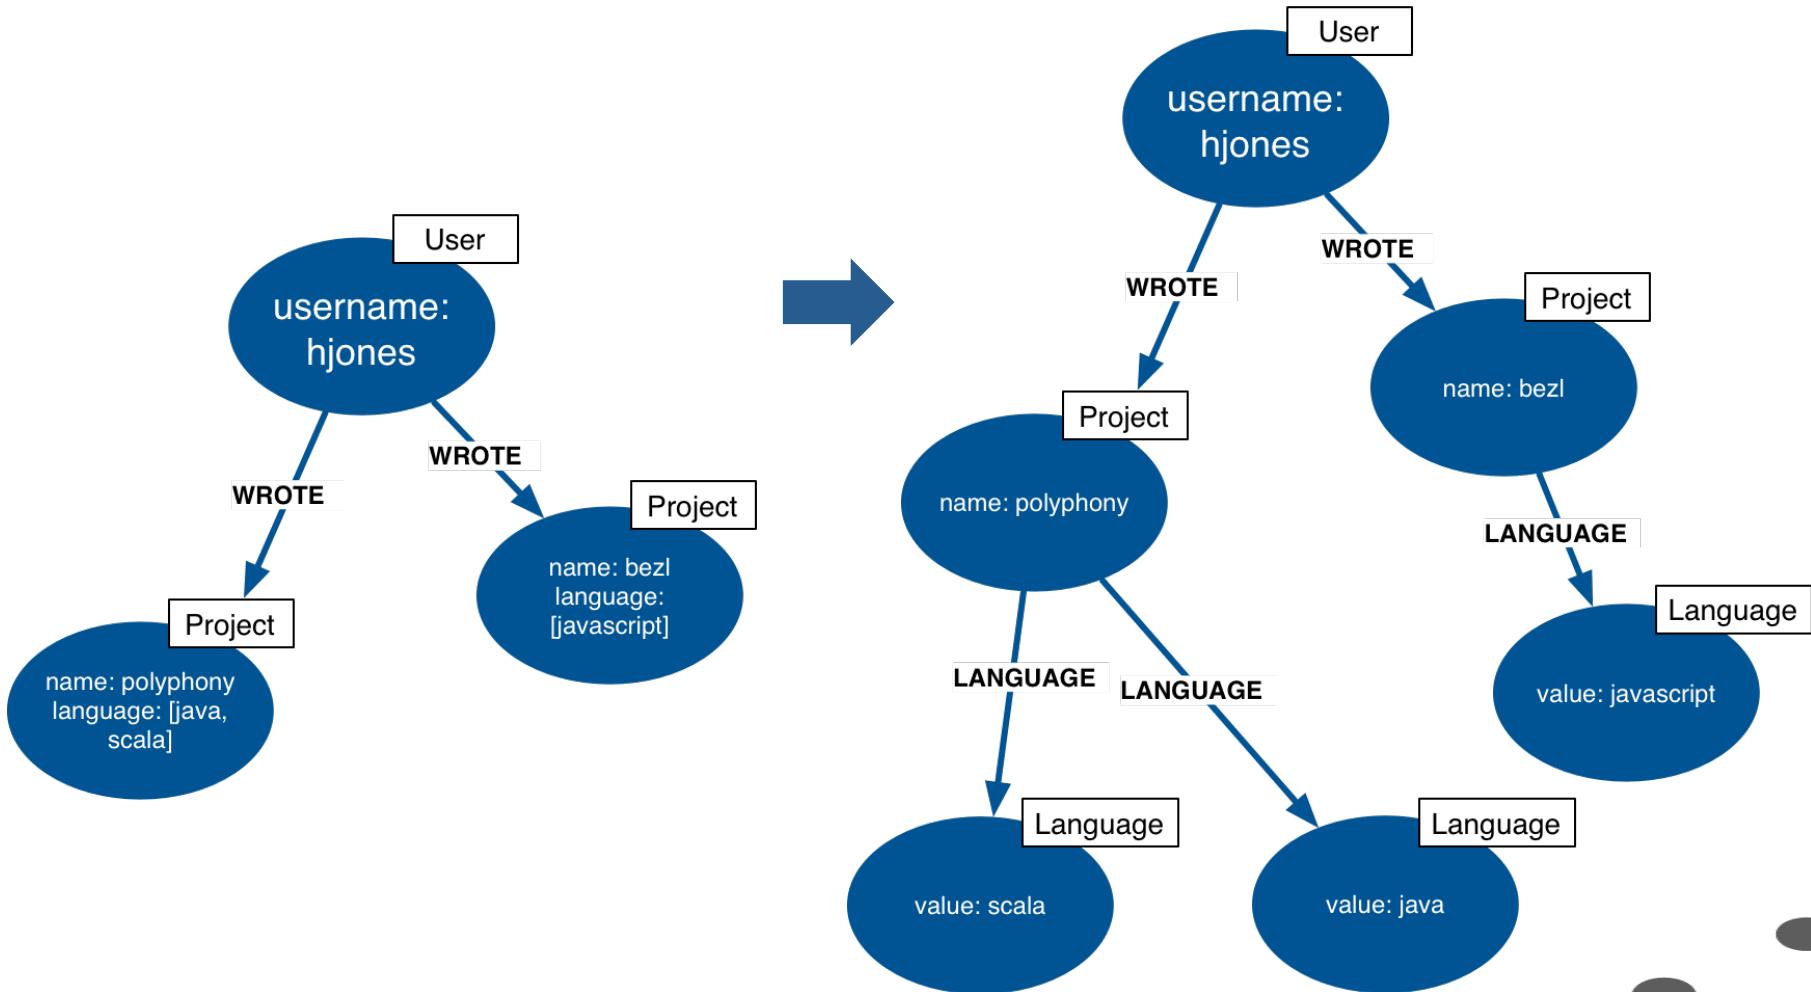
- You've modeled an attribute as a property with an array value, but now need to:
 - Qualify the attribute semantics AND/OR
 - Introduce a complex value AND/OR
 - Reify the relationship represented by the value

Solution

- Create a new node per unique property value
- Connect existing nodes to the new property nodes
- Remove the old property

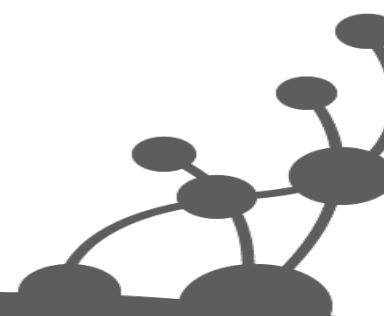


Example: Extract Language Nodes



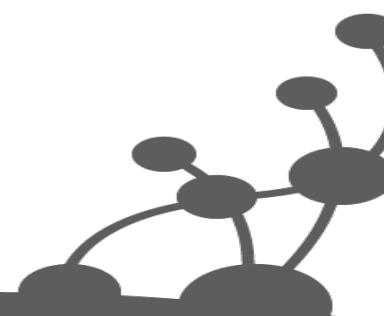
Exercise 6

Extract Node From Array Property



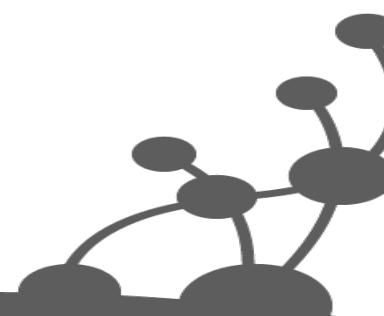
Your Turn

- Clean the database
- Execute *setup.txt*
- View the results
 - MATCH (n) RETURN n
- Execute *update-1.txt* repeatedly, until *numberRemoved* is zero
 - At each stage, view the results



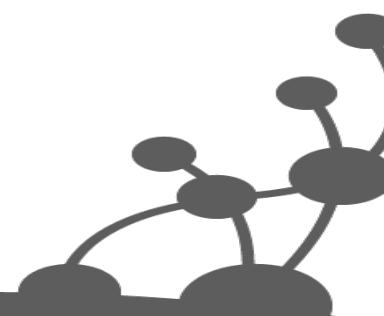
Extract Node From Array Property

```
MATCH (project:Project)
WHERE has(project.language)
WITH project LIMIT 2
FOREACH (l IN project.language |
  MERGE (language:Language{value:l})
  MERGE (project)-[:LANGUAGE]->(language))
REMOVE project.language
RETURN count(project) AS numberRemoved
```



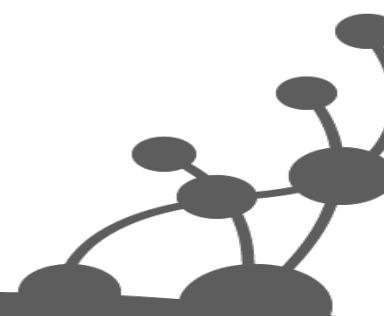
Select Batch of Nodes With Property

```
MATCH (project:Project)
WHERE has(project.language)
WITH project LIMIT 2
FOREACH (l IN project.language |
    MERGE (language:Language{value:l})
    MERGE (project)-[:LANGUAGE]->(language))
REMOVE project.language
RETURN count(project) AS numberRemoved
```



Loop Through Values in Array...

```
MATCH (project:Project)
WHERE has(project.language)
WITH project LIMIT 2
FOREACH (l IN project.language | 
    MERGE (language:Language{value:l})
    MERGE (project)-[:LANGUAGE]->(language))
REMOVE project.language
RETURN count(project) AS numberRemoved
```



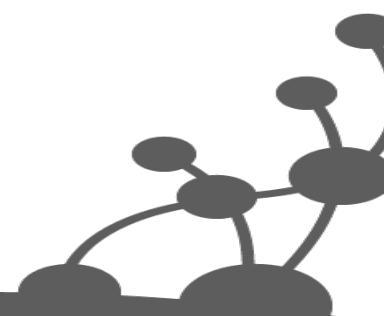
Create New Unique Node Per Value

```
MATCH (project:Project)
WHERE has(project.language)
WITH project LIMIT 2
FOREACH (l IN project.language |
  MERGE (language:Language{value:l})
  MERGE (project)-[:LANGUAGE]->(language))
REMOVE project.language
RETURN count(project) AS numberRemoved
```

Copy value from
current iteration

Relate Existing Node to Value Node

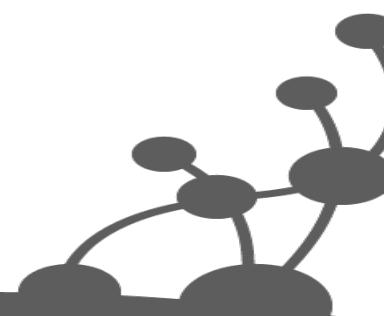
```
MATCH (project:Project)
WHERE has(project.language)
WITH project LIMIT 2
FOREACH (l IN project.language |
    MERGE (language:Language{value:l})
    MERGE (project)-[:LANGUAGE]->(language))
REMOVE project.language
RETURN count(project) AS numberRemoved
```



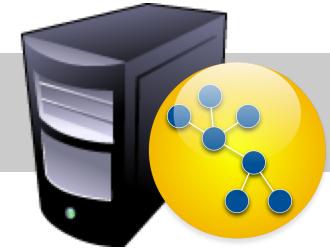
Remove Array Property

```
MATCH (project:Project)
WHERE has(project.language)
WITH project LIMIT 2
FOREACH (l IN project.language |
    MERGE (language:Language{value:l})
    MERGE (project)-[:LANGUAGE]->(language))
REMOVE project.language
RETURN count(project) AS numberRemoved
```

Repeat until
numberRemoved
is zero



Extract Node From Relationship

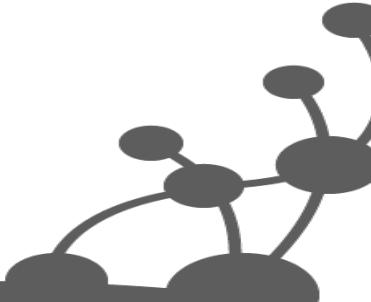


Problem

- You've modeled something as a relationship (with properties), but now need to connect it to more than two things

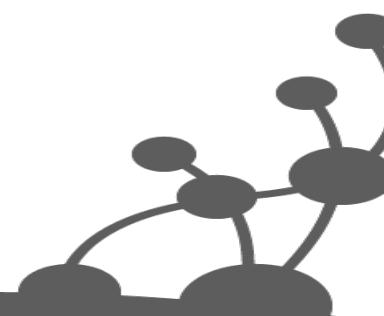
Solution

- Extract relationship into a new node (and two new relationships)
- Copy old relationship properties onto new node
- Delete old relationship

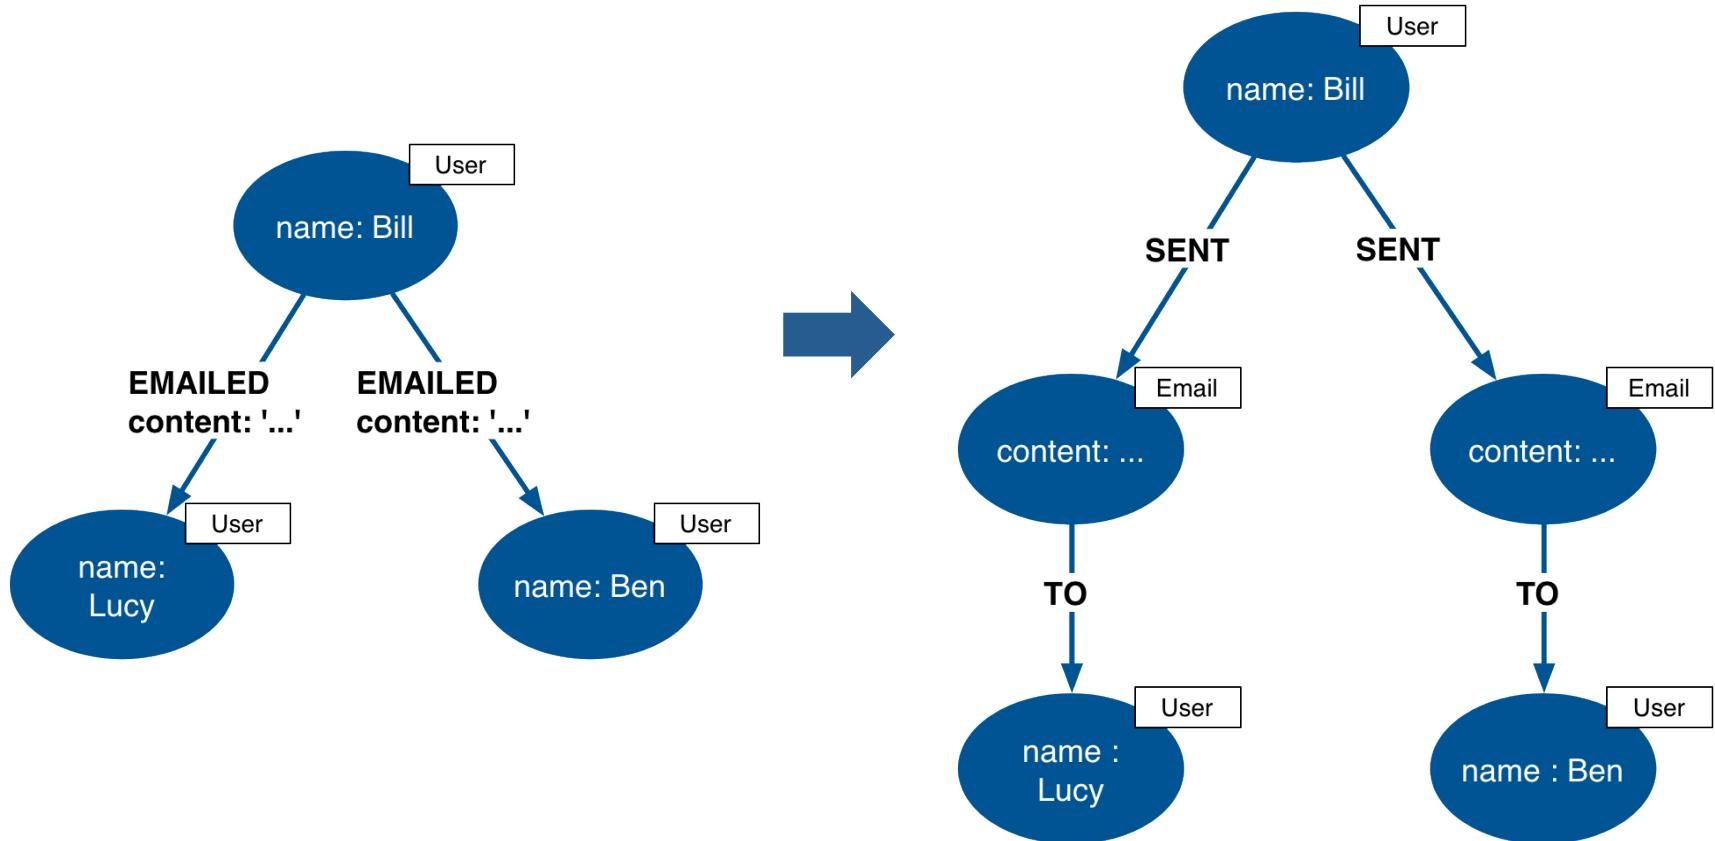


Exercise 7

Extract Node From Relationship

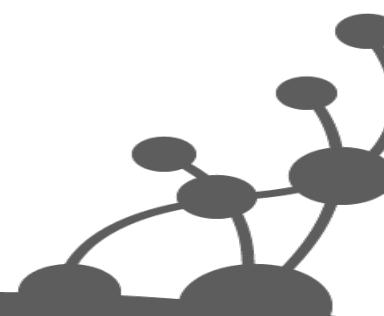


Example: [:EMAILED] to (:Email)



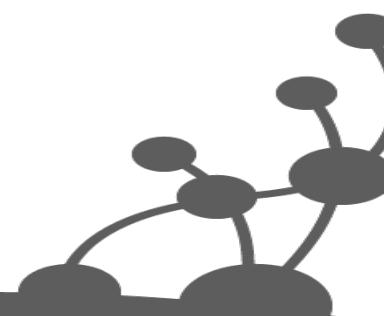
Your Turn

- Clean the database
- Execute *setup.txt*
- View the results
 - MATCH (n) RETURN n
- Execute *update-1.txt* repeatedly, until *numberDeleted* is zero
 - At each stage, view the results



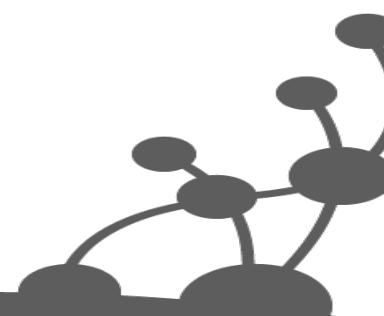
Extract Node From Relationship

```
MATCH (a:User)-[r:EMAILED]->(b:User)
WITH a, r, b LIMIT 2
CREATE (email:Email{content:r.content})
MERGE (a)-[:SENT]->(email)-[:TO]->(b)
DELETE r
RETURN count(r) AS numberDeleted
```



Select Batch of Relationships

```
MATCH (a:User)-[r:EMAILED]->(b:User)
WITH a, r, b LIMIT 2
CREATE (email:Email{content:r.content})
MERGE (a)-[:SENT]->(email)-[:TO]->(b)
DELETE r
RETURN count(r) AS numberDeleted
```

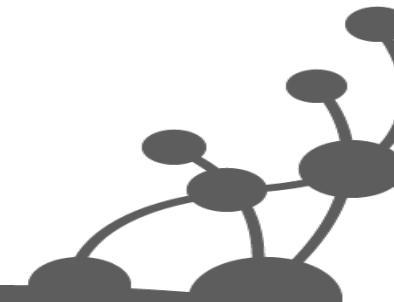


Create New Node and Relationships

```
MATCH (a:User)-[r:EMAILED]->(b:User)  
WITH a, r, b LIMIT 2  
CREATE (email:Email{content:r.content})  
MERGE (a)-[:SENT]->(email)-[:TO]->(b)  
DELETE r  
RETURN count(r) AS numberDeleted
```

Copy properties from old relationship

“Refactoring ID” ensures uniqueness



Delete Old Relationship

```
MATCH (a:User)-[r:EMAILED]->(b:User)
WITH a, r, b LIMIT 2
CREATE (email:Email{content:r.content})
MERGE (a)-[:SENT]->(email)-[:TO]->(b)
DELETE r
RETURN count(r) AS numberDeleted
```

Repeat until
numberDeleted
is zero

